

Morphological Stemming for Information Retrieval

Bob Carpenter
Alias-i, Inc.

March 21, 2007

Contents

Abstract	1
I Information Retrieval	2
1 What is Information Retrieval?	3
2 Modern Information Retrieval Engines	4
2.1 Document	4
2.2 Reverse Indexes	5
2.3 Matching Boolean Queries	6
2.3.1 Logical Operations on Queries	6
2.3.2 Phrasal and Proximity Queries	6
2.3.3 Document Ordering for Retrieval	6
2.3.4 Stemming in Boolean Queries	7
2.3.5 Pattern-Matching Queries	7
2.4 Matching Natural Language Queries	7
2.4.1 Term Vectors	8
2.4.2 Similarity Ranking	9
2.4.3 Cosine Ranking	9
2.4.4 Inverse Document Frequency Weighting	10
2.4.5 Coverage Weighting	10
3 Evaluating Information Retrieval Systems	12
3.1 Document Relevance to Queries	12
3.2 Precision, Recall and F-Measure	12
3.2.1 Precision	12
3.2.2 Recall	13
3.3 F-Measure	13
3.4 Mean Average Precision (MAP)	13

3.5	Precision/Recall Curves	13
3.6	Precision at n and Extrapolation Biases	14
II	Stemming	15
4	Natural Language Morphology	16
4.1	Orthography versus Phonology	16
4.2	Words, Stems and Modifiers	17
4.2.1	Inflectional Morphology	17
4.2.2	Derivational Morphology	18
4.2.3	Agglutinative versus Fusional	19
4.2.4	Synthetic versus Analytic	19
4.2.5	Incorporating versus Non-Incorporating	19
4.3	Realization of Morphemes	20
4.3.1	Prefixes and Suffixes	20
4.3.2	Infixing	21
4.3.3	Circumfixing	21
4.3.4	Vowel Changing	21
4.3.5	Templatic Morphology	22
4.3.6	Reduplication	22
4.3.7	Zero Morphology	22
4.3.8	Subtractive Morphology	23
4.3.9	Irregular Forms	23
4.4	Word Segmentation	23
4.4.1	Dictionary-Based Segmentation	23
4.4.2	Statistical Segmentation	24
4.4.3	n -Gram Indexing	24
4.4.4	Hybrid Systems	24
4.5	Stemming and Stoplisting	25
4.5.1	Stemming	25
4.5.2	Stoplisting	25
5	Evaluating Stemmers	26
5.1	How Much Stemming?	26
5.1.1	Etymological Stemming	26
5.1.2	Inflectional Stemming	26
5.1.3	Intentional Stemming	27
5.2	Evaluating Root Accuracy	28
5.3	Stemming as Clustering	28

5.4	Stemming Aggressiveness	30
5.5	The Problem of Word Sense	31
5.6	Case Study: Morpho Challenge 2005	31
6	Heuristic and Rule-Based Stemming	33
6.1	Prefix Stemming	33
6.2	The Lovins Stemmer	33
6.3	The Porter Stemmer	34
6.3.1	Step 1	34
6.3.2	Steps 2–6	35
6.4	Buckwalter’s Aramorph Arabic Stemmer	35
6.4.1	Tokenization	35
6.4.2	Word Segmentation	36
6.4.3	Dictionary Lookup	36
6.4.4	Compatibility Check	36
6.4.5	Analysis Report	36
6.4.6	Variant Spellings	37
6.5	Schulz et al., 2002	37
6.6	Ordered, Context-Sensitive Rewrite Rules	38
6.6.1	Feeding and Bleeding	38
6.7	Finite-State Transducers	39
6.7.1	Determinization and Ambiguity	39
6.7.2	Composition	40
6.7.3	Closure Properties	40
6.8	Two-Level Morphology	40
6.8.1	Two-Level Rules	41
6.8.2	Feeding in Two-Level Morphology	41
7	Previous Approaches to Statistical Stemming	43
7.1	Range of Morphology	43
7.1.1	Subsystems of Morphology	43
7.1.2	Morphological Operations	43
7.1.3	Quantity of Morphology	44
7.1.4	Ambiguity and Context Sensitivity	44
7.2	Types of Training Data	44
7.2.1	Supervised vs. Unsupervised	44
7.2.2	Word Lists, Dictionaries and Text Corpora	45
7.3	Problem Formulation	45
7.3.1	Word Segmentation	45
7.3.2	Extracting Stems and Affixes	45

7.3.3	Paradigm Extraction	46
7.3.4	Part of Speech Information	46
7.3.5	Clustering, Stemming and Similarity	46
7.4	Survey of Published Research	46
7.4.1	Harris, 1967	46
7.4.2	Hafer and Weiss, 1974	47
7.4.3	Krovetz, 1993	47
7.4.4	Church, 1995	49
7.4.5	Deligne and Bimbot, 1997	53
7.4.6	Jacquemin, 1997	55
7.4.7	Nevill-Manning and Witten, 1997	57
7.4.8	Déjean, 1998	57
7.4.9	Manning, 1998	58
7.4.10	Xu and Croft, 1998	60
7.4.11	Gaussier, 1999	62
7.4.12	Kit and Wilks, 1999	66
7.4.13	Yarowsky and Wicentowski, 2000	67
7.4.14	Yu, 2000	68
7.4.15	Goldsmith, 2001	70
7.4.16	Peng and Schuurmans, 2001	73
7.4.17	Schone and Jurafsky, 2001	74
7.4.18	Baroni et al., 2002	77
7.4.19	Creutz and Lagus, 2002	78
7.4.20	Neuvel and Fulop, 2002	79
7.4.21	Snover and Brent, 2002	80
7.4.22	Johnson and Martin, 2003	81
7.4.23	Argamon et al., 2004	82
7.4.24	Monson et al., 2004	84
7.4.25	Namer and Zeigenbaum, 2004	84
7.4.26	Wicentowski, 2004	85
7.4.27	Berhnhard, 2005	86
7.4.28	Bordag, 2005	89
7.4.29	Creutz and Lagus, 2005	90
7.4.30	Freitag, 2005	91
7.4.31	Goldwater and Johnson, 2005	94
7.4.32	Hu et al., 2005a	96
7.4.33	Hu et al., 2005b	98
7.4.34	Jordan et al., 2005	100
7.4.35	Keshava and Pitler, 2005	101
7.4.36	Chan, 2006	102

7.4.37	Hammerström, 2006	103
7.4.38	Karagol-Ayan et al., 2006	104
7.4.39	Xanthos et al., 2006	105

III Statistical Experiments 107

8 Corpora 108

8.1	Tokenization	108
8.2	Words	108
8.3	English Gigaword	108
8.4	Counts and Independence	109
8.5	Binomial Hypothesis Testing for Paradigms	110

9 Concatenative Stem/Suffix Model 113

9.1	Model Overview	113
9.1.1	Stemming Model	113
9.1.2	Maximum Likelihood Model Estimation	114
9.1.3	Expectation-Maximization (EM) Algorithm	114
9.2	Naive Suffixation Model	114
9.3	Stem and Suffix Models	115
9.3.1	Stem and Suffix Length Models	116
9.3.2	Character String Models	116
9.4	Training Data	117
9.5	Maximum a Posteriori Model Selection	117
9.5.1	Estimating Length Parameters	119
9.5.2	Estimating Character Language Models	121
9.6	Expectation Maximization (EM)	123

10 Affixation Boundary Effects 128

10.1	Problem for Pure Suffix Model	128
10.2	Modeling Boundary Effects	129
10.3	Results	130
10.4	Conclusions	130

IV Future Directions and Conclusions 131

11 Generalized Sequence Model 132

12 Recursive Categorical Model 134

13 Annotation Tool	136
14 Conclusions	140
References	141

Abstract

This document provides a survey of the use of stemming algorithms in information retrieval systems. We first describe the query ranking behind modern information retrieval engines, with an emphasis on evaluation. We then provide a high-level overview of natural language morphology and spelling. Next, we review the historical approaches to stemming using rule-driven heuristic stemmers such as the Porter stemmer and Soundex. The majority of the survey focuses on statistical stemming systems. We discuss both supervised and unsupervised approaches to the problem.

Part I

Information Retrieval

Chapter 1

What is Information Retrieval?

For the purposes of this survey, we employ the following definition:

Definition 1.0.1 (Information Retrieval Problem) *The information retrieval problem is that of returning a rank ordered (and optionally scored) set of documents in response to a user's query.*

The notion of document here is meant to be general; a document may be a single sentence, a passage, or even a whole collection of documents, with or without structure. The one thing all documents have in common is that they are made up out of textual components. In the simplest case, a document is just an unstructured sequence (or set of sequences) of characters.

Note that the return result is a rank ordered list of documents. It is not an answer to a question. Of course, a user may have a question in mind when they formulate a query.

In the simplest case, queries are just sequences of characters. In more complex systems, queries may also contain boolean constraints, boolean operators over subqueries, prefix matching, approximate matching proximity constraints, and so on. For instance, a query in WestLaw might look something like the following.

```
accident!  \S3 "gross negligence"
```

This query matches documents containing a term starting with the letters *accident* (e.g. *accidents*, *accident*, *accidental*, *accidentally*) appearing within three sentences of the phrase *gross negligence*.

Chapter 2

Modern Information Retrieval Engines

Almost all of the modern, large-scale information retrieval engines are organized around reverse indices of terms in documents. This holds for those that allow structured boolean queries and those that allow approximate “natural language” queries. We survey both varieties of search engine, with an emphasis on their underlying use of stemmed terms as the basis for reverse indices.

2.1 Document

Both boolean and natural language search engines are founded on the notion of a *term* in natural language. While there are no a priori constraints on what constitutes a term, typical choices involve words, pieces of words such as prefixes or n-grams, or simple transformations of words such as pronunciations.

For the purposes of this survey, we will consider a document to be an ordered collection of terms. In the simplest case, these terms are simply tokens found in a text. In the simplest tokenized case, tokens are defined starting from the left as the longest non-empty contiguous stretches of alphanumeric ASCII characters; that is, the successive matches of the regular expression:

$$\text{simpleTokenizer} = ([a-z] | [A-Z] | [0-9])^+$$

For instance, consider the following single-sentence document.

Smith-Barney’s released its earning reports for the first quarter,

in which it reportedly earned 3 cents/share.

The sequence of terms generated by the simple tokenizer is

⟨Smith, Barney, s, released, its, earning, reports, for, the, first, quarter, in, which, it, reportedly, earned, 3, cents, share⟩

As we will see later, it is typical to case normalize tokens, reduce them to a canonical (stem or lemma) form, and remove any remaining common function words. This would result in the following sequence of terms for the query above:

⟨ smith, barney, release, earn, report, first, quarter, report, earn, 3, cent, share ⟩

Note that the two forms of report, *reports* and *reportedly*, are both reduced to *report*. Also note that common words such as *its*, *for*, *the*, and *in* are removed; the tokenizer already removed punctuation.

Most of the rest of this document is concerned with how the tokenization, stemming and stopping processes are carried out. For now, we will simply assume a mapping from word sequences to sequences of underlying terms.

2.2 Reverse Indexes

All of the modern, scalable information retrieval systems are based on a mapping from terms to documents. Such an index reverses the natural mapping from documents to the terms it contains, and is thus called a *reverse index*. A reverse index usually includes a count of the term in a document. In order to support phrasal queries, such as our example "**gross negligence**", the position of the term in the document must also be stored. There is a substantial literature devoted to storing reverse indexes in compressed form on disk in such a way that resolving queries is efficient (see Witten et al.).

With stemming, the number of terms indexed in the reverse index is reduced compared to the strategy of just indexing every token. On the other hand, each term will now likely return more documents. Which setup is more efficient depends in part on the task (e.g. return the first 10 documents versus returning all the documents), and partly on the distribution of term occurrences in documents and queries.

As we discuss later, stemming provides accuracy benefits if users want to see the same documents for queries expressed as inflectional variants, such as *reports* and *report*.

By stoplisting, the absolute size of the index is reduced. This reduction is often substantial because of the high frequency of the stopped words. Stoplisting very frequent words is always a win in terms of computational efficiency. One problem with stoplisting is that if it is applied to the index as well as to queries, phrasal queries such as "King of France" cause problems, because it will be reduced to the term sequence $\langle king, france \rangle$. This may erroneously match documents containing phrases such as *England loved its King, but France was reeling from revolution.* will wind up removing the comma after *King* and the *but* before *France*, allowing it to match. This is a common problem with names on web search engines (take your pick, they all work like this), where a query such as "Fred Tyler" will match a list of names in another order, as in ... *Smith, Fred; Tyler, Sam; ...*, or just a list of first names, as in *Fred, Tyler and Sam....*

2.3 Matching Boolean Queries

Reverse indexes are sufficient to find all matches for *boolean queries*.

2.3.1 Logical Operations on Queries

A single term matches the set of documents that contain it. The logical conjunction (and) of two queries intersects their results. The logical disjunction (or) of two queries unions the results. Logical negation (not) simply complements the set of returned queries. In order to reduce the size of the result set, search engines restrict negative queries to differences. For instance, a difference query $q_1 - q_2$, matches documents that match q_1 but do not match q_2 .

2.3.2 Phrasal and Proximity Queries

If the indices also store document positions (e.g. paragraph, sentence and position-in-sentence), then they will be able to satisfy phrasal and proximity queries. Proximity queries require a pair of terms to occur within some distance of each other (a sentence, a given number of words, a paragraph, etc.).

2.3.3 Document Ordering for Retrieval

How efficient such queries are will depend on how quickly sets of hits can be intersected, unioned, and so on. Stemming causes result sets to be larger

and thus intersection/union operations to be slower. Usually, documents are stored in a canonical order to allow these operations to be simple linear merges. Such an order may be arbitrary (as in Lucene), or may be based on external metrics such as PageRank (as in Google), or date. By ordering documents by page rank and then by searching for the top-ranked documents, Google is able to avoid doing a full intersection for searches (which in the unmarked case, it treats as a large boolean conjunction, requiring all terms to exist on the page, or in text contained in links to the page).

2.3.4 Stemming in Boolean Queries

In a boolean setting, stemming has the following predictable effects. For a positive query (conjunction, disjunction or a must appear term), stemming increases the number of matching documents. For a negative query (negation or must-not-appear), it reduces the number of documents returned, because the set of documents excluded will be larger. For instance, a query **-earnings** will not match any document containing the words *earned*, *earn-ers*, etc., depending on how liberal the stemming process is.

2.3.5 Pattern-Matching Queries

Pattern matching queries typically work by expanding the query into a disjunction of terms. A prefix query, such as West's **"defen!"**, should match any document containing a term that starts with *defen*, such as *defend*, *defense*, *defenses*, *defensible*, *defendant*, *defender*, *defenders*, etc. The usual implementation strategy is to find all of the terms with the prefix and reduce the result to a disjunction.

Other pattern matching queries, such as Lucene's fuzzy queries, use other measures to pull terms out of indices. Lucene matches terms against fuzzy queries using a thresholded average edit distance per character. For instance, a fuzzy Lucene query like **" Smith"** might match terms *Smith*, *Smyth*, *Smythe*, *Smits*, *myth*, *Nesmith*, etc., depending on the matching threshold.

2.4 Matching Natural Language Queries

Natural language queries are simply those queries without (implicit or explicit) boolean structure. They are often, but not necessarily, expressed by users in the form of whole-sentence questions, such as the following example:

Who was the king of France in 1993?

Just to reinforce the point that search engines do not return answers, Google’s first hit for this query is a song called *King of France* written in 1993, and the 1993 film version of *The Three Musketeers*. Seeing the question mark, they recommend Google’s “Answers” service, which uses human searchers.

The goal of an information retrieval system in response to a natural language query is to return the documents ranked in order of relevance to the user’s query. In the best case, the query above would match relevant documents about France’s parliamentary system and lack of royalty, and maybe about regicide during the revolution.

It is much more difficult to scale systems to handle large numbers of natural language queries. The top web search engines (Google, Yahoo, Amazon) do not support them; they treat a sequence of words as a conjunction of boolean queries requiring each (non-stopped) word to appear. These large search engines also do not support proximity queries or approximate-match queries. Excite supported natural language queries in the past, but no longer does. Doug Cutting, the search architect for Excite went on to build the Apache Lucene search engine.

2.4.1 Term Vectors

The standard approach to scaling natural language queries is based on the following assumption.

Assumption 2.4.1 (Bag of Terms Assumption) *A document containing high counts of a large percentage of terms in a natural language query is more likely to be relevant than one containing a smaller percentage or lower counts of the query terms.*

With this guiding assumption, both documents and natural language queries are represented in modern information retrieval systems as bags of terms. A *bag*, also known as a multi-set, is just a set of elements with counts attached to the elements. We can thus think of a bag of terms as a *term vector*, with the vector dimensions corresponding to the set of terms. Although the number of terms may be infinite, typically the dimensionality is limited to the terms in the document collection, as all other terms tend to drop out of the equations.

2.4.2 Similarity Ranking

The final ranking of documents against a query is then reduced to generating a real-valued similarity measure between query term vectors and document term vectors. An obvious first approach to measuring the similarity of vectors $x = x_1, \dots, x_n$ and $y = y_1, \dots, y_n$ is to use a Euclidean measure of distance, as defined below.

$$d(x, y) = \left(\sum_i (x_i - y_i)^2 \right)^{1/2}$$

Another early approach was to use their dot product $x \cdot y$, as defined below.

$$x \cdot y = \sum_i x_i \cdot y_i$$

Both of these techniques suffer from a bias toward long documents. If a document is 1000 pages long and contains lots of counts for all of the words in the English language, it will have a relatively high score against just about any query.

2.4.3 Cosine Ranking

The simplest way to normalize for length is by only measuring the angle between two vectors, which is what vector cosine does by dividing by the vector lengths. The length of a vector x is written $|x|$ and defined using the usual Euclidean distance from the zero vector 0, which is zero in all dimensions:

$$|x| = d(x, 0) = \left(\sum_i (x_i - 0)^2 \right)^{1/2} = \left(\sum_i (x_i \cdot x_i) \right)^{1/2} = (x \cdot x)^{1/2}$$

Finally, the cosine is just the length-normalized dot product.

$$\cos(x, y) = \frac{x \cdot y}{|x| \cdot |y|}$$

Vectors running in the identical direction receive the maximal cosine value of 1.0 (e.g. $\cos(\langle 1, 1 \rangle, \langle 17.2, 17.2 \rangle) = 1.0$). Vectors running in opposite directions receive the minimal cosine value of -1.0 (e.g. $\cos(\langle 2, 2 \rangle, \langle -17.2, -17.2 \rangle) = -1.0$). Vectors running perpendicular (that is orthogonally, or at right angles) to one another receive the cosine value of 0.0 (e.g. $\cos(\langle 1, 0 \rangle, \langle 0, 1 \rangle) = 0.0$). Gerald Salton pioneered the use of cosine-based information retrieval, later realizing it in the SMART system (Salton 1968, Salton and Buckley 1975).

2.4.4 Inverse Document Frequency Weighting

As presented above, similarity is just the angle between term vectors. The entries in the term vectors are counts of the term in either a document or a query, and are known as *term frequency* (TF).

Term frequencies are often scaled in a way consistent with the following principle:

Definition 2.4.1 (Inverse Doc Frequency Assumption) *The utility of a term for search is inversely proportional to the number of documents in which it occurs.*

This paradigm is popular enough that the combination of term frequency, inverse document frequency and cosine scoring is usually simply referred to as *TF-IDF* scoring. To say there have been many formulations of IDF scores is an understatement. For instance, the Lucene search engine’s default scoring applies the following inverse-document frequency multiplier:

$$\text{idf}(t) = \left(1 + \log \left(\frac{\text{numDocs}}{\text{docFreq}(t) + 1} \right) \right)^2$$

where numDocs is the total number of documents in the collection, and docFreq(t) is the number of documents containing the term t . The addition of 1 to the document frequency is to make sure the minimal value is 1, so the denominator does not go to zero. The addition of 1 on the outside is so that the log doesn’t go to zero. The application of a logarithm tends to dampen the effects of large discrepancies in term frequency, whereas the final square applied on the outside has the opposite effect. There is no other theoretical motivation for this formula or any of its thousands of close relatives in the research literature.

Note that TF-IDF scoring is sensitive to features of the collection as a whole, and cannot be expressed as an operation on the term and document vectors alone. As a result, as the underlying collection changes, so may the search ordering of existing documents. It is particularly sensitive to stemming in the sense that the document frequency of a stem is typically much higher than that of an inflected form, particularly a rare one (e.g. the term *search* appears in 20 times as many documents as *searching* according to Google; *faith* is 500 times more likely than *unfaithfulness*.)

2.4.5 Coverage Weighting

With simple TF-IDF, the overall document score is heavily affected by one term showing up repeatedly in a document, lending a high TF-IDF score.

In fact, if there are two query terms t_1 and t_2 with the same IDF, then a document containing two instances of t_1 scores as well as one containing one instance of t_1 and one instance of t_2 . To put more emphasis on matching all terms, it is common to multiply by a coverage term. For instance, the Lucene search engine multiplies TF-IDF scores by the percentage of query terms that appear in the document (a term they call $\text{coord}(q, d)$ in the documentation for their **Similarity** class for mysterious reasons).

No matter which weighting scheme is chosen, the problems for stemming are similar.

Chapter 3

Evaluating Information Retrieval Systems

The information retrieval scoring metrics in common use are all based on the same idea of assessing a document’s relevance to a query.

3.1 Document Relevance to Queries

System scoring is built on top of the following general premise.

Definition 3.1.1 (Relevance and Scoring) *Documents are classified as either relevant or irrelevant to a query. The more relevant documents appear highly ranked, the better the score.*

This definition is almost as significant for what it leaves out, namely that there is no notion of one document being better or more relevant than another to a query in any of the standard IR system scoring metrics. Some systems do allow a “maybe relevant” category, but then either ignore those documents altogether or treat them as relevant in the scoring.

3.2 Precision, Recall and F-Measure

3.2.1 Precision

Given an assignment of relevance, a typical measure of an information retrieval system’s accuracy is *precision*, which is just the percentage of documents returned that are relevant. Often, precision is measured at predefined points, the top 10, top 50 and top 100 documents being typical numbers. For

example, precision at 100 is the fraction of the top 100 documents returned by the system that are relevant. The beauty of precision at n document scores is that human users need only evaluate n documents for relevance.

3.2.2 Recall

On the other hand, *recall* is the percentage of relevant documents found by a system. This is technically impossible to measure without knowing for each document in the collection whether it is relevant for the given query. This often leads to what we call *recall denial*. High precision systems demo well, and their lack of recall is essentially hidden by the high cost of evaluation. The only way such recall denial is found by users is if they know there is a relevant document in a collection which is not being found.

Recall is dual to precision in the sense that recall is the same as precision with the roles of the reference and response sets reversed.

3.3 F-Measure

In an attempt to reduce precision and recall to a single number, their geometric mean, known as the *F-measure*, is often used:

$$F_{\beta} = \frac{(1 + \beta^2) \cdot p \cdot r}{r + \beta^2 p}$$

The parameter β is used to control the balance of precision and recall; it's usually set to 1, which results in the so-called *balanced F-measure*.

3.4 Mean Average Precision (MAP)

Another metric used is mean average precision (MAP). The MAP score is the average of the precision scores computed for all documents returned from 0 to each relevant document. For instance, if a system returns d_0, d_1, d_2, d_3, d_4 and d_1 and d_2 are relevant, we compute the average of the precisions for two result sets d_0, d_1 and d_0, d_1, d_2 . Like precision at n , this measure favors results where relevant documents are ranked more highly than irrelevant ones.

3.5 Precision/Recall Curves

The same subsets as defined for MAP may be used to plot precision-recall curves. Each initial segment of results up to a relevant document induces a

precision/recall value. In TREC, these values are plotted with interpolation to guarantee they're monotonic.

The area under the precision-recall curve is another popular metric for information retrieval evaluations. It is closely related to the next metric we discuss.

3.6 Precision at n and Extrapolation Biases

The text retrieval conferences (TREC) attempt to get around the recall denial problem in a competitive evaluation by evaluating precision at 100 for dozens of systems. In a typical TREC evaluation, 20 groups submit query results consisting of 1000 documents in rank order. Humans then evaluate the top 100 return results from each group (a maximum of 2000 documents, but usually much less because of duplication). The relevant documents in all 20 top-100 lists are then considered the relevant documents in the collection.

These *known relevant documents* are then used to produce recall evaluations for all 1000 documents returned by a given group's systems. The resulting precision numbers at 1000 documents will be underestimates of true precision, as there may be relevant documents that are not known to be relevant because they were not returned in any group's top 100 list. Recall estimates at 1000 documents under this scheme may be biased either way, depending on how many relevant documents are not known. For instance, a system might return 75 documents in its top 1000 that are known to be relevant. Now suppose there are 175 known relevant documents. The TREC recall estimate is $75/175$, or about 42%. This estimate is biased to the high side if there are relevant documents that the system is missing, but that are not known to be relevant. The estimate is biased to the low side by any documents in the 101-1000 range are relevant but not in the known-relevant collection.

Part II

Stemming

Chapter 4

Natural Language Morphology

Morphology in the linguistics sense is the study of how words are formed. The notion of a *morpheme* is that of a minimal meaningful unit in a language. For instance, in the English word *unfaithfully*, we see morphemes *un* (meaning not), *faith* (meaning loyalty in this case), *ful* (converting a noun to an adjective), and *ly* (converting an adjective to an adverb).

4.1 Orthography versus Phonology

Linguists typically study morphology as it relates to the sounds of a language. The study of language sound is known as *phonology*. One reason for this is that not every language even has a written form. Another is that children acquire the morphology of a language before they learn to write. Furthermore, almost everyone speaks grammatically according to their culture's dialect, whereas many speakers are not able to write grammatically. All of this leads linguists to conclude that speech is more primary than writing. Search engines have, in fact, been built to do search over sound units, typically as uncovered by an automatic speech recognition system.

We are more concerned with search engines over written documents. The study of language writing is known as *orthography*. There has been much less attention paid to writing systems in the linguistics literature (but see Sproat ??). In practice, we are not even concerned with written characters per se, because we are dealing with electronic documents that provide a well-defined encoding of written characters as sequences of bytes (e.g. ASCII, Big5 or Unicode UTF-8).

4.2 Words, Stems and Modifiers

Although we can define morphemes as minimal meaningful units, the higher-level notion of “word” is much more subtle. In languages such as English, words are typically considered to be the units derived from writing that have spaces in between them. Of course, if you’re studying sound, this is a little trickier to pin down. Languages like Chinese are written without spaces between units, so they can not look to their written language for a definition of words. In practice, there are varying opinions as to how words should be construed in Chinese; the 2006 SIGHAN word segmentation bakeoffs presented four word segmentation standards that varied in how they treated compounds and some morphemes. Other languages, like Hebrew, don’t even bother writing the vowels, thus rendering the written forms of words more ambiguous than the spoken ones.

4.2.1 Inflectional Morphology

The world’s languages typically mark their nouns, verbs and often other parts of speech with grammatical information. These words often have the same core meaning, with different realizations depending on grammatical context. For instance, English noun inflection involves case marking (e.g. nominative (*he*) vs. accusative (*him*)), number marking (e.g. singular *engine* vs. plural *engines*), gender marking (e.g. masculine *he* vs. feminine *she* vs. neuter *it*). English verb inflection involves tense (e.g. past *ran* versus present *runs*), aspect (e.g. infinitive *eat*, past perfect *eaten*, and present progressive *eating*).

Inflectional classes are typically organized into *paradigms*, which are basically tables of the ranges of inflectional variants of words in the class. For instance, the common noun paradigm in English (as opposed to the pronoun paradigm) only varies in number, whereas the regular verbal paradigm (as opposed to the auxiliary verb paradigms) includes variation in tense, aspect and number.

Whether the suffixes and other markings used to mark inflectional paradigms count as morphemes or not is a contentious issue in linguistics. They are unlike other morphemes in that inflectional marking is not optional. Furthermore, in most contexts, only one instance of a paradigm is well formed due to grammatical *agreement*, where verbs and nouns (and often adjectives, determiners and other elements) must have the same marking. For example, we find *he runs* (singular nominative noun, singular verb), not *he run* (singular noun, plural verb) or *him runs* (singular accusative noun,

singular verb). Luckily, the morphemic status of inflection is not an issue we need to resolve in order to perform effective search.

For each inflectional class, a representative is designated as a canonical instance, often known as a *lemma*. Quite frequently, lemmas are chosen to be the “simplest” member of a class. For English verbs, lemmas are usually taken to be the infinitive form (e.g. *to write*), and for nouns, usually the singular form (e.g. *egg*). In general, the lemma is just a stand in for its full paradigm, and need not be a legal fully-inflected form.

Many search engines attempt to only stem based on inflectional morphology. For instance, the only stemming in the English versions WestLaw is of plural and possessive nouns to their singular forms; Google, on the other hand, does no stemming at all in English.

4.2.2 Derivational Morphology

Derivational morphology is differentiated from inflectional morphology by the fact that derivation stems are (a) optional, (b) don’t just mark a finite set of paradigmatic differences, and (c) change the underlying meaning of the root to which it applies. For instance, the English prefix *un-* attaches to adjectives and adverbs to denote the lack of the property denoted by the root, as in *unwelcome* versus *welcome*; whereas the verbal prefix *re-* means something done again, as in *rewritten* versus *written*. Examples of grammatical category changing derivational morphology include the suffix *-ly*, which converts an adjective into an adverb (e.g. *professional* to *professionally*), the suffix *-ation* which changes a verb to a noun (e.g. *distill* to *distillation*), the suffix *-al* which turns a noun into an adjective (e.g. *fiction* to *fictional*), or the suffix *-ness* which turns an adjective into a noun (e.g. *happy* to *happiness*). These operations typically stack freely, so that the suffix *-ize*, which converts an adjective to a verb, may apply to the derived adjective *fictional* to produce *fictionalize*, which may in turn take the *-ation* suffix to produce *fictionalization*.

For languages with extensive derivational morphology, dictionaries do not typically list inflected forms. The exception is cases where there is meaning drift of a derived form. For instance, *surveillance* has come to have a more specific meaning than its literal compound meaning of the act of surveying. One challenge for information retrieval systems is to not overly conflate derived forms during searches by stemming them to common underlying roots. Otherwise, a search for **electronic surveillance** will find many articles on questionnaire-type surveys delivered by computer and electronic devices for carrying out land surveys.

4.2.3 Agglutinative versus Fusional

Languages which append sequences of morphemes together in a productive fashion are said to be *agglutinative*. Examples include Turkish and Basque.

Languages for which morphemes are hard to separate out from words are said to be *fusional*. Latin and most of the Indo-European languages are examples of at least partly fusional languages.

4.2.4 Synthetic versus Analytic

Languages with a low morpheme to word ratio are said to be *analytic*. Analytic languages are often said to be *isolating*, because their concepts don't mix in words. Examples of analytic languages include Chinese, Thai and Vietnamese. In analytic languages, there is very little inflectional or derivational morphology.

Languages with higher morpheme to word ratios are said to be *synthetic*. Thus agglutinative languages tend to be synthetic. Languages which involve a great deal of synthesis are said to be *polysynthetic*. The typical example of a synthetic language is Latin, though the Germanic languages are also highly synthetic.

Romance languages tend to be synthetic when it comes to pronouns, allowing them to be expressed with the verb as clitics. On the other hand, the Germanic languages tend to be synthetic, with separate words for pronouns. Sometimes languages allow both possibilities. English comparatives may be expressed either synthetically (e.g. *tastier*) or analytically (e.g. *more tasty*).

4.2.5 Incorporating versus Non-Incorporating

Languages which allow multiple lemmas in a single word are said to be *incorporating*. The languages studied by morphologists tend to be polysynthetic, incorporating and agglutinating, like Chukchi and classical Ainu. For instance, in classical Ainu, the word *aejajkotujmasiramujukpa* ("I keep swaying my heart afar and toward myself over *X*") is synthesized from a sequence of nine morphemes *a-e-jaj-ko-tujma-si-ram-suj-pa*, two of which are lexical stems ("far" and "heart").

In English, we have noun compounding to produce terms such as *towel rack* or *towel rack designer*, or even *towel rack designer training course instructor*. In German, these compounds are written without spaces, and with boundary effects such as epenthetic vowel insertion. For instance, *Erdbeertorte* is a combination of *Erdbeere* (strawberry) and *torte* (cake). A

more extreme example is *Ueberseed Deutsch-lehrerinternetmailinglistenfragenstellundantwortkundigen* (People well versed in asking questions and supplying answers on the Internet Mailing List of German teachers abroad). Agglutinative morphology of this form presents roughly the same problem as breaking raw Chinese text into words.

4.3 Realization of Morphemes

Morphemes are realized differently in different languages, as we saw in the synthetic versus analytic distinction above. In this section, we consider some of the more common modes of morpheme expression.

4.3.1 Prefixes and Suffixes

The most common form of morphemes are simply synthetic prefixes and suffixes. For instance, English plurals for most words are simple suffixes, with singular *paper* and plural *papers*. Unfortunately, creating a plural in English isn't quite as easy as attaching an *s*, either in spelling or in pronunciation. The problem arises that languages have constraints on syllable structure and complexity. For instance, the word *box* cannot have an *s* attached to the end to form a single syllable in English (though for some reason *sixth* can). In this case, an *epenthetic vowel* is inserted to restore syllable structure, resulting in *boxes*; such epenthetic vowels are schwas in that their pronunciation is that of a reduced (non-stressed) vowel. In addition to pronunciation effects realized in spelling, there are pure spelling effects with morpheme attachment. Consider the relation between *sky* (singular) and *skies* (plural) and between *boy* (singular) and *boys* (plural). In English, a final *y* pronounced as a single vowel is respelled as *ie* before attaching the plural suffix. There are also irregular pluralizations, such as *man* (singular) and *men* (plural), which involve vowel changes (see section 4.3.4).

Another common boundary effect is consonant doubling, as in the relation between *run* and *running*. Yet another is “unpronounced” vowel dropping, as in the relation between *write* and *writing*.

The relation between spelling and sound in languages like English is tenuous at best. For regular plural nouns, there are three phonemic cases: unvoiced *s* as in *kits*, voiced *s* as in *kids*, voiced *s* plus epenthetic, as in *bosses*.

Prefixes behave just like suffixes, but attach to the front rather than the rear of a word. Languages tend to be divided as to whether their morphology is predominately suffixing or prefixing. For instance, Basque, despite

being highly synthetic, involves exclusively suffixing. In English, grammatical category-changing morphemes are always final, whereas prefixes tend to have semantic but not syntactic effects (e.g. *un-*, *re-*, *pre-*, etc.)

4.3.2 Infixing

Many languages present morphemes which are neither simple suffixes nor simple prefixes. One example is infixing. An *infix* is inserted into the interior of a word. For instance, *-um-* in Tagalog appears before the first vowel in a word to nominalize it to an agent (e.g. *sulat* (writing) vs. *s-um-ulat* (one who wrote)).

Infixes in other languages are often peripheral in their effects, as in Tagalog (see Sproat 1992). In the language Ulwa, infixes such as *-ka-* and *-ni-* mark number and person. The position of infixes in Ulwa is prosodically conditioned to occur after the first metrical foot (a unit of prosody consisting of a stressed syllable and one or more unstressed syllables).

4.3.3 Circumfixing

A *circumfix* is a pair of sequences that wrap around words. An example is Indonesian *ke-* *-an*, which wraps around an adjective to produce a nominal form (e.g. *besar* (big) vs. *ke-besar-an* (bigness)), and another is German *ge-* *-t*, which forms a past participle by wrapping around a verb (e.g. *spiel-en* ([to] play) vs. *ge-spiel-t* ([has] played)).

4.3.4 Vowel Changing

In addition to prefixes, suffixes, infixes and circumfixes, morphological productions can induce word-internal vowel changes. In English, we see this with both nouns (e.g. *goose* vs. *geese*, *man* vs. *men*) and verbs (e.g. *take* vs. *took*). Latin and other Romance languages (e.g. Spanish) typically involve regular vowel changing.

Traditionally, there was a distinction between so-called *item-and-arrangement* theories and *item-and-process* theories of morphology. The item-and-arrangement idea is that free morphemes combine to produce bigger words. The item-and-process notion is more general, allowing a base form to be transformed into an inflected form by a more general process (such as vowel change).

Binyan	Active (A)	Passive (UI)	Template	Gloss
I	<i>kAtAb</i>	<i>kUtAb</i>	CVCVC	write
II	<i>kAttAb</i>	<i>kUtAb</i>	CVCCVC	cause to write
III	<i>kAAAtAb</i>	<i>kUUtAb</i>	CVVCVC	correspond
VI	t <i>kAAAtAb</i>	t <i>kUUtAb</i>	tVCVVCVC	write to each other
VII	n <i>kAAAtAb</i>	n <i>kUUtAb</i>	nCVVCVC	subscribe
VIII	<i>ktAtAb</i>	<i>ktUtAb</i>	CtVCVC	write
X	st <i>kAtAb</i>	st <i>kUtAb</i>	stVCCVC	dictate

Figure 4.1: Paradigm for root *ktb* in Arabic

4.3.5 Templatic Morphology

Arabic and other Semitic languages construct words by overlaying orthogonal constraints on the consonant sequence, the vowel sequence and a syllable template. Because the surface realization is driven from the template, this approach is often known as *templatic morphology*.

Figure 4.1, copied from (Sproat 1992:51) shows the paradigm for a single verb, *ktb*, which fixes the consonant sequence. The syllable templates, known as *binyanim*, make up the rows of the table, and contribute to the meaning of the derived form. Note that a template may contain additional symbols, such as the initial **n** in binyan VII.

The surface forms are derived through a combination of constraints on the interaction of the three levels of constraint: consonants, vowels and syllabic template. Each of these constraints contributes morphological meaning: root, voice and binyan respectively.

4.3.6 Reduplication

Many languages, including Chinese, allow widespread morphological *reduplication* (the reason for the *re* in the name is a mystery). In Indonesian, *orang* means man and *orang orang* is the plural form, meaning men.

Reduplicated forms are not always exact duplicates. Sometimes the duplication is only of the a prefix, as in Warlpiri (see Sproat 1992), in which *wantimi* is the base verb meaning to fall, whereas *wanti+wantimi* is the non-past form.

4.3.7 Zero Morphology

Some theoreticians have stipulated *zero morphemes*, which are morphemes with no overt realization. Sproat (1992) mentions the verbalization of nouns

(punningly known as “verbing”) in English, and provides the following examples: *book a flight*, *table the motion*, *xerox this article*).

4.3.8 Subtractive Morphology

Sproat (1992) points out that some languages allow what appears to be *subtractive morphology*. In this situation, an inflected form is assumed to remove material from its base form. He provides the Muskogean language Kosati as an example, in which the singular/plural distinctions are *pitaf+fi+n/pit+li+n*, *lasap+li+n/las+li+n*.

We believe this problem stems as much from assuming that singulars provide the canonical lemma form as anything else. If the role of plural and singular were reversed, this would look like additive morphology. Another way around the problem is to consider the lemma form to be abstract *las*, then both forms, singular and plural, appear consistent with an additive model.

4.3.9 Irregular Forms

A final variety of morphological realization is a completely irregular form. For instance, the forms of the English auxiliary *to be* are highly irregular (e.g. *are*, *am*, *is*, etc.). There is often a positive correlation between the frequency of a lexical item and its irregularity. For instance, the English pronoun system is also irregular (e.g. *she* (feminine nominative) versus *her* (feminine accusative)).

4.4 Word Segmentation

Segmenting languages like Chinese and Japanese, which are written without spaces, presents unique challenges to word-based information retrieval systems.

Similar problems arise within words in highly agglutinative languages. For instance, if German noun compounds are to be considered matches for queries involving one of their nouns, then the noun compounds need to be accurately segmented in word-based information retrieval systems.

4.4.1 Dictionary-Based Segmentation

Segmentation systems often use large dictionaries of known words, taking longest matching sequences working through the characters. A strategy is

required in this case for sequences that do not match any dictionary entries.

For languages such as German, which have morphological boundary effects, dictionary matching schemes must take these into account.

4.4.2 Statistical Segmentation

Although it is possible to write statistical tokenization systems that are 97–98% accurate at finding word boundaries in Chinese text, they are not *stable* in the sense of always producing the same tokens from the same sequence of characters; instead, which tokens are produced depends on context. For instance, a statistical tokenizer might produce two tokens c_1c_2 c_3c_4 for the four-character Chinese query in one context, and only one token in another context. For instance, in English, the characters *runstuck* can be tokenized as either *runs+tuck* or as *run+stuck*; if *he* preceded the word, we might prefer *he+runs*, and if *they* preceded the word, we might prefer *they+run* as the tokenization. System errors on tokenization would result in serious problems for both recall (missing lots of documents with the correct token) and precision (finding erroneous documents with the incorrect token).

4.4.3 n -Gram Indexing

Search engines often do not even try to index words, but rather simply index n -grams of characters (e.g. see Lucene’s **CJKAnalyzer** for Chinese, Japanese and Korean). Under this scheme, a Chinese word $c_1c_2c_3c_4$ consisting of four characters would generate the 1-grams c_1 , c_2 , c_3 , and c_4 , the 2-grams c_1c_2 , c_2c_3 , and c_3c_4 , and the 3-grams $c_1c_2c_3$ and $c_2c_3c_4$.

Experiments in English have shown character 3-gram indexing to be as effective in terms of precision and recall in TREC-like evaluations as token-based indexing (?? 19??).

4.4.4 Hybrid Systems

It is possible to use more than one indexing scheme and combine their results. For instance, MSN reports good performance from a combination of large dictionary and unigram indexing for Chinese (Nie et al. 2000). Although not quite as effective as bigrams, unigrams resulted in a smaller index under MSN’s indexing scheme.

4.5 Stemming and Stoplisting

4.5.1 Stemming

The operating hypothesis of information retrieval engines that use stemming is the following.

Definition 4.5.1 (Stemming Hypothesis) *Documents containing words with stems that are stems of words in the query should match more highly.*

Evaluating this hypothesis has proved fairly tricky given the available testing resources, for all the same reasons as IR evaluations are tricky in general, mainly that too much human effort is involved around answering the vaguely stated question “is this document relevant to this query”.

Formal results using existing resources have shown at best a marginally small benefit from stemming English queries. In some domains, such as biomedical text, applying simple stemmers has detrimental effects. In highly inflected languages, stemming has helped with recall.

4.5.2 Stoplisting

The operating hypothesis behind stoplisting is the following.

Definition 4.5.2 (Stoplisting Hypothesis) *High-frequency function words do not matter when matching a query.*

Even if we do not believe the stoplisting hypothesis, we may use stoplisting simply as an efficiency measure. High frequency words such as English *the* take up a great deal of space to store in reverse indexes and also return large hit-sets when merging results.

Chapter 5

Evaluating Stemmers

Without an end-to-end application, it is difficult to decide how to evaluate a stemmer.

5.1 How Much Stemming?

The immediate question arising is that of how to decide how much stemming should be done.

5.1.1 Etymological Stemming

At one extreme, we could assume that every pair of words with a shared etymological root should have the same stem. Consider the words in figure 5.1, all of which derive etymologically from **author** (this set was drawn from the *New York Times* section of the Gigaword corpus (with dozens of misspellings removed, but alternate British/American spellings retained)

Although it has a similar prefix, **authentic** is not etymologically related to **author**.

This level of conflation seems far too aggressive for the purposes of information retrieval; someone searching for preauthorization information is unlikely to care about authorhood in the same search.

5.1.2 Inflectional Stemming

One approach to circumscribing the stemming problem in a linguistically motivated way is to only allow inflectional morphology. This would rule out variants for **author** such as **authorized** and **authorization**, because the first is the third-person singular verb inflection of **authorize**, whereas

antiauthoritarian, antiauthoritarianism, antiauthority,
 author, authoratative, authoratatively, authordom, authored,
 authoress, authoresses, authorhood, authorial, authoring,
 authorisation, authorised, authorises, authoritarian,
 authoritarianism, authoritarians, authoritative,
 authoritatively, authoritativeness, authorities, authority,
 authority, authorization, authorizations, authorize,
 authorized, authorizer, authorizers, authorizes, authorizing,
 authorless, authorly, authors, authorship, authorships,
 coauthor, coauthored, coauthoring, coauthors, cyberauthor,
 deauthorized, multiauthored, nonauthor, nonauthoritarian,
 nonauthorized, preauthorization, preauthorizations,
 preauthorized, quasiauthoritarian, reauthorization,
 reauthorizations, reauthorize, reauthorized, reauthorizes,
 reauthorizing, semiauthoritarian, semiauthorized,
 superauthoritarian, unauthorised, unauthoritative,
 unauthorized

Figure 5.1: Words with etymological root **author**

authorization is a noun derived through the derivational process of nominalization. Note that the base verb **authorize** is itself derived from the noun **author**. We return to two evaluations that compare stemmers with and without derivational components.

Inflectional stemming seems far too weak for the purposes of information retrieval; someone looking for **authorize** is most likely to be happy with responses for **authorization**.

5.1.3 Intentional Stemming

The holy grail of stemming is to capture the intentions of searchers. The requirement is that words should be conflated by stemmers if they should be conflated during search. As we will see later, this is a difficult question to answer on anything other than a word-by-word basis.

Dictionaries, although very incomplete with respect to live data, typically separate out words with different meanings and often list other words with similar meanings in the listings. For instance, **coauthor** has its own entry in the *American Heritage Dictionary*, but that entry includes **coauthored**, **coauthoring**, and **coauthors**. The entry for **author** also includes the plural **authors**, the verbs **authoring**, and **authored**, and also the

	<i>Response Same</i>	<i>Response Different</i>
<i>Reference Same</i>	True Positive (TP)	False Negative (FN)
<i>Reference Different</i>	False Positive (FP)	True Negative (TN)

Figure 5.2: Confusion Matrix for Stemming as Equivalence Relation

derived adjective **authorial**. There is a main entry for the verb **authorize**, which includes inflected forms **authorized**, **authorizing** and **authorizes**, as well as the derived nominal **authorizer**. And while **authoritarian** contains its own entry, **antiauthoritarian** and other derived forms are simply not in the dictionary. Unfortunately, dictionaries have different standards on how to separate meanings and how many to list, reflecting that this is basically a subtle semantic judgement call.

5.2 Evaluating Root Accuracy

Suppose we have a reference that indicates for each word what its intended stem is. A simple way to evaluate is accuracy. That is, what percentage of words are given the correct stem by a stemmer relative to the reference stemming.

5.3 Stemming as Clustering

Usually we don't care about the stems themselves *per se*. We are only interested in the question of whether two words have the same stem or not. So we define an equivalence relation over words where two words are equivalent if and only if they have the same stem. An equivalence relation partitions the set of words into disjoint classes, with the classes consisting of sets of words with the same stem.

There are a range of traditional evaluations that compare reference (e.g. gold standard) partitions with response (e.g. system) partitions; see (Jain and Dubes 1988; Jain et al. 1999). Most of these measures are computed from the simple confusion matrix shown in figure 5.2. For instance, the true positive count is the number of pairs of words that are equivalent in both the reference and response relations. False negatives are pairs of words that the reference considers equivalent but the response fails to conflate; false positives are words that the response conflates that are not equivalent in

the reference. True negatives are pairs that are equivalent in neither the reference nor the response relation. Note that we include comparisons of a word with itself; otherwise, singleton clusters would have no impact on scoring.

The simplest measures to compute are simple precision and recall type numbers.

$$\text{precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

Precision is the percentage of pairs that are equivalent in the response that are also equivalent in the reference.

$$\text{recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

Recall is the percentage of pairs that are equivalent in the reference that are also equivalent in the response. Note that recall is equivalent to precision with the roles of reference and response reversed.

There are two ways to combine results into a single number. The first is raw accuracy, which counts the number of correct decisions:

$$\text{accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

F-measure provides a single measure combining precision and recall as usual:

$$F_{\beta} = \frac{(1 + \beta^2) \cdot \text{recall} \cdot \text{precision}}{\text{recall} + \beta^2 \cdot \text{precision}}$$

The term β indicates the weight bias toward precision, with the balanced F measure being F_1 :

$$F_1 = \frac{2 \cdot \text{recall} \cdot \text{precision}}{\text{recall} + \text{precision}} = \frac{2 \cdot \text{TP}}{2 \cdot \text{TP} + \text{FN} + \text{TN}}$$

A traditional alternative to the F-measure for equivalence scoring is the Jaccard coefficient:

$$\text{jaccardCoeff} = \frac{\text{TP}}{\text{TP} + \text{FN} + \text{FP}}$$

The only difference from F-measure is the weight of the true positives. Other metrics in this family have been proposed with multipliers of 0.5.

Like precision, recall and F-measure, the Jaccard coefficient does not consider true negatives. The following two components of the receiver operating characteristic (ROC) curve do take the false negatives into account. Specifity measures the ability to reject false negatives:

$$\text{specifity} = \frac{\text{TN}}{\text{TN} + \text{FP}}$$

The dual measure to specifity is known as selectivity:

$$\text{selectivity} = \frac{\text{TN}}{\text{TN} + \text{FN}}$$

Note that specifity and selectivity are dual to recall and precision (respectively) with the notion of positive and negative reversed. That is, the specifity of a response equivalence against a reference equivalence is the same as the recall of the complement of the response equivalence against the component of the reference equivalence.

With these numbers we can further compute values such as the κ chance-adjusted agreement statistic, but this is not much use in comparing systems against each other over the same task, as it only varies from accuracy by a constant. Also, because in this task the number of true negatives completely overwhelms the number of true positives (by many orders of magnitude), κ and accuracy will only show significant changes in fairly deep decimal places.

If the clusterings are richer than simple partitions, more sophisticated evaluations may be possible. For instance, if the response consists of a probability estimate for pair equivalence, we can provide expected precision, recall and F-measures. If the response contains hierarchical structure, we can evaluate different thresholds of breaking them into pieces.

5.4 Stemming Aggressiveness

Paice (1996) introduced an evaluation method for stemmers that is almost like precision and recall. His overstemming index is just $(1 - \text{precision})$, except with the count of words equivalent to themselves being removed from the true positive count. His understemming index is similar for $(1 - \text{recall})$.

Paice introduces a notion of stemming weight, where a “light” stemmer does fairly minimal stemming whereas a “heavy” stemmer provides aggressive stemming. Paice defines stemming weight as the ratio of overstemming to understemming. Unfortunately, this statistic does not really capture the intension of the definition, which is to measure the aggressiveness of a stemmer. As the number of understemming errors approaches zero, the metric

tends to infinity; as the number of overstemming errors approaches zero, the metric tends to zero. As both approach zero, the value tends toward 1; for an error-free system the value is 0/0. A more interesting feature of Paice’s article is that he provides two reference clusterings, one corresponding to “light” stemming and one to “heavy” stemming.

A better measure of aggressiveness is simply the number of positive system responses. Lennon et al. (1981) proposed a simple measure based on the number of clusters. An alternative would be the average size of clusters.

5.5 The Problem of Word Sense

Even if we can define the notion of semantic equivalence well enough to make a sensible reference partition of words into related classes, we are left with the residual problem of word sense ambiguity. A word may be equivalent to another word under one sense, but not under another. For instance, the word **can** may be used three different ways: as a noun (meaning a kind of container), as an auxiliary verb (meaning a kind of ability), or as a verb (derived by zero morphology from the noun and meaning the process of putting something into the container kind of can). The auxiliary verb **can** would be related to past tense **could**, whereas the noun **can** would be related to its plural **cans**, whereas the verb **can** would be related to its present participle form **canning**. Note that the closely related word **cane** is similarly ambiguous between several meanings (a noun meaning a stick, a noun meaning a plant that’s like a stick, and the verb meanings of making a chair out of cane or flogging with a cane). Thus it is not clear whether the various verb forms should be related to the various noun forms. Sugar cane is not particularly related to corporal punishment.

The only way to tackle the problem of word sense would be to evaluate word tokens in context, not word types.

5.6 Case Study: Morpho Challenge 2005

In 2005, Mikko Kurimo, Mathias Creutz and Krista Lagus organized a competition (Kurimo et al. 2005) organized around unsupervised morphology learning.

Participants were provided with word lists with frequencies for Finnish, Turkish and English, as well as examples of the linguistically motivated “gold standard” reference results.

Of the 12 participating groups, 7 of them were from Leeds. We survey the systems that introduced new techniques in the literature section below.

There were two components to the evaluation. The first was a direct comparison to a linguistically motivated morphemic segmentation “gold standard”. The gold standard was exact stems as derived from a two-level morphology analysis. The second evaluation involved speech recognition with models composed of the generated morphemes rather than words, and is thus much harder to connect to the morphology itself, especially as the organizers ran the language modeling part of the bakeoff.

Chapter 6

Heuristic and Rule-Based Stemming

The “classical” approach to stemming is to write a system of “rules” that can be applied to strings to stem them. Rule-based systems are either written directly in code (e.g. the Porter stemmer), or are written in a higher level rule language (e.g. two-level morphology and regular transducers).

6.1 Prefix Stemming

The simplest stemmer is one that stems a word to its first k letters, or itself if it is less than k letters long. Surprisingly, this form of stemming is remarkably similar in performance to the more complex stemmers (Paice 1996).

6.2 The Lovins Stemmer

Lovins (1968) introduced the notion of stemming. Her stemmer was based on a list of 260 suffixes. The stemmer simply removes the longest matching suffix from a word. This straightforward approach has been shown in many applications to do nearly as well as the more sophisticated Porter stemmer we turn to next.


```

private final void step1()
{
    if (b[k] == 's')
    {
        if (ends("sses")) k -= 2; else
            if (ends("ies")) setto("i"); else
                if (b[k-1] != 's') k--;
    }
    if (ends("eed")) { if (m() > 0) k--; } else
    if ((ends("ed") || ends("ing")) && vowelinstem())
    {
        k = j;
        if (ends("at")) setto("ate"); else
        if (ends("bl")) setto("ble"); else
        if (ends("iz")) setto("ize"); else
        if (doublec(k))
        {
            k--;
            {
                int ch = b[k];
                if (ch == 'l' || ch == 's' || ch == 'z') k++;
            }
        }
    }
    else if (m() == 1 && cvc(k)) setto("e");
}
}

```

Figure 6.1: Step 1 of Porter's Stemmer

6.3 The Porter Stemmer

The most commonly applied stemmer for English is the Porter stemmer, named after its designer (Porter 1980). The Porter stemmer consists of a series of analysis stages, each of which inspects a suffix of the current state of the string and possibly applies a transformation to remove a particular suffix. In total, the Porter stemmer uses about 60 or so suffixes. Here's a description of the various steps from the Porter's current Java port of the system, which may be found at:

<http://www.tartarus.org/martin/PorterStemmer/>

6.3.1 Step 1

Removes plurals and *-ed* and *-ing* suffixes. The code is provided in Figure 6.1. The undeclared variables are all globals, and functions such as `ends()` have the obvious interpretation. It is fairly easy to see what's going on. The first test is if the last character is *s*. If it is, then if the ending is *sses* it is replaced with *ss* (e.g. converting *grasses* to *grass*, if it is *ies*, it's replaced with *i* (converting *skies* to *ski*), and otherwise, if it's not *ss*, the *s* is simply

removed (leaving *grass* untouched, but converting *cans* to *can*). The second test checks to see if the current string ends in *eed*; if it does and if there is at least one vowel before, the final *d* is removed (converting *agreed* to *agree*, but leaving *greed* untouched). The next set of tests are only used if the ending is *ed* or *ing*, in which case, if the remaining suffix is *at*, an *e* is inserted (converting *sated* to *sate*). The penultimate test handles the doubled consonants, one of which is removed unless the consonant is *l*, *s* or *z* (converting *fitted* to *fit*, but *filled* to *fill*).

6.3.2 Steps 2–6

The remaining stages of the Porter stemmer are similar in both form and concept.

Step 2 Turns terminal *y* to *i* when there is another vowel in the stem.

Step 3 Maps double stems to single ones (e.g. strings with suffix *-ization* to strings with suffix *-ize*). This is done with a fixed list of 21 compounds tems indexed by penultimate character rather than a loop to remove arbitrary pairs.

Step 4 Similar to step 3 with a different set of cases conditioned on the final letter (e.g. suffix *-icate* replaced by *ic*).

Step 5 Remove suffixes *-ant*, *-ence* etc. in cases where the word is long enough in terms of number of vowels.

Step 6 Removes a final *-e* suffix if the resulting stem is long enough.

6.4 Buckwalter’s Aramorph Arabic Stemmer

Ad-hoc stemmers have been written for many languages. A leading example that has found wide use is Tim Buckwalter’s Arabic stemmer, Aramorph (Buckwalter 2004). Aramorph operates either directly on Arabic characters, or by means of Buckwalter’s own transliteration to Roman characters.

6.4.1 Tokenization

The first stage of processing is tokenization. Tokenization is quite simple, taking maximal contiguous sequences of Arabic characters. There are a few side-conditions to deal with issues such as confusing word-final suffixes, use of characters for punctuation, and run-on words.

6.4.2 Word Segmentation

The word segmentation stage explores all possible splits of a word into three components. This is done quite simply by taking all choices where the prefix is 0–4 characters, the suffix is 0–6 characters, and the stem is 1 or more characters long. Among the segmentations of **wAlgAz** are $\langle \mathbf{w}, \mathbf{AlgA}, \mathbf{z} \rangle$, $\langle \epsilon, \mathbf{wAlg}, \mathbf{Az} \rangle$, $\langle \mathbf{wAl}, \mathbf{gAz}, \epsilon \rangle$, and $\langle \epsilon, \mathbf{wAlgAz}, \epsilon \rangle$.

6.4.3 Dictionary Lookup

For each word segmentation, Aramorph checks to see if the putative prefix, suffix and stem are in the lexical dictionary. To this end, Aramorph makes extensive use of three dictionary files, one for prefixes, one for stems, and one for suffixes. For each entry, the lexicons provide the entry without short vowels and diacritics and associate a part-of-speech tag. Other information, such as English glosses and entries with short vowels and diacritics are included, but not used in stemming.

An example of an entry in the prefix dictionary is $\langle \mathbf{wAl}, \mathbf{NPref-A1} \rangle$ (gloss: *and+the*). The part-of-speech tag for the prefix indicates that it's a noun prefix of subcategory **A1**. An example of an entry in the stem dictionary is $\langle \mathbf{ktAbxAn}, \mathbf{NapAt} \rangle$, where the category **NapAt** is that of a noun with subcategory **apAt** (gloss: *library/bookstore*). The stem dictionary also maps each form to a unique lemma form. For instance, entries **ktAbxAn** and **ktbxAn** both map to the lemma *kutubiy~_1*.

6.4.4 Compatibility Check

If all three components are in their respective dictionary, then they are checked for compatibility. Compatibility requires each pair of prefix/stem, stem/suffix and prefix/suffix to be compatible with each other.

Aramorph uses three *compatibility tables*, indicating valid combinations of prefix/stem, stem/suffix and prefix/suffix. For example, the prefix/stem table contains an entry $\langle \mathbf{NPref-A1}, \mathbf{N-ap} \rangle$, meaning that the prefix **wAl**, listed as category **Npref-A1**, is compatible with nouns of category **N-ap**.

6.4.5 Analysis Report

If all three compatibility checks pass, the resulting analysis is a possible result. Because the word segmentation pass is ambiguous, there might be more than one valid analysis of the input string.

6.4.6 Variant Spellings

Aramorph attempts to mitigate the existence of orthographic variation and non-canonical spellings by trying to vary commonly varied suffixes if there are no analyses found in the first stage. For instance, **Ely** and **Ely** return identical analyses, namely the proper noun **Ali** and the preposition **à la**.

6.5 Schulz et al., 2002

Schulz et al. (2002) propose an approach that varies from stemming by returning all of the components of a word, not just the stem. They claim this applies to inflectional morphology, providing **search+ed** and **search+ing** as examples, and derivational morphology, with examples **search-ers** and **search-able**. They then note that the same approach will work for compounds, citing the German **Blut-hoch-druck** [high blood pressure].

Schulz et al. are interested in indexing biomedical data, which for English, provides its own forms which are much rarer in the language as a whole (e.g. **-itis** in the disease domain and **-ase** in the protein domain). Several specialized medical morphosemantic processors are cited in the paper, ranging in scope from **-itis** forms to terms for surgical procedures, collectively referred to as “neo-classical” compounds due to their Greek and Latin roots (McCray et al. 1988).

Schulz et al. note that words like **diaphysis** should be decomposed into **diaphys-is** rather than the more etymologically motivated **dia-phys-is**, their point being that the precision will be much higher in the former given the prevalence of roots **dia** and **phys**. This analysis is based on the assumption that the information retrieval engine will be carrying out independent per-token retrieval, rather than say, using n-grams of units.

Schulz et al. evaluate over a corpus of German clinical medicine handbook (about 2.5 million tokens). First, they normalize orthography. For example, they normalize umlauts by mapping **ä** to **ae**, and they normalize for ambiguity of Latinate spellings in German by mapping **ca** to **ka**.

Their approach to stemming Then they extract a number of dictionaries, which they use in a longest-match heuristic morphology process. They create lists of stems (which they call subwords), prefixes, infixes, derivational suffixes and inflectional suffixes. They treat short words of four characters or less separately, and also break out acronyms as non-decomposable. They then use these dictionaries in a longest-match fashion.

For search, they don’t just use single subwords, but use locality to prefer documents containing **append-ectomy** over **append-ic-itis** and **thyroid-**

itis over thyroid-ectomy for a query append-ectomy and thyroid-itis.

6.6 Ordered, Context-Sensitive Rewrite Rules

Chomsky and Halle (1968) introduced an approach to morphology based on ordered, context-sensitive rewrite rules. A *context-sensitive rewrite rule* is of the form

$$\alpha \rightarrow \beta / \gamma _ \delta$$

where α, β, γ and δ are character sequences. Such a rule is read as saying α may be replaced by β if α is preceded by γ and followed by δ . Thus $\gamma _ \delta$ represents the *context* and $\alpha \rightarrow \beta$ the *rewrite*.

Chomsky and Halle’s scheme requires additional assumptions before it is well-defined. In the simplest case, the set of rules is linearly ordered by precedence, and each rule is obligatory. Then in each “cycle”, the list of rules is scanned and the first applicable rule applied, and then the cycle is repeated until a fixed-point is reached in which further cycles do not transform the input. Of course, this scheme is wildly underconstrained. For instance, it’s easy to write rules that either oscillate or grow strings indefinitely. Further refinements to Chomsky and Halle’s original proposal often relaxed assumptions by making some rules optional, requiring only a partial order, limiting the number of cycles, and so on.

We illustrate the application of such rules with an example from (Karttunen and Beesley 2001). Consider the following two context-sensitive rules:

$$\begin{aligned} N &\rightarrow m / _ p \\ p &\rightarrow m / m _ \end{aligned}$$

In the first rule, the left context is empty, and in the second rule, the right context is empty. Karttunen and Beesley note that given input **kaNpat**, the first rule applies to produce intermediate result **kampat**, then the second rule applies to produce the final result **kammat**.

6.6.1 Feeding and Bleeding

Because rules are ordered, the output of a rule may turn on or turn off a rule ordered after it. One rule is said to *feed* another rule if its output allows the second rule to apply. A rule is said to *bleed* another rule if the reverse is true.

Antworth (1991) provides the following example of a pair of rules where the first feeds the second. The first rule is *vowel raising*:

$$e \rightarrow i / __\text{CONSONANT}^*i$$

and the second is *palatalization*:

$$t \rightarrow c / __\text{i}$$

The letters all stand for themselves with CONSONANT^* referring to a sequence of zero or more consonants. Antworth notes that if vowel raising is ordered before palatalization, the input *temi* produces *timi* by application of the first rule (*mi* matches CONSONANT^*i), which in turn produces *cimi* by application of the second rule (*i* matches *i*).

6.7 Finite-State Transducers

Johnson (1972) realized that not only rules, but whole cycles in Chomsky and Halle's scheme could be formalized by means of a single finite-state transducer.

A *finite-state transducer* (FST) is a finite-state automaton where each transition is labeled with a pair of symbols rather than a single symbol. For convenience, one of the symbols is treated as input and the other as output. In order to transduce an input sequence to an output sequence, the input sequence is matched against the automaton using only the input symbols. As this matching takes place, the output symbols are accumulated from the transitions.

6.7.1 Determinization and Ambiguity

Although it is possible to determinize a finite-state transducer if both the input and output are known, residual transduction ambiguity may remain. For instance, a very simple transducer might have a single state with an arc back to itself labeled **a:a** and another labeled **a:b**. In this case, an input *a* may be transduced into either an output *a* or output *b*. For inputs consisting of *n* repetitions of the letter *a*, there are 2^n possible outputs!

The set of all input/output pairs that satisfy the transducer determines a *regular relation* *R* such that $R(\sigma, \tau)$ if and only if the FST transduces σ to τ .

6.7.2 Composition

Johnson's FST compilation technique depends on first representing each rule as a finite-state transducer. For instance, the first rule above may be represented by the following transducer:

Initial State: 1

Final States: 1, 3

Transitions: (1,1,x:x) (x != N)
(1,2,N:m)
(2,1,p:p)
(1,3,N:N)
(3,1,x:x) (x != p)

The automaton itself encodes the obligatory nature of the rule.

The composition of R and R' , traditionally written $R \circ R'$, is defined by taking $R \circ R'(\sigma, \rho)$ if and only if there is some τ such that $R(\sigma, \tau)$ and $R(\tau, \rho)$. Schützenberger (1961) proved that if R and R' are relations defined by FSTs, then the relation $R \circ R'$ is also defined by some FST.

It's relatively straightforward to directly construct the composition. The states of the composed transducer consist of pairs of states, one drawn from the first automaton and one from the second. There is an arc from (q_1, q'_1) to (q_2, q'_2) labeled $a : c$ if and only if there is some symbol b such that there is a transition from q_1 to q_2 in R labeled $a : b$ and a transition from q'_1 to q'_2 in R' labeled $b : c$.

6.7.3 Closure Properties

FSTs may define arbitrary finite relations. In addition, they are closed under union, concatenation and Kleene star (repetition). In all of these cases, the construction is identical to the ordinary finite-state automaton case. Unlike the finite-state language case, finite-state relations defined by FSTs are *not* closed under negation or intersection.

6.8 Two-Level Morphology

Koskenniemi (1983) proposed an alternative finite-state architecture for morphological rule systems. Rather than cascading finite-state transducers one after the other, Koskenniemi's system applied a set of finite-state

transducers in parallel, each expressing a relation between the input and output (typically called the *lexical level* and *surface level* in the two-level morphology literature).

6.8.1 Two-Level Rules

Sproat (1992) classifies the four types of rules used in Koskenniemi's system, expressing them in the notation of (Dalrymple 1997):

An *exclusion rule*:

$$a : b / \Leftarrow \text{LC} _ \text{RC}$$

prohibits an input a from producing output b in the specified context LC $_$ RC.

As before, LC is the left context and RC the right context. In two-level morphology, the contexts are specified as regular expressions over context pairs (note that this is not as general as making them regular relations).

A *context-restriction rule*:

$$a : b / \Rightarrow \text{LC} _ \text{RC}$$

prohibits input a from producing output b in any context other than the one specified.

A *surface coercion rule*:

$$a : b \Leftarrow \text{LC} _ \text{RC}$$

requires input a to produce an output b in the specified context.

A *composite rule*:

$$a : b / \Leftrightarrow \text{LC} _ \text{RC}$$

is a conjunction of context restriction and surface coercion; in other words, input a is realized as surface b if and only if it is in the specified context.

6.8.2 Feeding in Two-Level Morphology

Antworth (1991) notes that the feeding of palatalization by vowel raising may be captured in two-level morphology with the following pair of rules.

Palatalization is captured by the composite rule:

$$t : c / \Leftrightarrow _ @ : i$$

which states that t is realized as c if and only if it is followed by a pair of which the output is i . The symbol $@$ is a *wildcard* and matches any symbols.

Vowel raising is captured by the composite rule:

$$e : i / \Leftrightarrow _ (\text{CONSONANT} : \text{CONSONANT})^* @ : i$$

which states that input *e* becomes output *i* if and only if they precede an arbitrary sequence of consonant pairs followed by a pair the output of which is *i*.

Along with an identity rule that allows an input character to produce itself on output, these rules constrain the input sequence *temi* to produce the output sequence *cimi*. It is easy to see that this mapping satisfies both rules. The *t/c* input/output pairing has its context satisfied by the following *e/i* pairing, which has *i* as its output. The *e/i* pairing is itself licensed by the *m/m* and *i/i* pairings, which matches $(\text{CONSONANT} : \text{CONSONANT})^* @ : i$.

Chapter 7

Previous Approaches to Statistical Stemming

There have been at least as many approaches to statistical stemming as there have been papers written on the topic. In this section, we consider some of the main variations among systems, focusing on the model of the problem being solved and the method of estimating models. We briefly summarize the main modeling decisions and then proceed with a paper by paper survey.

7.1 Range of Morphology

7.1.1 Subsystems of Morphology

Some approaches have only considered inflectional verb morphology, although most neither distinguish inflectional from derivational morphology, nor distinguish among the word classes to which morphemes apply (e.g. nouns, verbs, or adjectives).

7.1.2 Morphological Operations

The range of morphological operations considered in statistical models has varied from very simple single stem concatenative models to ones that estimate circumfixes, point-of-affixation stem changes and stem-internal vowel changes.

We are not aware of any statistical approaches to the complex paradigmatic infixing of Arabic or Hebrew.

7.1.3 Quantity of Morphology

Although languages allow multiple affixes to stack on a single stem, most of the approaches we survey restrict their analyses to single suffixes. Some allow an affix as well as a suffix. Very few attempt to model multiple stems and/or affixes in a single word.

7.1.4 Ambiguity and Context Sensitivity

The same sequence of characters may correspond to different words in different contexts. For instance, the word *reflex* might refer to the root meaning a muscle reaction, or might be analyzed as the prefix *re-* applied to the verb *flex*. Thus any approach to morphology that treats words in isolation is bound to fail in general on all but one variant of a word. That is, we must choose to stem *reflex* to either *reflex* or *flex*.

More generally, a word like *runs* may be either a third person singular verb (e.g. *he runs*) or it may be a plural noun (e.g. *they scored three runs*). This may cause problems for systems that restrict attention to nouns or verbs.

Names pose a particular problem. Consider the nominal *singer* derived from the root verb *sing* versus names such as *Singer*. Similar problems occur with the names *John* and *Johns* and the common nouns *john* (a euphemism for toilet) and *johns* (its plural as well as a kind of underwear if preceded by the independent word *long*).

7.2 Types of Training Data

7.2.1 Supervised vs. Unsupervised

In the case of *supervised training*, a corpus of morphologically annotated words is used to induce a system for carrying out morphological analysis. In the unsupervised case, only a corpus of the text is available. This distinction is orthogonal to the other dimensions of the statistical stemming problem.

Because there are so few labeled corpora, most of the work in this area has been done with little or no supervised training.

Most of the unsupervised approaches have used heuristic approaches to bootstrap morphology rather than more mathematically rigorous forms of missing data imputation such as expectation/maximization (EM) or Gibbs sampling.

7.2.2 Word Lists, Dictionaries and Text Corpora

Many approaches begin with lists of words in a language. These may be extracted from dictionaries or from raw text. With dictionaries, it is common to have additional information such as pronunciation, example usages and parts of speech. Some dictionaries even consider morphological variants, but their lists are never comprehensive.

Sometimes, the underlying words also have associated frequency counts over a corpus.

The most general approach considers words in context as they appear in a text corpus.

7.3 Problem Formulation

In this section, we consider ways in which the statistical morphology problem has been formulated. Most of these formulations are approximations to the most general problem of providing a full analysis of a word into stems and affixes with part-of-speech information.

7.3.1 Word Segmentation

Although not typically considered part of the problem of morphology induction, many approaches have been applied to breaking unsegmented text into words. This kind of approach has been applied to Chinese, English with spaces removed, and most widely, to sequences of phonemes or other representations of sounds. Very often these approaches can be tuned to generate segmentations at a finer grain than words (e.g. syllables or bits thereof), or at a coarser grain (e.g. phrases or common compounds or collocations). Some approaches have generated a general hierarchical clustering of symbols, ranging from segments to phrases.

With all of these approaches, the question remains as to where to draw morpheme boundaries.

7.3.2 Extracting Stems and Affixes

Many of the approaches we discuss only extract a collection of stems and affixes from a corpus, without attempting to analyze individual words. Often these extracted lists are ranked.

7.3.3 Paradigm Extraction

Some approaches we will discuss treat the problem as one of extracting inflectional paradigms, generalizing the traditional notion to derivational paradigms, as well. In most of these cases, the paradigm is represented as a set of stems that fall into the paradigm along with the set of morphemes that apply to them.

7.3.4 Part of Speech Information

The morphological effect of a derivational suffix often changes the underlying stem's part of speech. For instance, the suffix *-ness* converts the adjectival stem *happy* into the nominal *happiness*; *-ly* converts it to the adverbial *happily*.

7.3.5 Clustering, Stemming and Similarity

Although our aim is to produce a stemmer that can take an arbitrary word as input and produce a stem, many statistical approaches to morphology have only attempted to induce a clustering over a finite set of words in a fixed corpus. Two words are in the same cluster if they have the same stem.

Many of the clustering-based approaches are based on estimating a similarity measurement between words. With a similarity measure, any of the standard hierarchical clustering methods may be used (complete-link, single-link and average-link).

More generally, similarity may be measured between a word and a cluster or between two clusters. With this approach, top-down clustering approaches may be used such as k-means or its generalization, soft model-based clustering (also known as EM clustering). We have not seen this kind of clustering in the statistical morphology literature.

7.4 Survey of Published Research

In this section, we survey the major papers published in the field of statistical morphology and word induction. The papers are presented in order of publication, not significance.

7.4.1 Harris, 1967

In perhaps the earliest widely cited work in automatic morphology induction, Harris (1967) suggested a method of inspecting successor frequencies (the

number of words that start with a given prefix) and predecessor frequencies (the number of words that end with a given suffix). For instance, for the prefix *wor-*, we find successors *word*, *worded*, *worm*, *worn*, etc. For the suffix *ning* we find predecessors such as *run*, *swoon*, *light*, etc.

7.4.2 Hafer and Weiss, 1974

Hafer and Weiss (1974) carried out a range of experiments studying Harris's notion of finding stems and affixes by inspecting predecessor and successor frequencies. They test hypotheses by setting decision-tree like thresholds, such as successor frequency plus predecessor frequency being above a threshold, or both being above independent thresholds.

Hafer and Weiss introduce a heuristic that favors cutting at a point that forms an independent word. That is, we would only allow *wordy* to be split into *wor+dy* if the sequence of characters *wor* formed a word (which it does not). This heuristic clearly under-filters, as many stems have prefixes which are themselves proper stems (e.g. the word *be* is the prefix of hundreds of words, and *ace* is the prefix of words such as *acetate*, *acetylene*, and *acetaminophen*).

In addition to Harris's suggested frequencies, they also consider estimated conditional entropy rates as the basis of cutoffs. The entropy of the conditional probability distribution $p(c|c_0 \cdots c_{n-1})$ for fixed $c_0 \cdots c_{n-1}$ is given by:

$$H(c_0 \cdots c_{n-1}) = - \sum_c p(c|c_0 \cdots c_{n-1}) \cdot \log_2 p(c|c_0 \cdots c_{n-1})$$

Note that the summation is over all possible next characters c , and given the probability weighting, the result is just the expected negative log probability, often written $E(-\log_2 p(c|c_0 \cdots c_{n-1}))$, with the averaging over c selected according to $p(c|c_0 \cdots c_{n-1})$ being implicit.

Most often, the entropy is bounded by above with an estimator \hat{p} of the true probability distribution, such as the maximum likelihood estimator.

As the number of equally likely outcomes increases, entropy increases. Entropy also increases as the distribution becomes more uniform. Although Hafer and Weiss only looked at characters, it is straightforward to also look at the entropy of the set of sequence continuations of a given prefix.

7.4.3 Krovetz, 1993

Krovetz (1993) introduces a suffix stripping stemmer that makes use of a dictionary in two ways to increase sensitivity to semantic meanings of words.

Krovetz uses the *Longman Dictionary of Contemporary English* to help control the over- and under-stemming problems arising from the Porter stemmer. Krovetz noticed that the Porter stemmer erroneously conflated the pairs:

university/universe, negligible/negligent, arm/army,

whereas it failed to conflate the variants:

analysis/analyzes, explain/explanation, urgency/urgent.

Krovetz's first stemmer is derived from a dictionary and a list of suffixes for a language. At run time, it applies a fairly straightforward suffix-stripping algorithm along the lines of Porter's. The difference is that before stripping a suffix, the dictionary is checked, and if it contains the word, the word is returned as the stem without further suffix stripping. For instance, Krovetz found 229 words that ended in **e** and were different than the same word with the **e** removed, such as:

suite/suit, fine/fin

Because of the ability to stop at a relevant dictionary item, Krovetz's stemmer could be more aggressive than Porter's in removing suffixes. Krovetz cites the word **suited** as one where it is not clear without the dictionary whether to stem it to **suit** or **suite**.

Krovetz includes a number of hand-tuned heuristics. For instance, words ending in **-ies** first have their **-s** stripped, and if the result is in the dictionary, that is returned; otherwise the **-ies** is replaced by **-y**, and the algorithm continues. This heuristic allows **calories** to stem to **calorie** and **fisheries** to stem to **fishery**.

For the derivational stemmer, Krovetz cites the following suffixes to be most frequent:

-er, -or, -ion, -ly, -ity, -al, -ive, -ize, -ment, -able,

Krovetz also uses a dictionary to allow more aggressive stemming than would be found with suffix-stripping in some cases. He notes that the definition of **cylindrical** contains the word **cylinder**, providing evidence they should be conflated.

Krovetz's experimental results over four collections to show that stemming generally improves performance. The largest collection was from West and contained less than 12,000 documents. The average query and document lengths in the collections were 7.1/43, 8.9/581, 9.6/3262, and 13.0/62;

the query lengths, in particular, are much longer than those found in typical web queries. Krovetz found particularly strong improvements due to stemming in the two collections with short documents.

Krovetz includes another heuristic that requires words to have the same initial three characters before attempting to stem them.

The West collection was much larger than the other collections (at only 36MB!). It included far more cases where spelling errors prevented a word from being stemmed. Even with this modest size corpus only 9% of the words in the corpus (counted by type, not token) appeared in the dictionary. On the other hand, 70% of the words in the dictionary appeared in the corpus. Longman's dictionary is very short, being intended for learners of English. Krovetz found the words whose stems were not in the dictionary tended to be proper names (e.g. *Algeria*) or technical terms (e.g. *exosphere*).

Krovetz evaluated two versions, one that only stripped inflectional suffixes using the dictionary to stop the process, and a second that also stripped derivational suffixes and used dictionary meanings. Although the results were broadly similar between the two versions, Krovetz found that at higher levels of recall, derivational morphology became more important.

Krovetz also noted that word sense can have an effect on stemming. For example, **media** is the plural of **medium** in the entertainment sense, but not in the spiritualist sense. He also notes that part-of-speech can play a role, because **intimate** may be an adjective, which is related to the adverb **intimately**, or may be a verb, in which case it is related to the nominal **intimation**.

Krovetz notes that dictionaries often contain so-called “run ons”, which are lists of suffixes that may apply to a word. This information may be used to help stemmers.

Finally, Krovetz poses a problem in the form of variant spellings. For instance, Briths and American English spellings vary, something with transistors is said to be **transistorized** in American English and **transistorised** in Briths English. And then there are just variant spellings, such as **judgement** and **judgment**.

7.4.4 Church, 1995

Church (1995) measures the correlation of pairs of words within documents over a given document collection. This measure could directly inform a language-model based information retrieval system, because if two words are positively correlated, then knowing that one shows up in a document provides positive evidence that the other shows up in the document. If they

	<i>+hostages</i>	<i>-hostages</i>	<i>Total</i>
<i>+hostage</i>	619	479	1098
<i>-hostage</i>	648	78,233	78,881
<i>Total</i>	1267	78,712	79,979

Figure 7.1: Contingency table for *hostage* versus *hostages*

are not correlated at all, they are independent, and knowing if one shows up indicates nothing about the other. The problem with using this information in practice is that the pairwise data is too sparse to gather the requisite estimates.

Term Correlation

Correlation may be computed with as simple an input as a confusion matrix over the document collection. To make this more concrete, consider the singular form *hostage* and its plural variant *hostages*. Church reports the contingency table in Figure 7.1 for these two words over a modestly sized collection of 79,979 documents. This says that 619 documents contained both terms, 479 documents contained *hostage* and not *hostages*, 648 contained *hostages* but not *hostage*, and 78,233 contained neither word. Church wants to test whether the two terms are independent of each other, and the answer in this case is a definitive *no*. Simple maximum likelihood estimates show $p^{\text{ml}}(\text{hostage}) = 1098/79,979 = 0.0137$, $p^{\text{ml}}(\text{hostages}) = 1267/79,979 = 0.0158$. If the terms were independent, we'd expect the joint probability to factor to the product of the independent word probabilities, $p(\text{hostage}, \text{hostages}) = p(\text{hostage}) \cdot p(\text{hostages})$. But in fact, the cooccurrence of the two terms is about 30 times higher than is expected if they were independent.

Church measures relatedness by a coarse-grained correlation coefficient derived from the collection contingency table. The formula for correlation adjusts the difference between the binomial probability estimate of the pair and the probability estimate given independence, adjusted for the variance of the term probabilities:

$$\rho(\omega_1, \omega_2) = \frac{p^{\text{ml}}(\omega_1, \omega_2) - p^{\text{ml}}(\omega_1) \cdot p^{\text{ml}}(\omega_2)}{\sigma(p^{\text{ml}}(\omega_1)) \cdot \sigma(p^{\text{ml}}(\omega_2))}$$

where the standard deviation $\sigma(p)$ is defined to be the square root of the

variance:

$$\sigma(p) = \sqrt{p \cdot (1 - p)}$$

Note that we have normalized everything down to a Bernoulli trial; correlations are often expressed in terms of empirical counts.

A more fine-grained approach would compute paired correlation over the documents. Another approach would be to do a binomial hypothesis test of independence. All of these approaches will sort pairs of terms by independence, although they measure it with different sensitivities.

Setting aside independence-testing methodology, the question remains as to whether dependent terms should be reduced to a single term. Church reports words the following words as having correlations between 0.4 and 0.5 (large correlations) between their singular and plural versions:

hostage(s), reactor(s), rebel(s), guerrilla(s), abortion(s), delegate(s)

On the other hand, the following pairs have near-zero correlations:

await(s), ground(s), possession(s), boast(s), belonging(s), compare(s)

Church's point is that if a document contains the word *hostage*, it's likely about hostages, whereas the word *await* is much less topic specific.

Church recognizes that nouns are better keywords than verbs in most cases. This may affect the correlation between a word and its *+s* form. A word like *await* acts only as a verb, whereas a word like *hostage* only acts as a noun, with words like *delegate* acting as both. In fact, Church reports that the term *delegate* and its plural have a 0.44 correlation in 1988 (a U.S. presidential election year, and the data was collected over AP newswire stories) versus 0.18 in 1989. As a verb, *delegate* seems much less specific than as a noun.

Church measures not only *-s* suffixes, but also *-ly* suffixation and capitalization. Highly correlated stem/*-ly* variants (correlations above .15) include:

sexual(ly), *racial(ly)*, *quarter(ly)*, *mental(ly)*, *alleged(ly)*

whereas low correlations (negative, but very close to zero) include:

New(ly), *hot(ly)*, *sole(ly)*, and *Week(ly)*

For capitalization versus non-capitalization, relatively high correlations were found with:

[Hh]urricane, *[Pp]ope*, *[Ll]ottery*, *[Zz]oo*, and low correlations with *[Tt]roy*, *[Pp]ath*, *[Ee]ditions*, *[Cc]ontinental*, and *[Bb]urns*.

Here, the terms seem to be proper names (the so-called “Shining Path Guerillas” were very active in South America in the 1980s when the documents were collected, and presumably *Editions* was a newspaper-specific term and *troy* referred to the measure of precious metals).

A difficult case for capitalization is a capitalized name that shares a spelling with a common word, such as *Burns*.

We can sum up Church’s correlation hypothesis as follows.

Definition 7.4.1 (Correlation Entails Similarity) *If two terms are highly correlated, they’re good candidates for conflation to a common stem.*

Church’s Burstiness Hypothesis

An even more intriguing point in Church’s paper concerns the burstiness of content words. The hypothesis is that a good indexing term like *hostage* is likely to show up more than once if it shows up at all. In fact, it’s very easy to reject the hypothesis that the distribution of words is a simple independent binomial or Poisson distribution because of burstiness. Since Church’s paper, various kinds of models have been proposed to model the counts of terms in papers using various kinds of two parameter models that estimate occurrence and number of occurrences separately, or ones that use overdispersed models instead of binomial or Poisson distributions (see, e.g., Jansche 2003).¹

Church considers the word *hostages*, which shows up six or more times in hundreds of documents in the five-year collection, which is millions of times more often than would be predicted by either a binomial or Poisson distribution. In contrast, function words like *the* or *be* fit a binomial or Poisson distribution much more closely, because the use of *the* in a document does not provide useful information for or against the appearance of *the* later in the document.

¹Recall that the binomial distribution is defined by a probability θ of a positive outcome per sample along with a number of samples n ; the chance of k positive outcomes in a sample size of k is given by:

$$\text{Bin}(k|\theta, n) = \binom{n}{k} \theta^k (1 - \theta)^{n-k}$$

For the Poisson distribution, the parameter is λ , the mean number of outcomes. The mean outcomes are scaled to a given interval, which we assume is one word for symmetry with the definition of the binomial. The chance of k outcomes given a text of length n is given by:

$$\text{Poisson}(k|\lambda, n) = \frac{(n \cdot \lambda)^k \cdot e^{-n \cdot \lambda}}{k!}$$

Church’s second search hypothesis can be boiled down to the following.

Definition 7.4.2 (Church’s Burstiness Hypothesis) *Terms that show high burstiness are likely to be good content terms.*

Church introduces a measure he calls *residual IDF* to denote the difference between the document frequency of a term and the expected document frequency of a term given its overall frequency under a Poisson distribution. Terms with high residual IDFs include *hostages*, *reactors*, *rebels*, *guerrillas*, *abortions*, and *delegates*. Terms with low residual IDFs include *awaits*, *grounds*, *possessions*, *boasts*, and *belongings*. Note that these are similar rankings to those found by comparing these *+s* variants to their non-suffixed forms, and thus the two measures (correlation and burstiness) are generating similar rankings.

Topic Mixture Models

Church notes that this kind of skewing would be expected if the documents were being generated by a topic-based mixture model. For a given topic, such as guerrilla warfare in South America, terms like *rebels* are much more common than they would be in a generic document. Function words and less topic-sensitive words, on the other hand, should be distributed much more evenly. This kind of topic mixture model is popular for language modeling in general (Chen 19??; Jurafsky 19??); language models have also been popular in direct application to information retrieval in recent years (Zhai 19??). Church hypothesizes at the end of his paper that both types of correlations (between a term and its variant, and between a term and itself repeated) are similar, and both are due to topic-like structure in documents.

7.4.5 Deligne and Bimbot, 1997

Deligne and Bimbot (1997) tackle the problem of extracting “regular” patterns from sequence data. They consider sequences consisting of either phonetic or graphemic units. Assuming a sequence of symbols (e.g. phonemes or characters), Deligne and Bimbot extract sequences of such symbols (e.g. syllables or words). We will restrict attention to the case where the training data is a sequence of characters without spaces.

Suppose $t = c_0 \cdots c_{n-1}$ is a sequence of n characters. We write $t = z_0 \cdots z_{n-1}$ if each z_i is itself a sequence of characters, and the concatenation of the z_i forms t ; we will call $z_0 \cdots z_{n-1}$ a segmentation of t ; Deligne and Bimbot refer to the z_i as multigrams. Deligne and Bimbot assume that a text

is generated by independently drawing segments from a multinomial distribution over multigrams. Note that they fail to account for the end-of-string probabilities properly, resulting in a defective distribution over segmentations; this problem is easily corrected by adding an end-of-string character to terminate sequences. Under this model, the likelihood of a string t given model parameters θ is the sum of the probabilities of its segmentations:

$$p(t|\theta) = \sum_{t=z_0, \dots, z_{k-1}} p(z_0 \cdots z_{k-1}|\theta) = \sum_{t=z_0 \cdots z_{k-1}} \prod_{i < k} p(z_i|\theta)$$

Deligne and Bimbot argue for a prior $p(\theta)$ over models θ based on a minimum description length (MDL) encoding of models. Under MDL, a prior represents the “cost” of encoding a model. As Deligne and Bimbot note, both the characters of the model and their probabilities must be encoded as part of $p(\theta)$. They then frame the problem as a search for the maximum a posteriori (MAP) model θ^* given a training sequence t :

$$\theta^* = \operatorname{argmax}_{\theta} p(t|\theta) \cdot p(\theta)$$

Unfortunately, Deligne and Bimbot never define the MDL prior $p(\theta)$, nor do they provide a search method for θ^* .

Instead of following the MDL approach, Deligne and Bimbot use a heuristically modified form of expectation/maximization (EM) to find an estimate $\hat{\theta}$. Without the prior term $p(\theta)$, the mode of $P(t|\theta)$ is generated from a model in which there is a single multigram t of probability 1.0.

Rather than finding a posterior mode θ^* , as EM typically does, Deligne and Bimbot impose three heuristics on the search:

1. Restrict the maximum length of multigram to k (5 in their examples);
2. Remove multigrams with low expected counts (3 in their examples);
3. Replace the standard fractional expectations computed in the E step with an assignment of probability 1.0 to the most likely outcome.

This last winner-take-all step leads to what is often called Viterbi-EM; in general, it has been shown to converge faster, but to be far less accurate in mode finding than standard EM. It is also much easier to implement and faster during each iteration. Because the prior distribution $p(\theta)$ is not defined, it is impossible to evaluate the effectiveness of Deligne and Bimbot’s heuristics. Unlike standard EM, which increases training data likelihood at each step, it is unclear whether the method will find a maximum subject to

0. buthisdelightisinthelawofthelordandinhislawdothhemeditatedayandnight
1. but his deli ghtis inthe lawof thelo rdand inhis law dothh eme
ditat eday and night
3. but his deli ghtis inthe lawof the lord and inhis law doth he
medit at e day and night
5. but his de light is inthe lawof the lord and inhis law doth he
medit at e day and night

Figure 7.2: Iterations of Deligne and Bimbot’s Approach (carried to convergence)

the constraints. It is also unclear under what conditions it will terminate, though pruning using the second heuristic above will also tend to drive the algorithm toward termination.

Deligne and Bimbot evaluate over the psalms drawn from the King James bible; based on 26 uncased letters, it contains 130,000 characters and 2500 distinct words. In figure 7.2, we repeat their dump of the iterations of their algorithm on a small sentence drawn from their corpus. Note that after 5 iterations after initialization, the algorithm has converged for this sentence; they state that in 10 iterations, it converges over the entire 130,000 character sequence. There is no evaluation of word induction beyond inspection of the output in figure 7.2; in the case of phonemic unit induction, they provide word accuracy figures resulting from using the model $p(t|\theta)$ as a language model in a speech decoder.

7.4.6 Jacquemin, 1997

Jacquemin (1997) provides a completely unsupervised technique for bootstrapping stems. The stated research goal is to find high-precision variants of terms. He uses a proprietary medical corpus called Medic for evaluation.

Jacquemin first induces an approximate notion of equality based on simple string matching. He then refines this crude equality measure based on higher-order cooccurrences in multi-word phrases.

Word Similarity

Word similarity is defined quite simply. First, common prefixes are removed and then similarity is measured by the smallest suffix left over. For instance,

for the pair `immunized` and `immunization`, the common prefix `immuniz-` is stripped off, leaving `-ed` and `-ation`. The shorter of these is of length 2, so the words are said to be 2-similar; they are also k -similar for any $k \geq 2$.

Similarity within Phrases

Jacquemin then searches for k -similar variants within a window of n terms. His examples here are `requires the use` and `user requirement`, where `requires` is 1-similar to `requirement` (suffixes `-s` and `-ment`), and `use` is 0-similar to `user` (suffixes `ε` and `-s`). The minimal window in which the terms are found is of length 3, so we say these are k - n -similar for $k \geq 1$ and $n \geq 3$. He then presents a table of precision and recall for various choices of k and n , with or without allowing order reversal (as in the example above). This basic idea is derived from (Martin et al. 1983), who consider values of n up to 5. The tables he presents are inconsistent, in that longer values of n sometimes lead to lower recall, which should be technically impossible, as longer values of n find everything shorter versions do. Further, his f-measures are not f-measures; for instance, for $n = 2, k = 3$, he reports 81% precision, 73% recall and 21% f-measure. The analysis in the paper (unlike the tables themselves), report the expected results: precision drops as the window size n increases, recall increases as the size of the suffixes k increases. He concludes the best f-measure (balanced for precision with $\beta = 3$) is obtained with 2-word windows without transposition, with the best value of k being 3. Therefore, the rest of the paper reports values for $k = 3, n = 2$ with no transposition.

Signatures

Jacquemin defines a signature to be the set of suffix pairs derived from phrase similarities. He provides the example of `continuous measure-ment` and `continuous-ly measure-d`, which have a signature of $\{\langle \epsilon, \text{ment} \rangle, \langle \text{ly}, \text{d} \rangle\}$.

In measuring similarity, Jacquemin hard-codes the fact that `-e` is a common alternation, considering `chromosom-e break` and `chomrsom-al break` to generate the same signature as `optic fiber` and `optic-al fiber`. He notes that this produces erroneous results for examples such as `rate` and `rat`.

Jacquemin extracts a total of 747 such signatures, then counts up the number of times each signature is realized; he calls the set of realizations of a signature a class. He notes that 92% of the errors occur in signatures for which every instance is incorrect. He then generates an ad-hoc formula

for the “quality of a class”. In short, quality goes up if the maximum suffix length stripped goes down, and goes up with the average length of similarity between shared prefixes. Ironically, given his goals, his final results favor recall over precision.

Finally, he presents an ad-hoc distance metric (with a two-page proof of the triangle inequality in an appendix) between classes. He then uses this distance to do greedy hierarchical clustering of the classes. He then presents some clusters for perusal, but does not indicate how to use such clusters for any kind of inference. As an example of a cluster, the pair **bend test** and **bend-ing test** with signature $\{\langle\epsilon, \epsilon\rangle, \langle\text{ing}, \epsilon\rangle\}$. is clustered with the pair **blot-ting assay** and **blot assay** with signature $\{\langle\text{ting}, \epsilon, \rangle, \langle\epsilon, \epsilon\rangle\}$.

7.4.7 Nevill-Manning and Witten, 1997

Nevill-Manning and Witten (1997) tackled the problem of decomposing a sequence (of text for the case with which we’re concerned) into hierarchical binary tree structures. Although not specifically concerned with morphology, it’s tempting to read morphology into their output. For instance, the word *beginning* gets the hierarchical structure $((((\text{be})\text{g})(\text{in}))((\text{ni})(\text{ng})))$ in which the top level split is between *begin* and *ning*, which is neatly morphologically segmented. Even with such a bracketing of a text, the question remains of where to draw morpheme boundaries.

7.4.8 Déjean, 1998

Déjean applies a version of Harris’s (1955) successor/predecessor frequency. It is not exactly clear what Déjean’s ultimate goals are for his model and the evaluation is vague. He says, for instance, “we do not try to discover all the morphemes contained in the corpus, since only the hundred most frequent ones are necessary in order to climb to the chunks level.”

Déjean outlines a three step process. The first step finds very frequent morphemes using a version of successor/predecessory frequency. The second step uses these morphemes to do further segmentations over the entire word list. Finally, the third step performs a final segmentation given the morphemes derived in the second step. There’s no guarantee with this kind of organization that the items extracted at any step are optimal for the subsequent stages of processing.

Déjean extracts the most frequent suffixes by looking for sudden growth in predecessor frequency. For instance, he notes that **-on** was preceded by 18 different letters, but 292/367 of the words that ended in **on** also ended in

-ion. This exceeds the 50% rejection threshold set by Déjean, so the suffix **-on** is rejected. Déjean retains suffixes with counts of 100 or greater in the analysis of the corpus. He reports the most frequent morphemes in English to be:

-e, -s, -ed, -ing, -al, -ation, -ly, -ic, -ent

Because of the heuristic noted above, this list looks slightly different than others of its kind for English. For French, the list was:

-s, -e, -es, -ent, -er, -és, -re, -ation, -ique, -,

In his second processing step, Déjean uses these frequent morphemes to extract stems. In extracting a stem such as **light**, he considers not only forms with known suffixes, such as **lights** and **lighted**, but also ones with unattested morphemes, such as **light-ness**, which introduces the potential morpheme **-ness**. If more than half the morphemes found are known, the new morphemes are kept. This provides a final list of 50 suffixes (and prefixes) for English. The suffixes are:

-able, -ably, -age, -al, -ally, -ances, -ance, -ation, -ations, -e, -ed, -ely, -ements, -ement, -en, -ences, ...

The final word segmentation is carried out by stripping off longest matching prefixes and suffixes. The most frequent 5% of words are not analyzed. Déjean reports a 1.5% error rate, but it appears that his classification of errors was quite strict.

In order to handle the most frequent 5% of the cases and also to figure out compound morphemes, Déjean uses higher-order contextual co-occurrence statistics. In particular, he uses bigrams, such as **des Hauses** in German. Although the exact technique is not described, Déjean suggests that because **des R-es** is a common segmentation for stems *R*, it should be possible to infer that **Haus-es** is the correct segmentation.

7.4.9 Manning, 1998

Manning (1998) takes as a starting point the numerous morphology induction studies in the linguistics and psycholinguistics literature centered around the connectionist (Rumelhart and McClelland 1986) versus symbolic (Ling 1994) debate. That literature has focused on inducing English past tense verb patterns.

	<i>foot</i>	<i>house</i>
<i>first person</i>	kepina	yotna
<i>second person</i>	kepika	yotda

Figure 7.3: Similarity and Difference Matching Example

	<i>thong</i>	<i>stake</i>
<i>nominative</i>	himas	skolops
<i>accusative</i>	himanta	skolopa
<i>genitive</i>	himantos	skolopos

Figure 7.4: Phonological Conditioning Example

Manning provides a system that induces traditional context-sensitive morpho-phonological transformations from training data consisting of phonological transcriptions of words along with their “semantics”, which is really just a unique identifier for the “predicate” and a collection of morpho-syntactic features (e.g. past tense or third-person singular). Note that this unique identifier uniquely picks out stems – two words have the same stem if and only if they have the same predicate identifier.

Manning applies a mixture of rule-based heuristics and statistical decision tree induction. The approach first applies the usual suffix and prefix stripping and matching, then attempts to reduce the resulting rule sets by incorporating point-of-affixation transformations such as consonant doubling and vowel insertion/deletion as well as sandhi (word internal transformations) such as vowel lengthening in Hungarian.

Manning introduces a notion of similarity matching and difference matching in order to derive candidate sets of stems and suffixes. He uses the example from Greek in figure 7.3. He notes that in terms of similarity, first person may be analyzed as **-na** and second person as **-a**, because those are the maximal shared suffixes reading across the features. Reading down, the differences between first and second person are **-na/-ka** for *foot* and **-na/-da** for *house*. He uses these two dimensions to sort out where the boundary effects occur.

Allomorphs are detected by a heuristic similarity measure over a traditional set of phonological features. For instance, he considers an example from Greek in figure 7.4. As an example of an initial analysis, the various forms of *thong* may be realized with stem **hima-** and suffixes **-s**, **-nta**, and

-ntos. Then the forms of *stake* would be realized with stem **skolop** and suffixes **-s**, **-a**, and **-os**. The main insight in Manning’s paper is moving from this analysis to one that is more “economical”. As usual in linguistics, the notion of “economy” is not explicitly defined and we are left to infer that Manning’s final result is more economical than the analysis above. The final result uses one set of suffixes **-s**, **-a** and **-os**, a single stem **skolop** for *stake*, with *thong* being realized with the pair of context-sensitive rules:

$$\begin{array}{lcl} \text{hima} & \text{thong} & / \text{ ___ s} \\ \text{himant} & \text{thong} & / \text{ elsewhere} \end{array}$$

with the usual caveat that the most specific rule applies first (though Manning does not define an ordering). He calls the result *conditioned allomorphy*. The system for measuring this includes very specific rules such as one allowing vowel-lengthening to be a better match than other transformations.

7.4.10 Xu and Croft, 1998

Xu and Croft (1998) introduce an approach to stemming that tunes an aggressive stemmer like the Porter stemmer by using corpus co-occurrence statistics. They adopt the same motivation as Church, namely that words variants that should be conflated by stemming are likely to co-occur in documents.

To measure co-occurrence, Xu and Croft introduce a heuristic weight:

$$\text{sim}(X, Y) = \max(0, \frac{\text{freq}(X, Y) - \text{E}(\text{freq}(X, Y))}{\text{freq}(X) + \text{freq}(Y)})$$

The terms $\text{freq}(X)$ and $\text{freq}(X, Y)$ measure individual and joint frequencies. These frequencies are collected within a fixed width window of size κ , introducing a tunable hyperparameter into the model. The expectation is computed using an independence assumption:

$$\text{E}(\text{freq}(X, Y)) = \text{corpusSize} \cdot p^{\text{ml}}(X) \cdot p^{\text{ml}}(Y)$$

where the maximum likelihood estimates are also based on frequencies:

$$p^{\text{ml}}(X) = \frac{\text{freq}(X)}{\sum_Z \text{freq}(Z)}$$

This makes the numerator of their term the same as in a significance test of independence of the terms X and Y , which is:

$$z = \frac{\text{freq}(X, Y) - \text{E}(\text{freq}(X, Y))}{\text{corpusSize} \cdot p^{\text{ml}}(X, Y) \cdot (1 - p^{\text{ml}}(X, Y))}$$

This variance-adjusted value has the usual statistical motivation for testing the independence of X and Y .

Xu and Croft note that **bond/bonds** have a 0.35 similarity, **animation** and **animators** a 0.14 similarity and **arm/army** and **desirable/desires** have near zero scores.

Xu and Croft use this pairwise comparison between terms in a connected component algorithm to pick out subclasses of terms conflated by an aggressive stemmer. They first define a graph whose vertices are the terms and where there is an edge between terms X and Y if $\text{sim}(X, Y) > \tau$ for threshold τ , introducing a hyperparameter. They then conflate any pair of terms that are connected by a path. This is just a transitive closure of a boolean matrix, which may be computed efficiently by means of a maximal connected components algorithm.

They also suggest an optimal partition algorithm that tries every partitioning to see which one has the least total cost. The cost of two words appearing together in a partition is defined to be $\text{sim}(X, Y) - \delta$, introducing another hyperparameter δ . They were only able to optimally partition up to size 12 clusters due to combinatoric issues. In practice, they speed this up by first using the connected component algorithm with a generous threshold to pick out smaller subclasses of the the conflation classes produced by the aggressive stemmer.

Xu and Croft train over four corpora containing tens of millions of words each. These were the West legal document corpus, two *Wall Street Journal* TREC corpora, and a Spanish corpus, ISM from TREC 4. Query lengths were roughly 5 and 10 for ISM and West, whereas they were about 35 for the TREC *WSJ* corpora. Documents averaged around 3000 tokens for WEST and 275 for the others.

They evaluate the Porter stemmer for the three English corpora. For both the English corpora and Spanish corpus, they evaluate a conflation strategy that conflates words if they have the same initial three characters. Precision gains were modest, but significant. On the West corpus, the basic 3-gram prefix-match stemmer achieved a 50.0 average precision at 100 recall, the Porter stemmer was 50.1 and the optimal partition algorithm variant 50.7. On WSJ, 3-stem was 27.6. For the Spanish corpus, average precision at recall 100 was 21.2 for the initial 3-gram stemmer and 21.9 for the initial 3-gram stemmer with optimal partitions.

7.4.11 Gaussier, 1999

Gaussier (1999) introduced a two-stage approach to developing a stemmer that is trained using dictionary part-of-speech information. The approach first derived relational families (or paradigms) and then used these to train a stemmer.

Suffix Pairs

The first stage in Gaussier’s construction extracts a set of suffix pairs heuristically. Suffixes τ_1 and τ_2 are said to be a suffix pair if there exist at least two character strings σ_1 and σ_2 of at least five characters in length such that all combinations $\sigma_i \cdot \tau_j$ are words. For instance, **-y** and **-ily** constitute a suffix pair, because the above condition is met with $\sigma_1 = \text{ordinar}$, $\sigma_2 = \text{summar}$, $\tau_1 = \text{y}$, $\tau_2 = \text{ily}$, because **summary** (5596), **summarily** (957), **ordinary** (28,886) and **ordinarily** (2200) are words (with counts derived from the *New York Times* section of LDC’s Gigaword corpus). Note that these examples also allow us to conclude that **-ry** and **-rily** is also a suffix pair.

We are simplifying Gaussier’s presentation, which also associated parts-of-speech with each word from a dictionary and then used that information along with suffixes to determine equality. That is, the French suffix pair **-ation** and **-er** cited by Gaussier, has the full form **-ation:N** and **-er:V**, indicating that the word produced by adding the first suffix is a noun, whereas the one produced by the second suffix is a verb. This will likely improve analyses of cases where a word coincidentally ends in a sequence that matches a suffix but does not occur in the correct part of speech for the suffix.

Relational Families

For Gaussier, a relational family is a maximal set of words that share the same stem. Gaussier infers these families by agglomerative clustering over a pairwise similarity measure between words. Two words ω_1 and ω_2 only have non-zero similarity if there is a suffix pair τ_1, τ_2 and some string σ such that $\omega_1 = \sigma\tau_1$ and $\omega_2 = \sigma\tau_2$. Their relatedness is then given by the productivity of the suffix pair τ_1, τ_2 :

$$\text{sim}(\sigma \cdot \tau, \sigma \cdot \tau') = \text{card}(\{ \rho \mid \rho \cdot \tau \text{ and } \rho \cdot \tau' \text{ are words} \})$$

where $\text{card}(s)$ is the cardinality of the set s . Note that this definition, as presented by Gaussier, is not well-defined, because he assumes there will be

a unique suffix pair linking a pair of words; as we saw in the last section, this is not necessarily the case. In cases such as the above, the longer pair of suffixes can have no more matches than the shorter pair.

With a similarity measure defined, Gaussier explored single-, average- and complete-link agglomerative clustering, concluding that complete link provided the best clusters. As examples of derived clusters, he cites:

deprecate, deprecation, deprecator, deprecative, deprecativeness,
deprecatively, deprecativity, deprecatorily, deprecatory,
deprecatingly

department, departmentality, departmental, departmentalness,
departmentally

depart, departure, departer

Note that the words that begin with **depart-** are broken into two classes, because the suffixes **-ure** and **-er** do not alternate frequently enough with **-ment**, **-mental**, **-mentality**, and **-mentally**.

Unfortunately, Gaussier does not tell us how to cut off the clustering procedure to provide a final partitioning of the set of words. He does say he explored multiple “versions”. Presumably his accuracy report is for the most favorable threshold, though it’s possible that the results arise from an even more generous selection of the best consistent set of clusters from the dendrogram resulting from clustering.

Gaussier reports an 85% accuracy figure for words being placed into the correct class using complete-link clustering, with group-average clustering being 77% and single link 47%. He considers a word to be in the correct class if two conditions hold. First, the majority of the words that have the same stem as the word in question must be in the class. Second, the majority of the class must be made up of the words with the same stem. Gaussier reports 65% accuracy for the Porter stemmer and 82% accuracy for the SMART stemmer using the same reference evaluation.

Deriving a Stemmer

Gaussier assumes that a stemmer is made up of a number of suffixation operations. Each such operation contains two components, a morpho-syntactic operation and a suffix. The morpho-syntactic operation maps one syntactic category into another. In English the following suffixation operation could

be used to transform the base verb **eat** into the singular noun **eater**:

$$S = \langle V_b \rightarrow N_s, -er \rangle$$

Although a suffixation operation contains only a suffix, Gaussier allows the application of that suffix to differ from the result of concatenation by up to 3 characters. For instance, the English plural suffix **-s** might attach to a singular noun as **box** to produce **boxes**, in which an extra **e** is inserted; similarly, the suffix **-ing** might apply to the verb **pace** to produce the form **pacing** in which the final **e** of the original form is removed.

Gaussier estimates the probability of a realized suffixation by its distance from pure concatenation. For instance, **boxes** differs from the concatenation **box · s** by one character. Gaussier estimates the probability $p(n)$ of a difference of n characters by taking $\hat{p}(0) = 0.5$, $\hat{p}(1) = 0.25$ and $\hat{p}(2) = 0.125$. This produces a defective probability distribution because the differences themselves are not being generated. For instance, we might insert 26 different letters in English, each with a probability of 0.25, leading to a total edit probability estimate of 6.5, which exceeds the limit of 1.0 imposed by a proper distribution. This model is clearly defective.

This distribution could be made proper by treating it as a stochastically weighted edit distance, but then inference would need to be done by summing over all edits that lead to the same result. Languages do not use a wide range of edits at suffixation boundaries, so most of these estimates should probably be very close to zero. non-zero probability (e.g. for English, duplicating consonants (e.g. inserting **m** in **swimming**, inserting **n** in **running**, deleting **e** in **pacing**, etc.)

A uniform estimates of the inverse of the number of characters would provide a proper estimate for insertion (or deletion). The exponential decay then arises naturally as the product of edits. Unfortunately, this estimate would seriously underestimate edits. For instance, the insertion probability of an **e** would likely be around 1/100 in a uniform estimate for English, because the model must allow for matches, insertions, deletions and substitutions of arbitrary characters up to 3 in length. The underestimation arises from ignoring relevant contextual information about the form of the stem and suffix. Thus we believe the estimates must essentially be conditional. For instance, we could predict the edit based on the final rhyme of the stem (the final vowel complex until the end) and the suffix:

$$p(\text{insert}(\mathbf{m})|\mathbf{im}, \mathbf{ing})$$

Presumably this estimate would be much higher than 1/100.

Gaussier also assumes a probability distribution $p(S)$ over suffixes S , which he estimates using the expectation/maximization (EM) algorithm. Within the maximization step of the EM algorithm, Gaussier fudges Bayes's law with a multiplicative constant, presumably to make up for the improper distribution introduced for boundary edits. He provides no information on convergence.

Paradigm Extraction

The last step in a single pass of Gaussier's modeling involves extracting paradigms from the resulting stem probability estimates. The paradigm extraction step operates on a single relational family at a time. For each relational family, word pairs in the family are considered in decreasing order of likelihood (as estimated in previous steps).

If a pair relates words already related through previous pairs, the pair is skipped. For pairs being considered, the set of stems that could lead to the pair through suffixation is constructed.

After the last pair is added, all the sets of stems found are intersected to produce the final resulting set of possible stems.

Gaussier provides an example involving French words **produire**, **production**, and **producteur**. He assumes the probability estimates are such that the pairs considered are $\langle \text{produire}, \text{production} \rangle$ and then $\langle \text{produire}, \text{producteur} \rangle$. For the first pair, the possible suffixes allowing **production** to be derived within the maximum allowable number of edits are **ction**, **tion** and **ion**, with corresponding root "allomorphs" **produ**, **produc**, **product** and **producti**. For the second pair, the possible suffixes are **cteur**, **teur**, **eur**, and **ur**, leading to allomorphs **produ**, **produc** and **product**. Intersecting the two sets eliminates the allomorph derived from only the first pair, namely **producti**.

Iteration

After extracting paradigms, Gaussier repeats the whole process with the suffix pairs in the next step restricted to those used in the paradigm extraction phase of the current iteration.

For instance, we initially assumed that **on** was a possible suffix. If, in fact, all other families show the same behavior as **produire**, then all of the instances of **on** may be eliminated from consideration in the next iteration.

It would be interesting to see if a model could be built in which edits are conditional probability estimates and paradigms are estimated probabilisti-


```
|the| fulton county| grand jury| said |friday | |an| investation of|
atlanta| 's |recent| | primary |election| produced| ‘‘ no| evidence|
| ’’ |that any| irregularities| took place |.|
```

Figure 7.5: Kit and Wilks’s segmentation of the Brown corpus

cally rather than heuristically. If so, the entire process could be maximized into a point estimate using EM.

7.4.12 Kit and Wilks, 1999

Kit and Wilks (1999) approach the same problem as Deligne and Bimbot (1997), that is segmenting a corpus of text without word boundary information into words. Kit and Wilks introduce an approach based on a greedy heuristic rather than a statistically sound search strategy.

Their approach is based on starting with an unsegmented text and then measuring the reduction in description length (which they confusingly call “description length gain”) arising from adding a new multigram to the model. They assume that the new multigram will be used to analyze every instance which is still available (that is, not part of a previous multigram).

Unlike Huffman coding, Kit and Wilks’s greedy search strategy is unlikely to uncover minimum description length segmentations because of the extensive dependencies in natural language data. Although they call it “Viterbi segmentation”, they do not demonstrate the critical decomposition property underlying the dynamic programming condition of the Viterbi algorithm.

Kit and Wilks report the segmentation of the corpus listed in figure 7.5. The location of spaces to the left or right of segmentation points is not significant in their evaluation; their inclusion of spaces in locations such as before the possessive ’s leads us to believe they used the segmented form of the Brown corpus as their basis. Rather than evaluating words, they instead evaluate the more generous metric of word boundary detection, reporting models trained on 500,000 characters to have precision=67.8%, recall=70.4% and those trained on 6,000,000 characters having precision=79.3% and recall=63.0%.

They claim prediction based on more data is more reliable, but they do not notice that as data grows, MDL becomes problematic, because the cost of encoding something in a dictionary becomes insignificant, so that longer and longer segments are entertained. In effect, this should reduce their recall, just as they observe.

7.4.13 Yarowsky and Wicentowski, 2000

Yarowsky and Wicentowski (2000) factor the suffixation process into a stem-final change and a concatenative suffixation. For instance, the word **defy** is mapped to the word **defied** by modifying the final **y** in the stem to an **i** and then concatenating the usual suffix **ed**. For highly irregular pairs like **sing** and **sung**, the final stem change is extensive, mapping **ing** to **ung** with no suffix.

The major benefit of Yarowsky and Wicentowski’s model is that the suffix change operation absorbs the point-of-suffixation complexities such as epenthetic vowel insertion, leaving the stems and suffixes as simple and high frequency as possible.

Yarowsky and Wicentowski’s algorithm is “semi-supervised.” It requires a list of parts of speech (e.g. past participle verb) and a list of canonical suffixes for the part of speech (e.g. **-en**, **-ed**, **-t**, and **-e** for past participle verbs in English). They also require a part-of-speech tagger for the language in question. In addition, they require a classification of the letters into vowels and consonants. Their algorithm allows, but does not require, a list of function words. It also allows data from similar languages (e.g. the output of the algorithm on Italian when training Spanish). They then combine this supervised data with a large corpus of unannotated data.

Yarowsky and Wicentowski introduce a relative frequency measure that compares a potential relation based on frequency. They note that $\text{freq}(\text{sang})$ divided by $\text{freq}(\text{sing})$ is 1.19, close to the average ratio of 0.85 of past tense verbs to base verbs. They contrast this with $\text{freq}(\text{singed})/\text{freq}(\text{sing})$, which is only 0.007, and $\text{freq}(\text{singed})/\text{freq}(\text{singe})$, which is 4.5. They approximate these numbers take supervision from the part-of-speech tagger and the list of regular affixes. They consider a range of such estimates, but somewhat surprisingly, never consider adjusting any of them for variance; we know that for small sample counts, the variance in the ratios will be much larger, so confidence intervals need to be much wider.

Yarowsky and Wicentowski also use context similarity. As usual, they create vectors of features based on local context. Rather than just using words, they use “context features”, which are derived from applying regular expressions of part-of-speech tags. They describe feature selection from among this set of patterns.

Yarowsky and Wicentowski use a form of edit distance to align potential stems with words and suffixes. They use this to generate baseline probabilities using mixture models in order to estimate the probability that the stem **solidify** and the suffix **-ed** producing a past participle verb will combine

to yield **solidified**. The result is fully determined by the stem change, so reduces to the probability of a stem change from **y** to **i** before **-ed**, which they estimate without recourse to the target part-of-speech.

Yarowsky and Wicentowski only evaluate their system on the past tense of English verbs, reporting 99.2% accuracy.

7.4.14 Yu, 2000

Yu (2000) takes on the same problem as Deligne and Bimbot (1997). Like Deligne and Bimbot, Yu applies a heuristic variant of expectation / maximization (EM) that is motivated by the minimum description length (MDL) principle.

Unlike Deligne and Bimbot, Yu provides a formula for his prior. As in (Deligne and Bimbot 1997), the probability model is a multinomial θ over multigrams z_0, \dots, z_{n-1} . Yu defines a prior based on a rough estimate \hat{h}_c of the encoding cost of the characters making up the multigrams with non-zero probability:

$$\log_2 p(\theta) \propto \sum_{i < n} \hat{h}_c \cdot \text{len}(z_i)$$

Yu defines \hat{h}_c as the estimate of single-character entropy:

$$\hat{h}_c = H(\hat{p}^{\text{ml}}) = \sum_c \hat{p}^{\text{ml}}(c) \cdot -\log_2 \hat{p}^{\text{ml}}(c)$$

derived from the maximum likelihood character multinomial estimate over training data t :

$$\hat{p}^{\text{ml}}(c) = \frac{\text{freq}(c, t)}{\sum_{c'} \text{freq}(c', t)}$$

Yu's prior ignores the cost of encoding the end-of-multigram character. More significantly, he also ignores the cost of encoding the model probabilities. In Bayesian terms, Yu's prior is improper, and is said to be uninformative about the distribution in the sense that any distribution is as likely as any other.

Yu introduces a novel heuristic that he applies as a weight during the estimation step of his EM-like algorithm. In MDL, the total description weight is minimized:

$$\text{descLen}(t, \theta) = \text{descLen}(t|\theta) + \text{descLen}(\theta)$$

Because of the duality between description length and probability:

$$\text{descLen}(t|\theta) = -\log_2 p(t|\theta)$$

and

$$\text{descLen}(\theta) = -\log_2 p(\theta)$$

So the minimum description length turns out to be the model θ that has the maximum a posteriori (MAP) likelihood in the Bayesian setting given the data t :

$$\begin{aligned} \theta^* &= \operatorname{argmax}_{\theta} p(\theta|t) \\ &= \operatorname{argmax}_{\theta} \frac{p(t, \theta)}{p(t)} \\ &= \operatorname{argmax}_{\theta} p(t, \theta) \\ &= \operatorname{argmax}_{\theta} p(t|\theta) \cdot p(\theta) \\ &= \operatorname{argmax}_{\theta} \log_2 p(t|\theta) + \log_2 p(\theta) \\ &= \operatorname{argmin}_{\theta} -\log_2 p(t|\theta) + -\log_2 p(\theta) \\ &= \operatorname{argmin}_{\theta} \text{descLen}(t, \theta) \end{aligned}$$

As usual, the model description length is $-\log_2 p(\theta)$, the negative log likelihood of the model prior, and the data description length given the model is $-\log_2 p(Z|\theta)$, the negative log likelihood of the data given the model. Yu then transforms the description length equations as:

$$\begin{aligned} \text{descLen}(t, \theta) &= \text{descLen}(\theta) + \text{descLen}(t|\theta) \\ &= \sum_i \hat{h}_c \cdot \text{len}(z_i) + \sum_i -\text{freq}(z_i, t) \cdot \log_2 p(z_i|\theta) \\ &= \sum_w \hat{h}_c \cdot \text{len}(w) - \text{freq}(w) \cdot \log_2 p(w|\theta) \\ &= \sum_w \text{freq}(w) \cdot \left(\frac{\hat{h}_c \cdot \text{len}(w)}{\text{freq}(w)} - \log_2 p(w|\theta) \right) \end{aligned}$$

Technically, the first reduction of the description length of t is wrong, unless the text has a unique decomposition into multigrams. Instead, this term should be as in Deligne and Bimbot's calculations:

$$\log_2 \sum_{t=z_0 \cdots z_{n-1}} \prod_{i < n} p(z_i|\theta)$$

Presumably Yu intends to use a Viterbi-style estimate of the best segmentation:

$$\log_2 \max_{t=z_0 \cdots z_{n-1}} \prod_{i < n} p(z_i|\theta)$$

Using his defective estimate, Yu then replaces the standard estimation step in the EM algorithm with one modified to include a “penalty term” for the cost of the codebook:

$$\hat{p}'(z|\theta) = \hat{p}(z|\theta) 2^{\frac{\hat{h}_C \cdot \text{len}(w)}{\text{freq}(z,t)}}$$

This has an enormous effect on the computation of the expectation step. Furthermore, like Deligne and Bimbot, Yu uses the first-best Viterbi approximation of EM. The result is no longer EM, and thus the standard convergence arguments no longer work, despite Yu’s claims to the contrary. Yu, like Deligne and Bimbot, place an upper bound on segment size. As shown in the final table of the paper, the algorithm does not find a maximum, as might be expected; in fact, the standard segmentation of words outperforms it.

Yu also modified Deligne and Bimbot’s estimates by allowing smoothing to unseen words using interpolation with a model based on character entropies. As in its earlier use, this distribution is improper because it does not code end-of-word information nor take into account that not all characters have the average entropy. He does not say how he estimates the interpolation ratio with the held out estimates, nor does he analyze the effect of this interpolation.

Yu notes these as search errors in section 3.2, but attributes it to pruning, which is another heuristic that Yu follows Deligne and Bimbot in adopting. It’s not clear what exact search strategy is being used, as Yu cites an unpublished paper by A. Raman as the source of additional “conservative” search heuristics.

7.4.15 Goldsmith, 2001

Goldsmith (2001) provides a method, which he names *Linguistica*, to find single split points between a stem and its zero or more suffixes. Like the Porter stemmer, simple segmentation results in pseudo-stems. Goldsmith provides the example *notice*, *noticed*, *notices* and *noticing*, all of which produce stem *notic*, with suffixes *-e*, *-ed*, *-es* and *-ing* respectively. The main danger with pseudo-stems is that they may overlap with real words, as in *care*, which patterns like *notice*, producing stem *car*, and the independent word *car*, which bears no morphological relation to *care*.

Goldsmith motivates his approach based on sets of stems that co-vary. Using the book *Tom Sawyer* as a corpus, Goldsmith extracts the signature ϵ , *-ed*, *-ing* and *-s*, all of which occur with the word **remain**. The word

`notice` on the other hand has stem `notic` and signature `-e`, `-ed`, `-es` and `-ing`.

Goldsmith introduces a notion of “prefix conditional entropy”, which is just the entropy of the prefix, computed with an n -gram model. The conditionality is that of the standard n -gram model which predicts the next character given a sequence of the previous characters; the sum of the entropy computed at each character is just the entropy of the prefix itself.

Goldsmith proceeds from a minimum description length (MDL) motivation. He notes that the maximum likelihood model would simply return each word unanalyzed according to its empirical frequency. So he imposes costs on the coding of words, hoping to force analyses into smaller parts.

Goldsmith’s initial model cost codes the number of stems using Ris-sanen’s (1989) universal prior. This only slightly favors models with fewer stems overall. Goldsmith includes both probability costs and character string costs. He notes that an outcome of likelihood p in a model requires $\log_2 p$ bits to encode. He also encodes the actual letters in each stem using a uniform distribution over all characters, for a cost of $\log_2 26$ bits per character for English.

Goldsmith’s models also include signatures. There is a distribution over signatures that is encoded, and then for each signature, the number of stems and suffixes is encoded, the stem and suffix probabilities are encoded, and the characters themselves are encoded.

After establishing this model, Goldsmith then takes a heuristic approach to minimizing the total cost of encoding a model plus the data. He notes that EM without model costs is driven into the trivial solution where there is no morphology.

Goldsmith’s first heuristic is to replace the estimates in EM with ones that favor longer suffixes and stems. When these estimates are fed into the maximization step, the results are not guaranteed to converge. He defines log probability proportional to:

$$\log p(R, S) \propto \text{len}(R) \log p(R) + \text{len}(S) \log p(S)$$

This formula doesn’t actually seem to encourage splitting into longer stems and suffixes, because the multiples introduced by length have to go on one side or the other and making all suffixes ϵ still seems to maximize this estimate.

Goldsmith’s second heuristic restricts stems to the ones with the highest

scores per character, which he calls “weighted mutual information”:

$$\frac{\text{freq}(c_0 \cdots c_{k-1})}{\sum_{c'_0 \cdots c'_{k-1}} \text{freq}(c'_0 \cdots c'_{k-1})} \log \frac{\text{freq}(c_0 \cdots c_{k-1})}{\text{freq}(c_0) \cdots \text{freq}(c_{k-1})}$$

He restricts attention to the top 100 results of this ranking.

Goldsmith defines a signature to be a set **R** of stems and set **S** of stems such that $R \cdot S$ is a word for every $R \in \mathbf{R}$ and every $S \in \mathbf{S}$. Goldsmith assumes that each stem will be in exactly one such signature. Goldsmith enumerates numerous heuristics and means for evaluating whether a signature is a good one or not; he also supplies a nice data analysis of the mistakes made for signature extraction.

Goldsmith first splits suffixes if both of their parts show up and if the resulting description length is shorter (though that would seem unlikely unless the model was adjusted and the replacement made global). The second heuristic looks to see if all the suffixes begin with the same letter, in which case it is heuristically moved into the stem. A third stage heuristic is punningly called “triage”; it attempts to remove low count signatures. It looks not only at local counts, but at counts of related signatures.

Goldsmith provides a range of evaluations, using his own judgements as to accuracy. He first notes that the top signatures extracted, according to some heuristics, look quite good, saying: “it is almost as if the textbook patterns have been ripped out and placed in a chart.” (Goldsmith 2001, p. 178). He also presents numerical results, with 82.9% of analyses being considered “good”, which for Goldsmith, is a matter of both syntactic and semantic features. There were no phonotactic criteria, so he notes either **aboli-tion/abol-ish** or **abol-ition/abol-ish** would have been judged as good. In cases where there was a stem and multiple suffixes, but it was analyzed as containing two stems and two sets of suffixes was penalized. Thus if the system grouped **accompni-ed**, **accomini-ment** and **accompni-st**, and also grouped **accompany** and **accompany-ing**, he treats the first group as wrong (it is not clear why other than that **accompany** is the base form). In any case, these judgements would be unstable under the addition of new data, where new signatures may show up that contain existing ones.

In addition to the good cases, Goldsmith breaks the 17.1% of the cases involving errors into three groups. Goldsmith judged 5.2% of the analyses to be wrong in the sense that there was a correct stem and suffix and the analysis returned the wrong stem and suffix. 3.6% of the words consisting of two morphemes were wrongly not decomposed. A final 8.3% provided a morphological decomposition of a word that should not have been analyzed.

Goldsmith indicates that he is working on extending groups of signatures into paradigm; for example, each non-empty subset of the four suffixes ϵ , **-ed**, **-ing** and **-s** leads to a different signature, resulting in 15 total signatures for a single paradigm.

Goldsmith counterintuitively suggests that syntax and grammar would not help in resolving morphological issues. He then goes on to suggest that his model might not be such a bad one for child language acquisition.

7.4.16 Peng and Schuurmans, 2001

Peng and Schuurmans (2001) propose a heuristic hierarchical EM-derived model to extract words from unsegmented training data. In their introduction, they put their finger on the primary obstacle to using simple multinomial EM to extract words from unlabeled data. The problem is that breaking down sequences into smaller chunks tends to underestimate their probabilities. This leads EM into solutions where **ofthe** and **computerscience** are good word candidates. Others have approached the problem from a minimum description length perspective, penalizing the model for each word in the resulting multinomial. The problem with this approach is that as more data is seen, the likelihood overtakes the prior, again producing words like **ofthe**.

Peng and Schuurmans mount a two-pronged heuristic attack on the problem. First, they prune their resulting models, removing words like **computerscience** if the probabilities of **computer** and **science** are high enough. Their measure involves a pointwise (character length adjusted) mutual information measure:

$$pMI = \frac{1}{T} \cdot \log_2 \frac{p(S_1 \cdot S_2)}{p(S_1) \cdot p(S_2)}$$

If the measure is low enough, the words are split, if it is high enough, the word is left intact, and if it's in between, the a third of the probability is "shifted" to the split and two thirds is left for the whole word. From this operation, it seems that they are using a Viterbi-style EM, which as we've noted before, is not really EM in that it doesn't compute a proper expectation to maximize.

Although they claim pruning helps drive EM out of poor local minima, this should actually drive EM out of an optimal solution. That is, Peng and Schuurman's strategy introduces a significant bias into the estimate.

More interestingly, Peng and Schuurmans model segmentation hierarchically. Rather than generating a segment as a sequence of characters, they

generate a segment as a sequence of smaller segments. They apply a two-step inference process, first using their pruned EM to produce an inventory of short segments (up to 5 characters in their experiments). Then, they use these units in a second step to extract words. There is no guarantee that the units uncovered in the first stage will be optimal for the task required of them in the second stage.

Their evaluation scored boundary detection over the Brown corpus. With only the first level model (maximum word length 5), they report boundary detection performance of 0.57 precision and 0.91 recall for a 0.69 F measure. They compared this to a flat one-level model allowing longer segments with pruning, which scored 0.85 precision and 0.51 recall for segmentation for a 0.62 F measure. Clearly the flat model with a tighter upper bound on segment length worked better here. With their hierarchical model, trained over the units of the first model, the performance was around 0.75 F measure, with precision and recall varying between 0.68 and 0.81 depending on parameterizations.

Word level performance on their best model, the two-level model, was 0.49 precision and 0.60 recall, for a 0.54 F-measure.

7.4.17 Schone and Jurafsky, 2001

Schone and Jurafsky (2001) present a context- and frequency-based approach to clustering based on stems. Their approach involves multiple heuristic stages, many of which are based on reasoning over pairs of suffixes. We begin by restricting attention to their approach to suffixes.

As a first step, they remove what they call “pseudo-prefixes”, which are prefixes whose count exceeds a threshold t_1 (10 here).

In the second step, they apply a modified predecessor-frequency heuristic to extract potential suffixes that appear more frequently than a given threshold t_2 (10 here). They only count a suffix S in a word RS if there is another suffix S' such that RS' is also a word. They refer to R as a pseudo-stem. A “candidate suffix” is a suffix that is shared by at least t_3 (10 here) pseudo-stems that also occur with other candidate suffixes. This makes the candidate suffix calculation self-referential, but not inconsistent; presumably they take the maximum fixed point. Finally, they extract “rules”, which are pairs of suffixes S_1 and S_2 that appear with at least t_4 (here 3) pseudo-stems. In English, the highest frequency rules involving only suffixes are $-s/\epsilon$, $-ed/-ing$, $-ing/\epsilon$, $-ly/\epsilon$, $-ers/-ing$, and $-d/-r$.

Schone and Jurafsky next apply a stage of semantic filtering. They create two vectors for each word: a vector of word frequencies in a 50-word

window before the word in question and a similar vector for the 50-word window following the word. These were combined by concatenation to form a longer vector; they were not added together. They then normalize each dimension of the vector to a z-score over the vector. That is, if the count n is replaced by $(n - \mu)/\sigma$ where μ is the average word count in the vector and σ is the deviation; this value is just the number of deviations above or below the mean count n falls. They claim this is because SVD is designed to work with normally distributed data, which is curious, because SVDs are fairly simple least-squares dimensionality minimizers; they do not compare to the unnormalized alternative. They use dimensionality of 300, which probably results in a very tight fit to the original (normalized) data.

To compare two words W_1 and W_2 to each other, Schone and Jurafsky take their vector representations under SVD at 300 dimensions, and compute a cosine between the two words' representations. They then take a sample of 200 other words. They then look at the 200 cosines between W_1 and the words in the sample and compute a mean μ_1 and deviation σ_1 with which to compute a z-score of $z_1 = (\cos(W_1, W_2) - \mu_1)/\sigma_1$. They then do the same thing for the set of cosines in the set of 200 words with W_2 to produce the z-score z_2 and take the minimum of the two z-scores $\min(z_1, z_2)$ to be their "normalized cosine score" (NCS). They then use this along with normality assumptions to derive the probability of a given NCS. It's then possible to set a threshold below which a pairing is rejected.

Schone and Jurafsky note that for their high frequency rules such as $-\mathbf{s}/\epsilon$, 99.7% of the stems would be correct according to CELEX (Baayen 1995). Their threshold on semantic relatedness, on the other hand, only retains 92% of the pairs. In order to raise the recall, they blend in a notion of orthographic similarity based on edit distance, which raises recall to 95% of the pairs based on orthographic and semantic similarity.

Schone and Jurafsky next blend in a measure based on local syntactic context. For each stem pair, such as \mathbf{s}/ϵ , they look at roots that have both extensions and then build models consisting of a distribution of words to the right and distribution of words to the left. They note that words such as **are** show up significantly more to the right of **s** variants of words and **is** shows up more for the ϵ variant. For example, this pattern holds for roots **book** and **lab**, where you find **books are** versus **book is**, and **labs are** versus **lab is**. It is used to define a probability estimate. They use these contexts as further factors in the model along with semantic and orthographic features.

When all of these measures are added together (they use the "noisy or" formulation, which is problematic given the correlation between their three terms), they are left with a graph over the set of words where edges are

labeled with probability estimates that the two words have the same stem. For example, they estimate $p^{\text{ml}}(\text{stem}(\textit{abusing}) = \text{stem}(\textit{abusive}) = 0.98$ and $p^{\text{ml}}(\text{stem}(\textit{abusers}) = \text{stem}(\textit{abuses}) = 0.26$.

In such a system, we would expect the following transitivity constraint to hold over estimates:

$$p^{\text{ml}}(\text{stem}(x) = \text{stem}(y)) \geq p^{\text{ml}}(\text{stem}(x) = \text{stem}(z)) \cdot p^{\text{ml}}(\text{stem}(z) = \text{stem}(y))$$

which reduces to the triangle inequality for log probabilities:

$$\log p^{\text{ml}}(\text{stem}(x) = \text{stem}(y)) \geq \log p^{\text{ml}}(\text{stem}(x) = \text{stem}(z)) + \log p^{\text{ml}}(\text{stem}(z) = \text{stem}(y)).$$

There are various ways to coerce defective pairwise estimates into proper collective estimates. Schone and Jurafsky use a “decayed” form of transitive closure that does not guarantee the above equations are satisfied. Swanson et al. (19??) provide a novel solution to the transitive closure problem that doesn’t simply increase probabilities that are too low, but may also reduce probabilities that are too high.

With the final graph, Schone and Jurafsky then set thresholds to determine if words W_1 and W_2 are considered to have the same root. That is, the weighted graph is turned into an unweighted graph by pruning all edges whose weights are too low. The resulting graph will not be transitively closed, and it is not clear whether the resulting system returns pairwise judgements that do not form an equivalence relation, or whether transitive closure is carried out, or only cliques are retained.

They evaluate on words occurring at least 10 times in their corpus of 6.7 million words of newswire for English (2.3 million for German, 6.7 million for Dutch). Only words occurring in CELEX are scored. They use a relational measure for scoring, computing the number of true positives (which they call correct), false positives (which they call insertions) and false negatives (which they call deletions). this allows them to compute F measures, which they do for their baseline system, which was 85.2/88.3/82.2 for English/German/Dutch. Adding the orthography based on edit distance increased F-measures to 85.7/89.3/84.5, with syntactic co-occurrence modeled by immediate left and right contexts it was further increased F-measures to 87.5/91.6/85.6, with the final addition of transitive closure increasing this to their best reported F-measures 88.1/92.3/85.8. They perform the same evaluation for a pre-publication version of Linguistica (Goldsmith 2001), reporting scores of 81.8/84.0/75.8.

Even if the same-stem selection problem could be solved perfectly over a corpus, it does not handle the case of novel forms that may be seen in

queries, which is particularly problematic for low-frequency words, which includes most words in highly-inflected languages.

To further muddy the water, Schone and Jurafsky then pose thresholds for which they decide whether x and y have the same root if and only if the probability estimate is high enough. The end result is thus inconsistent in that the relation of “same stem” is not transitive.

7.4.18 Baroni et al., 2002

Baroni et al. (2002) introduce a fairly straightforward ad-hoc statistical approach to morphology extraction. They frame the problem as taking a raw corpus in and returning a ranked list of “morphologically related pairs”. The only indication of what a morphologically related pair is, is provided by the fact that they evaluate against a reference analysis produced by the Xerox morphological analyzer (Karttunen et al. 1997).

Baroni et al.’s basic approach involves scoring pairs against each other independently of the rest of the corpus. Each score is the sum of an orthographic similarity score and a contextual similarity score.

The orthographic similarity between two pairs is their standard edit distance (presumably without transpose), divided by the length of the longer of the two strings. This division ensures scores fall in the range $[0, 1]$.

The contextual similarity is defined by mutual information estimate based on a maximum likelihood estimates:

$$\hat{I}(w_1, w_2) = \log \frac{\hat{p}^{\text{ml}}(w_1, w_2)}{\hat{p}^{\text{ml}}(w_1) \cdot \hat{p}^{\text{ml}}(w_2)}$$

where the \hat{p}^{ml} are maximum likelihood estimates of some kind; the article is not clear on how they’re estimated. Presumably, they follow Church and Hanks (1990), who take the individual word estimates over the whole corpus and the pairwise ones over windows of 500 words. Baroni et al. modify this approach by eliminating words within a window of 3 characters, thus requiring the two words to be from 4 to 250 characters away (assuming boundaries are inclusive; the article is unclear). In any case, the two numbers are additional free hyperparameters which are presumably heuristically optimized by the authors; they do not say how 3 and 500 were chosen.

Mutual information should be between marginal and joint probabilities, but here the marginals are estimated very differently than the joint probabilities, which may or may not cause problems. Ideally, the estimates $\hat{p}^{\text{ml}}(w)$ would not be per-word, but per 500 word window, thus normalizing them to the joint estimates $\hat{p}^{\text{ml}}(w_1, w_2)$ which are estimated this way.

No matter how the estimates are made, mutual information remains a kind of weak independence test that does not take into account variance, which is influenced both by the probability estimates and counts. They credit Brown et al. (1990) with the observation that high mutual likelihood pairs are often morphological variants.

In the final ranking, both the maximum likelihood and edit distance scores are “normalized” to the interval $[0, 1]$ by dividing each value by the maximum possible value. This clearly doesn’t make much sense in the case of mutual information values, which are a kind of log odds. That is, if z is the maximum value, then:

$$\frac{1}{z} \cdot \log \frac{\hat{p}^{\text{ml}}(w_1, w_2)}{\hat{p}^{\text{ml}}(w_1) \cdot \hat{p}^{\text{ml}}(w_2)} = \log \left(\frac{\hat{p}^{\text{ml}}(w_1, w_2)}{\hat{p}^{\text{ml}}(w_1) \cdot \hat{p}^{\text{ml}}(w_2)} \right)^{\frac{1}{z}}$$

Because of the division by the maximum value, the base of the log will not matter.

Edit distances, because they are divided by the longest substring, are already “normalized” to the $[0, 1]$ scale. Although the paper states they are also divided by the maximum value, presumably that is 1.0, as that would result from comparing two words of length n that mismatch on every character.

For English, they use the rather small Brown corpus of roughly one million tokens. They evaluate on only the first 5000 pairs, taking the output of the Xerox toolkit as the reference. The authors then considered the pairs that the toolkit said were not related, and overruled the decision if the authors thought they were related enough; they say they were “conservative” in these decisions. Unfortunately, this is unfair, as they don’t consider the cases where the toolkit itself may have made false positives relative to the true, intended relations. They report results of roughly 50% for German and English at the 5000 pairs level. This seems very low, given that the overall number of relations that are required for a sample of 20,000 English words is at least an order of magnitude larger; they do not report results for the Xerox toolkit on the total number of morphologically related items in their corpus.

7.4.19 Creutz and Lagus, 2002

Creutz and Lagus (2002) motivate two heuristic approaches to unsupervised morpheme analysis using minimum description length.

Creutz and Lagus note that the description length of a text is the length of the model plus the length of the text being encoded given the model.

They provide a defective description length by only encoding the symbols of the model and failing to encode the end-of-symbol marker to normalize the symbol models properly. Their search algorithm operates heuristically in a mixture of online and offline modalities. As a word is scanned, all possible segmentations of it are considered, and if any reduce the overall cost, they are accepted. During learning, they occasionally revisit earlier words and attempt to resegment them given the current status of the codebook. They do not go back and iterate over the entire corpus. Note that the method is able to extract arbitrary numbers of morphemes, returning analyses such as *affect+ion+ate* in English and *eläin+lääkäri+lle* in Finnish.

Like the approach of Kit and Wilks (1999), Creutz and Lagus’s greedy algorithm is not guaranteed to find the segmentation leading to the minimum description length. Training over a tiny corpus of 100,000 words of Finnish news text, the resulting encoding is 22 bits/word.

The second method they consider is motivated by simple maximum likelihood, which is just like MDL, only without the cost of encoding the model. They approximately compute this result with a heavily heuristic variant of EM. They kick off the iteration by taking each word and then breaking off chunks by sampling lengths from a Poisson(5.5) distribution; this decision is not motivated. They then enter a reestimation loop. Rather than computing expectations, they take the most likely segmentation, leading to a Viterbi-like variant of EM that has been shown to be inferior in every setting for which it will converge. To sidestep problems created by the two previous assumptions, they also apply a heuristic that rejects segments if they are rare, if they consist of more than one single letter in a row.

They evaluate their results against a reference derived from running Koskenniemi’s (1983) rule-based stemmer, using a kind of EM-derived algorithm to align system responses against the reference results. They claim modest improvements over *Linguistica*, with all systems performing below the 50% level on the Finnish test data evaluated against the gold standard.

7.4.20 Neuvel and Fulop, 2002

Neuvel and Fulop (2002) work with a part-of-speech tagged lexicon and attempt to extract morphological relationships in order to generate unseen well-formed words. This goal is quite different than the more typical goal of analyzing unseen words into morphemes.

Their system is called the “Whole Word Morphologizer” and is based on the “whole word theory” of morphology. The heuristic matching shares many features with (Harris 1967), with various counts pushing matches over

thresholds given the part-of-speech contexts in which they are found, as well as aspects of Jacquemin’s (1997) phrase based account.

They evaluate on tiny 3000 word lexicons from novels.

7.4.21 Snover and Brent, 2002

Snover and Brent (2002) cite the desire to have a fully probabilistic model with properties like minimum description length. As usual, a Bayesian prior over the space of models does the same work as the model description length. But rather than defining a likelihood function and placing a prior over models, they introduce fixed hierarchical model that jointly generates paradigms, stems, suffixes and a corpus, which they confusingly call a “prior”.

Snover and Brent’s hierarchical model involves the following seven steps:

1. (a) Choose the number of stems $M \propto 1/M^2$
 (b) Choose the number of suffixes $X \propto 1/X^2$
2. (a) For each stem, choose its length $k \propto 1/k^2$.
 (b) For each suffix, choose its length $k \propto 1/k^2$.
3. For each stem and suffix, generate its characters according to maximum likelihood estimates.
4. Choose the number of paradigms D from the uniform distribution from 1 to X . (Note that the paradigms will partition the suffixes.)
5. Place suffixes in paradigm according to a uniform distribution over the paradigms.
6. Place stems in paradigms with proportional to the number of suffixes in the paradigm.

Note that there is no interesting posterior to reason about here, because the data z is fully determined by the model θ (as noted by Brent and Snover). Consider Bayes’s law for this case:

$$p(\theta|z) = \frac{p(z|\theta) \cdot p(\theta)}{p(z)}$$

The term $p(z|\theta)$ is 1.0, and hence the posterior $p(\theta|z)$ is just proportional to the prior $p(\theta)$, because $p(z)$ is fixed.

This trivial likelihood function and complex hierarchical “prior” do not lead to a natural Bayesian inference scheme. Thus Brent and Snover are left with a heuristic approach in which they attempt to use hypothesis testing to find paradigms with at least two stems and two suffixes. These “sub-hypotheses” are then combined into larger models. Finally, these larger models are rebalanced by allowing stems to move to new paradigms.

Brent and Snover evaluate their model using the relational method, which they apparently independently discovered, considering precision and recall over the conflation decisions made by the stemmer against a gold standard. Rather than requiring an equivalence relation over stems, they argue for an arbitrary graph. They consider **build**, **building** and **builds** to be related, because they share the verbal stem **build**. They also consider **building** and **buildings** to be related, because they share the nominal stem **building**. But they do not judge **buildings** and **build** to be related. This introduces a rather subtle judgement involving meaning, as the noun **building** is a nominalization of the verb **build** which refers to the product of the action; this morphological relation is quite regular for verbs of construction such as **carve** or **write**.

They report F-scores over this measure of roughly 0.8 for various lexicon sizes ranging up to 16,000. For Polish, the F-measure was roughly (0.7).

7.4.22 Johnson and Martin, 2003

Johnson and Martin (2003) extract morphemes with a version of Harris’s (1967) notions of successor frequency and predecessor frequency. They modify one of Hafer and Weiss’s (1974) heuristics which required a successor frequency and predecessor frequency above a constant to signal a morpheme break.

Johnson and Martin first construct the unique deterministic finite-state automaton that accepts exactly the training set of words. They then extract pairs of (not necessarily distinct) nodes such that (a) there is a unique path from the first node to the second node, and (b) the first node has in-degree greater than one, and (c) the second node has an out-degree greater than one. The initial and final nodes then constitute morpheme breaks, with the characters labeling the transitions on the unique path providing the string representing the morpheme extracted between the break points. They call the pair of nodes a *hub* and say that it is *stretched* if the first node is not equal to the second. It seems necessary to require initial nodes and final nodes to constitute possible boundaries, in general.

The notion of hub is not stable in the face of more data. As more data

is added, existing hubs may disappear, because the path between them is no longer unique. For instance, if we had **unhelpful**, **unhelpless**, **helpful** and **helpless**, we would have an automaton with a path from the start node along **un** back to the start node, a unique path along **help**, then a branching that allows completion with **less** or **ful**. We'd extract **help** as a morpheme. But what if we added in the word **helping**. Now all of a sudden **help** is no longer extracted.

After extracting the basic set of hubs, Johnson and Martin take an additional step of merging hubs. They merge states corresponding to final states and the initial node of any hub with an in degree of three or more. This strategy will mitigate the problem of hub instability as long as there is enough evidence. But it will also predict a large number of non-words, by allowing any word to be followed by any suffixes resulting from a high successor frequency hub.

Note that there is no way for this system to deal with new stems.

Their evaluation was over a tiny corpus: the text of Mark Twain's book *Tom Sawyer*. They created a reference morpheme extraction without describing their procedure other than to say it was based on dictionary entries and the judgements of two native speakers of English. They outline a very generous notion of acceptable breaks, citing **muddy** as an example that would be analyzed as **mud-d-y**, which would match either **mudd-y** or **mud-dy**; this is where most of the errors arise in segmentations and this may allow the conflation of elements ending in **e** such as **pal** and **pale**, depending on the reference standards. Over that corpus, they report Goldsmith's Linguistica (2001) to have recall 0.57 and precision 0.91, with their first hub searching strategy having recall 0.45 and precision of 0.92, with their modified merged system having recall of 0.59 and precision of 0.92. The merged system has slightly better performance than Linguistica, but it's unclear how performance would scale to larger corpus, despite the author's claims that more data would improve results.

They provide a one-paragraph sketch of an application to Inuktitut, a native Canadian language with highly productive morphology. Drawing data from the Canadian Hansards, they showed a precision of 0.31 and recall of 0.08, citing underextraction of morphemes as the main problem.

7.4.23 Argamon et al., 2004

Argamon et al. (2004) present a minimum description length motivated model of multisegment morphological analyses, such as **inter-nation-al-ist**. They seem not to have noticed the extensive previous literature doing ex-

actly this (e.g. Deligne and Bimbot 1997, Nevill-Manning and Witten 1997, and many papers thereafter reviewed above).

Like most of the other heuristic approaches to MDL, Argamon et al. use an initial simple assignment followed by successive greedy rearrangements. Much of the paper is concerned with data structures for computing these greedy steps effectively; with a full subsequence count of character sequences in words, the necessary computations are trivial. They call this idea “local description length models”, which is just another way of saying words are generated independently given the model.

They consider a number of “models” of model length. The first of these is just the number of words in the dictionary:

$$\text{descLen}(\theta) = \text{numWords}(\theta)$$

Obviously, this is not sufficient for properly coding the model, as the lengths of words and the identity of their characters are also required.

Argamon et al.’s second model is a heuristic attempt at modeling something akin to likelihoods where the cost of modeling a term in the dictionary is inversely proportional to its frequency; a proper MDL description of a model should, of course, include codings of the symbols and their likelihoods. I couldn’t understand the details of their formulas, which appear to produce negative numbers for code costs.

Their third model is an attempt at a “proper” MDL model. They encode all of the symbols as part of the model using a constant encoding scheme. As usual in MDL approaches to morphology, they fail to encode the length of a word or an end-of-word symbol, leaving their symbol generating process defective. They also do not encode the model’s probabilities.

Their data description length is generated by independently generating the morphemes according to maximum likelihood estimates in the corpus:

$$\text{descLen}(m_0, \dots, m_{k-1}) = \sum_{i < k} \log_2 p(m_i) = \sum_{i < k} \log_2 \frac{\text{freq}(m_i)}{k}$$

With this formulation, the trivial solution taking each word to be a unique morpheme is not only the maximum likelihood solution, but also typically the MDL solution with any sizeable data and a fixed model encoding cost such as the one shown here (see Peng and Schuurmans 2001, or the discussion of their paper above). In short, even if the suffix *-s* is very common, the independence assumption means that for a given word such as *eats*:

$$p(\text{eat}) \cdot p(\text{s}) << p(\text{eats})$$

It's not clear why Argamon et al.'s algorithm didn't fall into that trivial solution.

7.4.24 Monson et al., 2004

Monson (2004) and Monson et al. (2004) provide a partially ordered set representation of a corpus. Monson's representation involves a node for each set of suffixes that shares at least one stem. Nodes are ordered by their sets of suffixes. Although Monson calls it a "lattice", it is not a lattice in the algebraic sense, because there are not unique joins. Note that an identical ordering of potential paradigms was described by Brent and Snover (2002, section 3.1).

The upper reaches of this poset are supposed to suggest morphological paradigms in much the same way as Harris's (1955) heuristics. Monson applies a χ^2 test for the independence of a pair of suffix by considering their distribution with stems. For instance, in Spanish, the consider the following tables:

	+a	-a		+a	-a
+as	199	205	+tro	2	14
-as	1038	21508	-tro	1235	216999

Although Monson suggests that these numbers support the dependence of suffixes **a** and **as**, whereas they support the independence of **tro** and **a**. Tests like these, or even simpler collocation tests are useful, in that they adjust raw counts for variance. Note that they are only good approximations when the counts in all cells are at least 5, a condition violated by Monson's second example.

Monson provides no suggestion of how to use these dependencies to select nodes as paradigms and no evaluation.

7.4.25 Namer and Zeigenbaum, 2004

Pierre Zweigenbaum has published extensively in medical terminology extraction from corpora and structured resources; this section summarizes a recent paper with Fiammetta Namer. They provide heuristic word formation rules that allow words to be analyzed hierarchically in terms of "combining forms", consisting of stems and morphemes. As an example, they provide

```
dsintoxication/N => [[[d [ in [toxique A] (er) V]] A]tion N],
  (dsintoxication/N, dsintoxiquer/V, intoxiciquer/V, toxique/A)
  "(Action|rsultat de) de dsintoxiquer"
```

("(Action|result of) detoxicate")

They carry out analyses by means of word formation rules, which are classical rewritings:

`dXiser V --> [d [X' N] +iser V]`

which maps a noun to a verb by circumfixing it with `d-` `-iser`.

7.4.26 Wicentowski, 2004

Wicentowski (2004) extends and generalizes his earlier model, described in (Yarowsky and Wicentowski 2000). The main difference between the later model and earlier one is that the later one models not only suffixes plus point-of-suffixation changes in the stem, but also models stem-internal vowel changes. For instance, the pairs in Spanish **acuerdo/acortar** are handled with an internal vowel change mapping **ue** to **o**; this pattern is regular, showing up in other pairs such as **apruebo/aprobar** and **muestro/mostrar**.

Wicentowski also extends the system to prefixes, and evaluates over dozens of languages.

Unlike the earlier model, the current model is supervised through a list of inflection/stem pairs. It optionally accepts lists of prefixes, suffixes and potential roots. It further requires a list of vowels, which it uses to target stem-internal vowel changes.

Inflection/root pairs are deterministically aligned to find their longest matching overlap given the prefix, suffix and vowel-change possibilities. We only consider suffixes, as prefixes behave similarly. A word consists of a stem, a vowel change, a suffix, and a point-of-suffixation change. For example, the word **acuerdo** is analyzed as **ac.ue.rt.o**, with **-ue-** being the internal vowel and **-o** the suffix, and **ac-** **-rt** being the fixed part of the stem. Similarly, **acortar** is analyzed as **ac.o.rt.ar**, sharing the two consonant strings, but differing in internal vowel **-o-** and suffix **-ar**.

Suppose we have a word analyzed as chunks R_1VR_2DS , where R_1 and R_2 are the fixed stem components, V is the internal vowel, D is the point-of-suffixation material and S is the suffix itself. Then stemming returns the stem $W = R_1V'R_2D'S'$ that is most likely to have generated it:

$$\text{stem}(R_1VR_2DS) = \text{argmax}_{W=R_1V'R_2D'S'} p(R_1VR_2DS|R_1V'R_2D'S')$$

Wicentowski assumes that suffixes and word-endings are all equally likely, the vowel change is context free, and the point-of-suffixation change is conditioned on the result thus reducing this problem to:

$$\text{argmax}_{W=R_1V'R_2D'S'} p(R_1VR_2DS|R_1V'R_2D'S')$$

$$= \operatorname{argmax}_{W=R_1V'R_2D'S'} p(V|V') \cdot p(D|R_1VR_2D'S)$$

Prefixes are handled in the same way as suffixes, and simply add another substring to the sequences and multiply the overall root estimate by the prefix change estimate.

Performance is excellent for the evaluation over verb inflections, averaging around 97.5% over a range of more than thirty languages. Wicentowski notes that the affix lists did not help much for most languages and even hurt some, such as Portuguese; he speculates this is due to the partial nature of these lists. He also notes that the Yarowsky and Wicentowski (2000) model outperformed the current model in some languages. He further notes that the vowel change estimates without context hurt the overall performance. He also notes that the performance on very irregular forms was not good, and suggests corpus co-occurrence counts to help with that problem. Finally, Wicentowski noted that if the heuristic of requiring roots to be in a broad coverage word list, coverage fell to 97.4% but precision increased from 97.5% to 99.1%.

7.4.27 Bernhard, 2005

Bernhard (2005) provided the winning entry for MorphoChallenge 2005 for Finnish and Turkish. Her approach integrates three strategies. First, she uses segment predictability, in the sense of Harris (1955), citing (Saffran 1996) as a variant based on transition probabilities. Second, she extends Neuvel and Fulop’s (2002) notion of word similarity at boundaries to arbitrary substring similarity. Her third strategy is a heuristic based on stem and suffix length motivated by Zipf’s comment:

“the length of a morpheme tends to bear an inverse ratio to its relative frequency of occurrence.” (Zipf 1968, p. 173).

She presumes affixes will be frequent and stems infrequent, concluding that affixes will be short and stems long. Surprisingly, Bernhard’s method does not use word frequency information directly in any way.

Bernhard extracts four categories of morphological segment: stems, prefixes, suffixes and linking elements. She cites the third *o* in *homonothe^orapy* as an example of a linking element. She then restricts the allowable sequences of these segments.

She first extracts potential segments from long words, breaking a word $W = U \cdot V$ into non-empty components U and V if the value of the average value of the maximum transition likelihood from a non-empty suffix of U to

a non-empty prefix of V or vice-versa.

$$f(U, V) = \frac{1}{\text{len}(U) \cdot \text{len}(V)} \sum_{U=U_1 \cdot U_2} \sum_{V=V_1 \cdot V_2} \max(p(U_2|V_1), p(V_1|U_2))$$

We are not aware of any probabilistic interpretation of this value. She notes that for **ultracentrifugation**, the local minima after plotting f suggest breaks **ultra**, **centrifug** and **ation**, because $f(\text{ultra}, \text{centrifugation})$ and $f(\text{ultracentrifug}, \text{ation})$ are local minima of the function f that are at least a constant value less

$$f(\text{ultra}, \text{centrifugation}) < k + f(\text{ultr}, \text{acentrifugation})$$

and

$$f(\text{ultra}, \text{centrifugation}) < k + f(\text{ultrac}, \text{entrifugation})$$

Bernhard sets k to be the standard deviation of all of the $f(U, V)$ values for the string. Using this technique on the corpus, she extracted the following prefixes:

in-, pre-, un-, natur-, inter-, counter-, dis-, over-,
mis-, psycho-, re-, ultra-, ex-, hyper-, pseudo-, con-

and the following suffixes:

-s, -ly, -e, -ble, -ed, -tion, -al, -es, -ally, -ately, -ing, -ity, -ation,
-l, -ness, -ism

Bernhard then extracts stems by stripping off all affixes from the previous lists to get a list of candidates. This list is then whittled down by requiring stems (1) to be at least three characters long, (2) to have at least two letters following them in the corpus, (3) to not contain hyphens, and (4) to occur at least once as an initial segment. This last condition assumes that the language does not typically require prefixes.

All of the words containing a particular stem are considered. For instance, Bernhard considers **hous**, with words **hous-ekeeping**, **hous-ing**, **hous-ehold**, **hous-e's**, **hous-e**, and **hous-ed** (note that in general, there may also be one or more suffixes or prefixes). Most of these suffixes show up on the suffix lists; the others lead to potential new stems **-ekeeping** and **-ehold** and a new potential suffix **-e's**. These new stems and suffixes are evaluated collectively, requiring for some hyperparameters a and b :

$$\frac{\text{knownSuff} + \text{newStem}}{\text{knownSuff} + \text{newStem} + \text{newSuff}} > a$$

and

$$\frac{\text{knownSuff}}{\text{knownSuff} + \text{newStem}} > b$$

Another heuristic iteration is undertaken where validated new stems and suffixes are added and the procedure repeats. Then, the number of instances of each segment is counted across all analyses of all words that contain it. Finally, the best segments are chosen using a best-first heuristic search.

With a final list of segments, Bernhard assigns a cost to each and uses A* search to find the lowest cost segmentation; dynamic programming seems a slightly more natural choice for the size of the problem. She evaluates two cost functions. The first is a standard maximum likelihood Viterbi-style cost, where the cost of a segmentation is the sum of the log probabilities of the segments, assuming each segment is generated independently. The second model assumes a similar form, but rather than the log of the segment frequency divided by the sum of all segment frequencies, it is taken to be the segment frequency divided by the maximum segment frequency. Obviously, the first measure will penalize each segment more, and thus prefer longer segments. The second is unmotivated probabilistically.

She provides a conflated-pairs style evaluation, for which she lists 62.5% precision, 43.0% recall, for a balanced F-measure of 51.0% over her own data set and gold standard annotation of the top 5,000 keywords in the corpus; keywords are identified as in (Rawson and Garside 2000). She cites **lymphedematous/lympoedema** as a true positive, **additive/addresses** as a false positive and **therapeutics/therapy** as a false negative.

She provides final evaluation measures (optimizing the three hyperparameters on the sample labeled datasets), and reporting both cost methods (maximum likelihood and heuristic). For English F-measure was 66.6% for method 1 and 62.4% for method 2; for Finnish, 63.3% for method 1 and 64.7% for method 2; for Turkish, 55.3% for method 1 and 65.3% for method 2. Unfortunately, she doesn't provide precision/recall figures here, but does note that method 2 has higher recall and lower precision on all data sets.

She did mention some lack of recall from not conflating pairs with orthographic variation such as **cancer/cancér**.

Although not reported in her paper, her system also did the best of all of the systems in the second evaluation, n -gram language modeling with the segments generated (Kurimo et al. 2005).

7.4.28 Bordag, 2005

We describe Bordag’s (2005b) MorphoChallenge system, which includes the morpheme extractor described in (Bordag 2005a) as a first component. Like Harris (1955), he counts successor/predecessor variety statistics. Unlike Harris, he does not consider the entire word list, but rather restricts attention to words that are similar to the word being analyzed in both form and corpus distribution.

Instead of using the entire word list for the statistics, he finds 150 words that are most similar to a given word by substring match statistics; this is a k -nearest neighbor (KNN) approach, with $k = 150$.

The first measure of similarity is substring similarity, as measured by edit distance. Under this measure, **clearly** and **early** are similar, and thus likely neighbors.

The second measure of similarity is contextual. Contextual similarity is measured by computing by representing a word by the set of words that appear with it in the corpus. Bordag mentions that such vectors may be collected from words within a sentence or just neighbors, but it’s not clear exactly which setup was used in the paper. Nor was it exactly clear which metric was used over the vectors, as both Euclidean distance and cosine were mentioned. Bordag does note that under his measure, **clearly** is similar to **closely**, **legally**, **weakly**, making them likely neighbors.

Working through an example, Bordag notes that for **early**, there are only four letters preceding **-ly#** among its 150 nearest neighbors. For **clearly**, there were 16 different letters appearing before **-ly#** among its 150 nearest neighbors.

Bordag introduces another modification, which he refers to as “substring frequency”. He uses the set of nearest neighbors for a maximum likelihood estimate of $p(\text{ly\#}|\text{y\#})$. High estimates favor **-ly** as a suffix, whereas low estimates favor its independence. Among the neighbors of **early**, 6/19 words ending in **-y** end in **-ly**, whereas 76/90 words in the neighborhood of **clearly** ending in **-y** also ended in **-ly**. This measure is like a confidence test for the independence of **-ly**, though it is neither adjusted for variance nor for the probability of **-l** itself.

As another heuristic adjustment, Bordag looks at the frequencies of character n -grams, noting that **th** acts more like a single character based on its frequency, and is thus unlikely to be broken. The sequence **rl** in **early**, on the other hand, is relatively rare as a bigram and thus a likely candidate

to be split. Decisions are weighted by inverse bigram weight, defined to be:

$$\text{invBigramWgt}(c_1c_2) = 1 - \frac{\text{freq}(c_1c_2)}{\max_{c'_1, c'_2} \text{freq}(c'_1c'_2)}$$

The final score Bordag uses is left plus right successor variety among the 150 nearest neighbors, multiplied by inverse bigram weight and substring frequency:

$$\text{score}(M_1 \cdot c_1, c_2 \cdot M_2) = p^{\text{ml}}(c_2M_2|M_2) \cdot \text{invBigramWgt}(c_1c_2) \cdot (\text{succ}(M_1) + \text{pred}(M_2))$$

Words are then segmented at their highest scoring points.

In order to apply to unseen data, Bordag builds very simple decision trees with features selected by ordering longest suffix matches. For insan instance, he cites a suffix tree trained on **clear-ly**, **strong-ly** and **early**, leading to the trie suffixes, namely entries **yl-**, **yl-** and **ylrae**. If we see a new word like **daily**, the node **yl** branches to either end-of-suffix or the word **early** with a 2:1 count ratio. Therefore, the suffix **-ly** is extracted. But if we see **early**, the longest matching node is **ylrae** and there are no alternatives other than to take it to be a token.

For the MorphoChallenge evaluation, the system did well for German and English, with straight-LSV having higher precision and much lower recall than the system stacked with the trie-based classifier. The combined system ahd 0.688/0.721 precision/recall for German and 0.529/0.526 for English. Performance was much lower for Finnish and Turkish, which combined much smaller corpora and much richer morphology.

7.4.29 Creutz and Lagus, 2005

Creutz and Lagus extend their (2002) approach with a range of heuristics. They note, as did Peng and Schuurmans (2000), that frequent words tend to remain unsplit even using minimum description length (MDL) priors.

Creutz and Lagus state their model as an HMM, with distinct states for stems, suffixes, prefixes and non-morphemic strings, with the expected constraints on transitions. For instance, **forward** is analyzed as a non-morphemic unit **for** and a stem **ward**. Their figures (e.g. figure 2) show hierarchical structure which would suggest a context-free grammar rather than an HMM. For instance, **straightforward** is composed of two stems, **straight** and **forward**, with the latter analyzed as above.

Creutz and Lagus set out to apply maximum a posteriori (MAP) estimates, but their heavy use of heuristics ensures that their point estimates are not local maxima.

Creutz and Lagus introduce a number of generative models, most of which are defective in one or more ways. For instance, the probability of a lexical item is defined to be the product of the probability of its meaning and its form. But the form fully determines the unit in their model, whereas the meaning is an extrinsic measure of context, frequency, etc. This does not constitute a proper prior; it also says nothing about the likelihood of different distributions over the lexical entries.

For the form term, they allow the form to be composed either of a sequence of characters or as a sum of smaller forms. This imposes an implicit exponential length model, which does not match word length distributions very well. When they allow segmentations, they only consider a single segmentation, when they should consider a sum over all segmentation. They also fail to normalize the sequence decompositions with end-of-sequence statistics, which adds another defective component to the model.

Creutz and Lagus derive the meaning features by multiplying a number of probabilities. First, the probabilities of the frequencies, for which they introduce an unusual inverse binomial prior that is highly uninformative. Another term is “left intraword perplexity”, which for a stem and suffix would be the entropy of the distribution over stems given the suffix; the rightward version is also multiplied in, and none of these values are normalized, though they mention Rissanen’s universal prior for positive numbers, which is only defined for integers, whereas their perplexities will be floating point.

Over these models, Creutz and Lagus then define a heavily heuristic inference procedure, which involves splitting and joining, all within a Viterbi reestimation loop.

7.4.30 Freitag, 2005

Freitag (2005) introduces an approach to bootstrapping general morphological relations from terms that are clustered on the basis of syntactic cooccurrence.

Freitag begins by clustering terms in a corpus based on their local contexts. For each word, he generates a separate count of words occurring immediately to the word’s left and immediately to the word’s right. As he notes, this typically results in a syntactic notion of cluster, with terms that play the same syntactic role (e.g. adjective or auxiliary verb) showing up in the same clusters. As Freitag also notes, this contrasts with the approach of Xu and Croft (1998) or Schone and Jurafsky (2001), which use sets of terms occurring at greater distances, resulting in a more topic-like clustering.

Freitag’s term clustering uses information-theoretic clustering (Dhillon et al. 2003). The main idea here is to maximize the mutual information between the terms in the clusters and the terms that show up in their contexts:

$$\text{MI}(X, Y) = \sum_{x,y} p(x, y) \log \frac{p(x, y)}{p(x) \cdot p(y)}$$

The random variable X generates a term from a cluster according to $p(x)$, and random variable Y generates terms from the contextual cluster according to $p(y)$. Freitag uses a kind of greedy initial clustering with a heuristic post-process that tries to move terms between clusters to increase mutual information. Freitag does not explain how he prevents the obvious overfit solution of one term per cluster.

Freitag then looks for rules that transform words in one cluster into words in another cluster. Such rules were introduced by Mikheev (1997) for inducing part-of-speech taggers.

Freitag represents affixes as general regular expressions. Freitag uses the regular expression syntax of the Perl programming language. For instance, the regular expression `e?d$` matches the suffixes `ed` or `d`; the question mark (?) indicates the `e` is optional, and the dollar sign (\$) indicates the end of a word. The prefix `re` appears as the regular expression `^re`; the carat (^) indicates the beginning of a word.

Transforms are represented as Perl substitutions, which use the binary syntax `s/././`. For example, `s/ed$/ing/` represents a transformation that replaces a word-final `ed` with `ing`.

Candidate transformations are generated by comparing prefixes and suffixes across term clusters, retaining mappings that apply to at least three members. Freitag employs the usual heuristic of removing mappings which are subsumed by other mappings because of redundant material; e.g. the general transform `s/s$/ed/` subsumes the transform `s/ts$/ted`, and thus the latter is removed. This presents opportunities to combine transforms. For example, the transforms `s/es$/ing/` and `s/s$/ing/` may be unioned into the single transform `s/e?s$/ing/`.

Rather than a stright MDL or Bayesian approach to deciding which candidate stems are best, Freitag employs a graph-theoretical model from social network analysis (SNA), citing Kleinberg’s (1998) application of the hubs and authorities social network analysis model to the web. The graph structure has links from stems to transforms in which they participate. For example, assuming a reasonable set of clusters and rules, the stem `eat` would be linked to `s/s$/ing/`, because there would be a cluster containing `eats` with

the rule `s/s$/ing/` mapping that cluster to a cluster containing `eating`. In the network flow algorithm, all weights start at a uniform value, then for each stem linked to a transform, the weight of the stem is added to the weight of the transform; I would have expected this to somehow be spread out over the outflow as in most SNA algorithms (such as PageRank). Then stem weights are all set to zero, and weights are propagated back from transforms to stems. This is iterated five times.

The highest-weighted transforms as a result of this SNA algorithm look very similar to everyone else’s lists for English, with the exception of the combined rules. The three highest weighted rules are `s/$/s/`, `s/e?$/ed/`, and `s/e?$/ing`; the three lowest weighted pairs are: `s/s$/ses`, `s/w$/ws`, and `s/^b/c/`.

In another heuristic pass, every word is segmented according to its highest scoring transform. Only transforms participating in at least one segmentation are maintained. (In some ways, this is like a winner-take-all round of EM training.)

Finally, the surviving transforms are combined into a stemmer. If there are competing analyses, heuristics are applied to select transforms to apply to reduce a word to its stem. In cases of multiple possible analyses, the steps of the analysis (referred to as a “trace” in the paper), are subjected to heuristic rules. These heuristics prefer shorter stems, fewer edits in the transforms, transforms that apply to a lot of words, and to prefer more steps in the analysis to fewer.

Freitag performed experiments on data sets of the most frequent words of English drawn from the *Wall Street Journal*. Truth is derived from the CELEX morphology database (Baayen et al. 1995). The largest corpus was 20,000 words, though this is still quite small. Frequency information is not used, though co-occurrence is used for clustering terms initially. Freitag assumed 200 clusters for the experiments, and does not mention what happens with other sizes. Using relational scoring that counts true positives as relations between words that share a stem, the performance on the top 20,000 words had a precision of 0.80 and recall 0.82. Freitag also reports an analysis against CELEX itself, measuring a 71% overall accuracy; 7% of the words were understemmed, 14% of the words were not stemmed at all when they should have been, 2% of the words had spurious stems in the sense of being single words in CELEX that the system thought had morphological structure, and finally, 5% of the results were other forms of incorrect cases.

In error analysis, Freitag notes that at 20,000 words, several spurious transforms enter the set, including `s/$/e/` and `s/$/o/`, grouping together names such as `Clark/Clarke`, `Brook/Brooke`, `Robert/Roberto`, etc. Freitag

@	g	r	i	m	@	n	t
-----	-----	-----	-----	-----	-----	-----	-----
Nuc_I	Ons_M_1	Ons_M_2	Nuc_M	Ons_F_1	Nuc_F	Cod_F_1	Cod_F_2
-----	-----	-----	-----	-----	-----	-----	-----
Rhy_I	Ons_M			Ons_F	Rhy_F	Cod_F	
-----	-----	-----	-----	-----	-----	-----	-----
Syll_I	Syll_M			Syll_F			
-----	-----	-----	-----	-----	-----	-----	-----
Word							

Figure 7.6: Goldwater and Johnson’s Syllable Model 1

considers possible ways to solve these problems, based on looking at the resulting cluster behavior.

7.4.31 Goldwater and Johnson, 2005

Goldwater and Johnson (2005) provide a statistically sound approach to estimating syllable structure from an unlabeled corpus of words. They also evaluate supervised models. They use the CELEX dictionary (Baayen et al. 1995) for German and English. The words themselves are sequences of phonemes, as in (Manning 1998).

Goldwater and Johnson’s thesis is that the structure of the model, what they call “representational bias”, is more important than factors such as initial parameters in achieving good learning results from data augmentation algorithms like expectation/maximization (EM).

They consider two specific model structures, both of which are probabilistic context free grammars. Their common structure allows the inside/outside algorithm to be used for the expectation step of the EM algorithm.

In both models, syllables are broken down into onsets (initial consonants) and rhymes (vowel and final consonants), with rhymes broken down into nuclei (the vowel or dipthong) and codas (the final consonants).

In the first model, each syllable, onset, rhyme, nucleus, and coda is marked as to whether it is part of the first syllable in a word, part of the last syllable in a word, part of a medial syllable in a word, or part of the single syllable in a one syllable word. An analysis of the word **agreeemnt**, with sequence of phonemes represented as **@grim@nt**, is in figure 7.6.

In the second model, the onset, nuclei and codas are not distinguished as to position in the word, whereas the syllables above them are subcategorized to indicate the syllable structure of their leftmost syllable. Figure 7.7 shows an analysis in the second model of the same word as in figure 7.6.

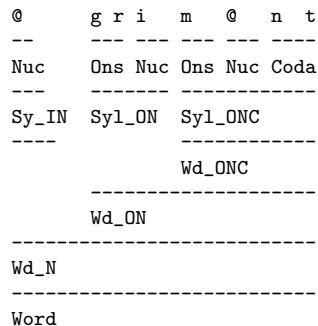


Figure 7.7: Goldwater and Johnson’s Syllable Model 2

For instance, **Syl_ONC** represents a syllable composed of an onset, nucleus and coda, such as **m@nt**. The category **Wd_ONC** also spans the same lexical material, representing a sequence of syllables beginning with a syllable composed of an onset, nucleus and coda. Note that **Syl_ON** spans **gri**, which contains an onset and nucleus; the category covering all of **grim@nt** is **Wd_ON**, representing the fact that the leftmost syllable is composed of an onset and nucleus. Only this leftmost syllable is relevant for further combination of the pair of syllables, because it will attach only with another syllable to the left. This is why Goldwater and Johnson call their second model a bigram model.

With supervised training for German, model 1 scores 97.4% and model 2 scored 97.2%. For English, model 1 scored 98.1% and model 2 scored 97.6%. Goldwater and Johnson point out that 49.161.2% of the English words were monosyllabic, which are guaranteed to be syllabified correctly. They only evaluated on words in CELEX.

Goldwater and Johnson propose a purely heuristic segmenter which they call a “categorical parser”. This parser prefers maximal onsets, but knows about the sonority hierarchy in order to eliminate illegal syllabifications. The maximal onset principle states that if a consonant may attach as the last consonant of the preceding syllable or the first consonant of the following syllable, then it should attach to the following syllable. The notion of legal attachment is then defined by sonority, with sonority decreasing away from the nucleus in both the onset and coda. As they note, languages sometimes rule out syllable structures that would otherwise be acceptable strictly on universal preferences for sonority contours and maximal onsets.

Goldwater and Johnson use the output of the categorical parser as the

starting point for an expectation/maximization (EM) algorithm. They show that accuracy after EM estimation drops precipitously with uniform initial estimates instead of the categorial parser’s output. The algorithm treats the parses of the words as missing data. The second model, with bigrams, outperformed the first. For German, the categorial parser had 92.7% and the EM trained parser 95.9% accuracy (supervised was 97.2%). For English, the categorial parser was 94.9% accurate and after EM, 97.1% accurate (supervised was 97.6%).

Goldwater and Johnson trace the difference between their models to a bias-variance tradeoff. The bigram model has a higher representational bias in that it builds into the model structure a shared (tied) distribution between onsets, nuclei and codas in every position. The positional model, on the other hand, has more variance in that it allows each syllable position to learn a different distribution of onsets, nuclei and codas. The positional model was particularly bad when used with uniform initialization in EM training; it found a local maximum consisting of coda maximization rather than onset maximization.

7.4.32 Hu et al., 2005a

Hu et al. (2005a) generalizes and introduces a few heuristics for operating in Goldsmith’s (2001) *Linguistica* framework. Recall that the *Linguistica* framework involves extracting signatures, defined as sets of stems and suffixes such that every pair of stem and suffix is in the corpus. Hu et al. generalize this notion to triples. Their goal is to find sets A, B, C of strings (possibly including the empty string) such that $a \cdot b \cdot c$ is a word for every $a \in A, b \in B, c \in C$. They say that a triple is a “template” if at most one of the sets is a singleton. The sets may contain prefixes, stems or suffixes; there is no typing, but extracted examples include, for example, a stem and two suffixes, or a prefix, stem and suffix.

Hu et al. represent templates A, B, C as sequential, four-state finite state automata. There are no complex automata operations required, and all automata and generated languages are finite.

Hu et al. use edit distance to find alignments between pairs of words extracted from a corpus. They do not seem to realize that alignments generated in this way are ambiguous. For instance, aligning **aabb** with **ab** according to a minimal edit sequence, allows the **a** to match either **a**, and the **b** to match either **b**, thus producing four equivalent alignments at edit distance 2:

aabb aabb aabb aabb

a b ab a b a b

They provide a heuristic set of edit costs with substitutions costing 1.5, deletions and insertions 1.0, and matches 0.0.

After generating the templates, these are collapsed using a series of heuristics. First, they score the sets by the difference in number of characters between the set of strings A, B, C and the set of all strings $a \cdot b \cdot c$ with $a \in A, b \in B, c \in C$.

$$score(A, B, C) = \sum_{a \in A} \text{len}(a) + \sum_{b \in B} \text{len}(b) + \sum_{c \in C} \text{len}(c) - \sum_{a \in A} \sum_{b \in B} \sum_{c \in C} \text{len}(a) + \text{len}(b) + \text{len}(c)$$

They call this score the “robustness” of a template.

With scores for templates, they consider pairs of templates A_1, B_1, C_1 and A_2, B_2, C_2 to collapse into a template $A_1 \cup A_2, B_1 \cup B_2, C_1 \cup C_2$. The requirement is that two of the three sets must be identical. This leads to the construction of further templates, because any $a \in A_1 \cup A_2, b \in B_1 \cup B_2, c \in C_1 \cup C_2, a \cdot b \cdot c$ is in the corpus.

In order to disambiguate a case where every string in a template shares a prefix or suffix, they score morphemes cumulatively by the sum of the scores of the templates in which they appear. They then disambiguate by looking at the sum of the cumulative scores of the strings in the analysis.

Finally, they allow templates to be further collapsed to form triples which are not templates. They do this if there is enough overlap between the sets being collapsed. They consider individual sets A, B and C to represent grammatical suffixes or prefixes if the set has five or fewer elements and a stem otherwise. They then union pairs under a complex restriction that if there are two affixes, then one set must be identical and the other must have at least two affixes in common. If the sets are stems, then they must have at least two strings in common.

Hu et al. evaluate over a tiny subset of a Bible translation into Swahili, consisting of 50,000 words, producing 7180 distinct words. They report roughly 89% precision and 46% recall, though it is not clear how the reference segmentation is determined.

Hu et al. provide a novel analysis of their last stage of collapsing; they submit the words that are generated to Yahoo to see if they have hits. They found 87% of the words did. It is difficult to interpret this result without taking some kind of sample to see if the words found were actually in the proper language and not typos, etc. General web text is particularly noisy and typo- and braino-prone.

7.4.33 Hu et al., 2005b

Hu et al. (2005b) wrote another paper in the signature-based Linguistica framework of Goldsmith (2001). Recall that a signature over a corpus (here, a set of words) is a set of stems and a set of suffixes, such that the concatenation of any stem with any suffix is in the corpus. This paper also introduces the notion of a signature transform of a word, which isn't actually a transform, but rather a pair consisting of a signature and a single suffix τ from the suffix list such that the word "is morphologically analyzed as a stem σ followed by suffix τ ." It is not clear what the "is morphologically analyzed as" is, as there is no labeled corpus under consideration.

Hu et al. introduce another heuristic into their MDL-motivated models in this paper, this time replacing high frequency words (more than 200 instances) with their transforms. Although not discussed in these terms by the authors, such a move is intended to reduce the overfitting of the optimal model according to the MDL criterion by imposing extrinsic constraints on model structure. That is, their solution under heuristics such as these are no longer the minimum description length solution to their MDL equations, thus contradicting their discussion about their MDL equation (1) (renormalized to description lengths; the original was a mix of entropy and description length):

$$\hat{\theta} = \operatorname{argmin}_{\theta} \operatorname{descLen}(\theta) - \sum_w \operatorname{descLen}(w|\theta)$$

This is the usual equation stating that the model parameters $\hat{\theta}$ are defined to be the minimizer of the sum of description length and the probability of the corpus given the parameters.

As usual, we can take the description lengths to be negative log probabilities. For the model parameters, that's:

$$\operatorname{descLen}(\theta) = -\log_2 p(\theta)$$

where $p(\theta)$ is the prior over parameters θ . Such a bound is possible through entropy coding, and is the tightest possible (smallest) such encoding.

Hu et al. define the probability of the data defectively, assuming that each word arises from a single signature under a single stem/suffix split point, assuming the analysis of a word $\omega = \sigma \cdot \tau$ as having stem σ and suffix τ and being drawn from signature s :

$$p(\omega) = p(s) \cdot p(\sigma|s) \cdot p(\tau|s, \sigma)$$

This should really be:

$$p(\omega) = \sum_{\omega=\sigma \cdot \tau} \sum_s p(s) \cdot p(\sigma|s) \cdot p(\tau|s, \sigma)$$

Given this understanding of their model, Hu et al. make several mistakes in encoding. For an unexplained reason, rather than encoding multinomial probabilities for stems σ using $-\log_2 p(\sigma)$ bits, they use $-\log_2 \text{freq}(\sigma)$ which is missing the normalizing term. It should be:

$$\text{descLen}(\sigma) = -\log_2 \frac{\text{freq}(\sigma)}{\sum_{\sigma'} \text{freq}(\sigma')}$$

And each of the three distributions in the word probability factoring needs to be encoded. Typically, the model probabilities are encoded, but Yu et al. use maximum likelihood estimates from the corpus. In addition, they encode the probabilities of all n outcomes, where clearly the n th is determined by 1 minus the sum of the first $n - 1$ probabilities (that is, there are $n - 1$ degrees of freedom, or an $n - 1$ dimensional simplex defining the space of parameters).

The interesting contribution of this paper is the introduction of a dependence on part-of-speech tags. The data is assumed to consist of part-of-speech tagged words, so that a word now consists of a pair (ω, c) of a word ω and part-of-speech category c . Note that it would also be possible to bootstrap part-of-speech categories, or work with soft assignments, such as would be provided by a probabilistic n -best part-of-speech chunker.

Their model then needs to generate the category along with the word, and they factor this in the obvious way for a fixed signature, category and split of the word into prefix/suffix:

$$p(\omega, c) = p(s) \cdot p(\sigma|s) \cdot p(\tau|s, \sigma) \cdot p(c|s, \tau)$$

They drop the dependence on σ in predicting the category, which seems natural. Again, they fail to encode the multinomials in the model properly and fail to consider alternative analyses in assigning probabilities $p(\omega, c)$.

Other than this additional encoding step, their model remains the same. Yu et al.’s search routine is even more elaborate and heuristic than that described in the previous section, though it is based on the same underlying principle of clique-based signature collapsing. For instance, high frequency words are left alone, the words in “reliable” signatures are replaced with a token indicating their signature, and all other words are normalized to “dummy” symbols. This effectively defeats any attempt at solving the MDL

equations given above, and thus defeating any attempt to characterize the search space for solutions.

More interestingly, Yu et al. introduce left and right context information. That is, they create term vectors (word frequency counts) for both the words immediately to the left and immediately to the right of a word. They then apply the inverse-document-frequency (IDF) heuristic to scale the vectors and then compare them by dot product. These contextual elements are used to compare context similarity before collapsing signatures.

Given that they’ve explicitly avoided building MDL solutions with their heuristics, it is unclear why they present their evaluations in MDL terms. Clearly much better MDL results would arise from dropping all the heuristics. They do not provide any descriptive or statistical evaluation of segmentation performance under the new scheme, leading one to conclude that perhaps this new model performed worse than the baseline without syntactic features.

7.4.34 Jordan et al., 2005

Jordan, Healy and Keselj (2005) provided an entry for the 2005 MorphoChallenge. They extract counts for all character subsequences in the training data. It is not clear whether they used corpus frequencies for the words. They then take a maximum likelihood multinomial estimator of character sequences given the corpus counts. They then use this estimator to split a word W into two subsequences $W = M_1^* \cdot M_2^*$, where:

$$\langle M_1^*, M_2^* \rangle = \operatorname{argmax}_{\langle M_1, M_2 \rangle} p^{\text{ml}}(M_1) \cdot p^{\text{ml}}(M_2)$$

A word W is only split if there is a pair $\langle M_1^*, M_2^* \rangle$ such that:

$$p^{\text{ml}}(M_1^*) \cdot p^{\text{ml}}(M_2^*) > p^{\text{ml}}(W)$$

Once a word is split into a pair of substrings, each of those substrings is itself considered for splitting. Thus the algorithm is recursive.

Jordan et al.’s splitting condition imposes a very strong requirement, as reflected by the resulting system’s evaluated performance, which shows much higher precision than recall.

Presumably this algorithm weights its counts by frequency. Otherwise, it would be almost impossible to ever split words, as their stem subparts would never have high enough frequencies. As is, the splitting condition is very conservative.

For running on unseen words, the new word could always be added to a dynamic word list and the maximum likelihood estimates readjusted. Otherwise, any unseen word would have a zero maximum likelihood estimate and hence must be split into some pair of substrings. Another alternative would be some kind of smoothing for the subsequence probabilities.

7.4.35 Keshava and Pitler, 2005

Keshava and Pitler (2005) provided the winning entry for English in the 2005 MorphoChallenge. They apply a heuristic system motivated by Harris's (1955) successor/predecessor variety heuristic. In particular, they count substrings to get maximum likelihood estimates probability $p^{\text{ml}}(c|X)$ of the conditional probability of character c following character sequence X . They do this in the forward (p_f) and backward (p_b) direction and then break $X_1c_1c_2X_2$ into parts X_1c_1 and c_2X_2 if $p_f^{\text{ml}}(c_1|X_1)$ is near 1.0 and $p_f^{\text{ml}}(c_2|X_1c_1)$ is less than 1. The backward direction is defined symmetrically. They then consider as morphemes any substring that occurs as a morpheme in the analysis of a word more than a constant times more than the number of times it is not a morpheme in a word.

In order to tune their entry for the MorphoChallenge, Keshava and Pitler remove any morpheme that can be analyzed as a sequence of shorter morphemes.

To carry out segmentation, Keshava and Pitler consider all possible matching suffixes. To break ties, they compare the conditional probabilities of characters at the break point. For example, `politeness` is correctly analyzed as `polite-ness` rather than `politenes-s` because $p(\text{n}|\text{polite})$ is less than $p(\text{s}|\text{politenes})$.

Keshava and Pitler used *Wall Street Journal*, a Linux dictionary file and the word frequency list from the MorphoChallenge, to collect approximately 185,000 words to use for training. The program extracted 808 prefixes and 987 suffixes.

Keshava and Pitler derive counts of morphemes from their corpus, with the top suffixes being:

`s ly ness, ing, ed, al, ism, less, ist, able`

with counts ranging from 24,351 for `s` to 1613 for `able`. The top prefixes are:

`un re, dis, non, over, mis, in, sub, pre, inter`

with counts ranging from 15,858 for **un** to 1189 for **inter**.

They cite an 80.9% balanced F-measure for English segmentation; the official results cite 76.8%, which is the highest reported for English in the competition.

7.4.36 Chan, 2006

Chan (2006) sets out to induce traditional linguistic paradigms from a corpus of unlabeled data. Inexplicably, Chan assumes that each suffix should belong to exactly one paradigm; this flies in the face of simple evidence from English where **-s** is used as both a plural marker on nouns and a third-singular present marker on verbs.

Chan first sets up a matrix D of stem/suffix counts where $D(R, S)$ is the count of stem R and suffix S . The dimensionality is the number of stems by the number of suffixes. A typical approach to latent paradigm induction would be something like latent semantic indexing (LSI). LSI reduces the dimensionality of the representation of the stem/suffix matrix by factoring it into a product of three matrices, one of which represents the stems, one of which represents the suffixes, and one of which represents a square matrix in the latent dimensions.

Instead, Chan proposes two matrices L and M for stems and suffixes. The entry $L(C, R)$ is the conditional probability $p(C|R)$ of the class or paradigm C given the root or stem R . The entry $L(S, C)$ is defined as containing entries equal to the conditional probability $p(S|C)$ of a suffix given a paradigm. Multiplying these together produces a matrix $X = M \cdot C$ whose entries marginalize out the classes to produce the probability $p(S|R)$ of a suffix given a stem:

$$\begin{aligned} X(S, R) &= \sum_C L(S, C) \cdot M(C, R) \\ &= \sum_C p(S|C) \cdot p(C|R) \\ &= p(S|R) \end{aligned}$$

Because the result X contains only conditional estimates. An additional vector of marginal probabilities $p(R)$ for stems would be necessary to compute expectations for joint counts. This would then represent the original matrix D up to a multiplicative constant corresponding to the total corpus count.

Chan applies latent Dirichlet allocation (LDA) to impute the paradigms. LDA was stated in terms of documents and words. The basic model assumes

that for each document, a number of words is generated according to a Poisson distribution with mean λ , then a topic multinomial θ is sampled from a Dirichlet distribution with parameter vector α , then for each word, a topic z is sampled from the multinomial θ , and then given the topic z , a word is sampled from the multinomial conditional distribution $\beta(w|z)$ of words given topics (thus β is really a family of multinomials indexed by topic z). Given a prior on the Dirichlet and prior on β , Blei et al. use a combination of variational methods and EM to compute point estimates for the posterior Dirichlet α and multinomials $\beta(\cdot|z)$.

Chan’s analogy is that the document/word/topic triple of LDA may be used for the suffix/stem/paradigm triple of linguistic morphology. Working through the LDA model, this means that a suffix here is really a set of stems generated by a suffix-specific multinomial over paradigms. As Chan notes, this model accounts for the fact that not every possible suffix/stem pair shows up in the corpus.

It’s unclear what Chan actually did at this point. He claims that the posterior Dirichlet under LDA produces the matrix M , whereas it actually produces a vector, whose interpretation as the Dirichlet prior over a multinomial of paradigms is as counts over paradigms. It’s simply not a matrix. His analogy does hold through the fact that the posterior β will represent $p(R|C)$, the probability of a root given a category.

Chan goes on to claim that he uses the matrix of Dirichlet posteriors to do clustering, but this posterior is a vector, as we noted in the previous paragraph.

Chan’s diagrams indicate that he does divisive clustering on the set of suffixes in an attempt to generate clusters that correspond to traditional linguistic paradigms. Given the amount of category and sense ambiguity, it’s unlikely that LDA would produce paradigms that have this “natural” interpretation.

7.4.37 Hammerström, 2006

Hammarström (2006) seeks to extract a ranked list of suffixes for a language. Note that many of the approaches described above take this problem on as a subproblem.

Hammarström computes a very simple statistic over the suffixes of words given a corpus. He defines a statistic F over suffixes S by:

$$f(S) = \text{freq}(S) - E(\text{freq}(S))$$

where $\text{freq}(S)$ is the frequency of the suffix S in the corpus and the expectation is computed by assuming the characters were generated independently at random:

$$E(\text{freq}(S)) = \text{numWords} \cdot \hat{p}^{\text{len}(S)}$$

where \hat{p} is the inverse of the empirical entropy of the unigram character distribution in the corpus (roughly 19 for English), and numWords is the number of words in the corpus and $\text{len}(S)$ is the length of the suffix S in characters.

Hammarström actually uses the slope of the statistic F evaluated at the various suffixes S possible for a word $W = R \cdot c \cdot S$

$$g(R \cdot c, S) = \frac{f(S) - f(c \cdot S)}{|f(c \cdot S)|}$$

Hammarström plots the values of z for the various suffixes of **playing**, noting that $Z(\text{play, ing}) = 17$ was the highest scoring split, with $Z(\text{pla, ying})$ and $Z(\text{p, laying})$ being the next two highest splits, scoring around 4 each.

Hammarström aggregates the scores by adding them across words in the corpus, providing a final score for suffixes:

$$h(S) = \sum_{W=R \cdot S} g(R, S)$$

Finally, Hammarström ranks suffixes S by their aggregate score $h(S)$. For English, Hammarström produced the following suffix ranking, running over a one million token newspaper corpus:

-ing, -ed, -s, -'s, -ation, -es, -e, -er, -ers, -ting, -ly,
 -ations, -ted, -able, -ated, -al, -ness, -ling, -ent, -ating,
 -ate, -an, -ies, -ts, -ically, -ment, -led, -ering, -er's, -y

Hammarström suggests removing from the list any suffix that is not the best suffix for some word.

7.4.38 Karagol-Ayan et al., 2006

Karagol-Ayan et al. (2006) provide a semi-supervised approach to learning morphology that is based on optical character recognition (OCR) over dictionaries. Their method is based on the assumption that such dictionaries contain examples of usage, and a usage example will contain the dictionary entry's head word in a variant morphological form. By extracting pairs of such words, they produce training data for a stemmer. Their stemmer then

operates by stripping off affixes that are learned from the training data until a word in the dictionary is found or no more affixes may be stripped; this use of dictionaries makes the runtime morphology similar to (Krovetz 1993).

Matching of words is carried out using a variety of methods: longest common substring, edit distance, and a kind of prefix edit distance that does not penalize initial deletions. They then apply a sequence of heuristics to extract the match and then to extract the alignment using a deterministic extraction of an edit path from an edit; note that Hu et al. (2005) would need to invoke a similar array of heuristics to make their edit operations deterministic.

They evaluate their approach on Cebuano and Turkish. They also compare Linguistica, which does not use the aligned data, but merely a list of words. In the end, Linguistica appeared to work better, though both systems were reported in the 40-50% accuracy range.

7.4.39 Xanthos et al., 2006

Xanthos et al. (2006) explore alternative codebook definitions for the Linguistica framework (Goldsmith 2001). They note in their conclusion that they were surprised that the minimum description length (MDL) principle did not specify the exact form in which a grammar is expressed. Technically, Kolmogorov presented a universally applicable coding scheme for objects in terms of the length of the minimal program to generate the object. In practice, any assumption of coding scheme will simply bound the universal complexity from above; that is, the coding will most likely be improvable. In theory, any such coding scheme corresponds to a Bayesian prior, with the MDL solution being the MAP solution in Bayesian terms.

Xanthos et al. refer to codings of symbols as “pointers”, presumably by analogy to concrete computer implementations of model representations. They properly realize that coding a multinomial probability distribution with parameter vector θ requires its entropy number of bits to encode:

$$H(\text{Multi}(\theta)) = - \sum_i \theta_i \cdot \log_2 \theta_i$$

They further assume that θ will have a Zipf-like distribution, which they use to estimate this entropy.

Xanthos et al. need to code sets of strings in their paradigms/templates. Their baseline is a defective encoding that approximates the coding cost for a set of strings to be the set’s size times the entropy of the underlying distribution. A proper model would use the actual string likelihoods in the

multinomial for the encoding, and would also have to encode the model probabilities themselves. In addition, they would need to code the number of outcomes or a stop symbol.

Xanthos et al. go on to consider alternative “codings”. First, they consider encoding the complement of a set of stems relative to the universal set of all stems. They refer to this process as “polarization”. Otherwise, the encoding of each stem is as above.

Xanthos et al.’s last model codes a set as a bit vector, which they call a “binary string”. For some reason, they assume the entropy associated with a bit vector is the number of true bits that are set. Assuming any size set is equally likely and all choices of stems for the set are equally likely means that the number of bits required to code it is the total number of bits, not the just the positive ones.

According to a footnote, their reviewers pointed out an analogy to Rissanen’s (1989) combinatorial codes, which are just uniform distributions over all the ways to choose an m -element subset of a set of n elements. There are $\binom{n}{m}$ such sets, so each set has $1/\binom{n}{m}$ probability, and thus $\log_2 \binom{n}{m}$ entropy.

They evaluate on tiny 100,000 token corpora in Finnish and English, but only description length reductions due to the above heuristics, not any kind of accuracy.

Part III

Statistical Experiments

Chapter 8

Corpora

We begin our discussion of statistical experiments by providing an overview of the data corpora we used for various languages.

8.1 Tokenization

Tokens are generated using the tokenizer:

```
com.aliasi.tokenizer.IndoEuropeanTokenizer
```

8.2 Words

Words are defined to be tokens consisting of lowercased letters (e.g. **a-z**) of less than 32 letters.

In some cases, we restrict the training data to a set of tokens whose counts exceed a specified threshold.

8.3 English Gigaword

The University of Pennsylvania’s Linguistic Data Consortium (LDC) distributes the English Gigaword corpus (Graff 2003). It contains roughly 1.75 gigawords of English. The corpus contains four subcorpora, corresponding to data source. We consider the *New York Times* subcorpus, drawn from the eight year span July 1994 to June 2002. Documents were drawn from all parts of the newspaper, and some of the data is tabular, and some of it is editorial instructions. The corpus consists of 1.3 million documents, the titles and bodies of which contain 900M tokens.

<i>Count</i>	<i>Tokens</i>	<i>Count</i>	<i>Tokens</i>
≥ 1	336,277	≥ 16	3,855
≥ 2	214,727	≥ 32	2,255
≥ 4	144,767	≥ 65	1,229
≥ 8	106,438	≥ 131	611
≥ 16	82,461	≥ 262	283
≥ 32	66,068	≥ 524	134
≥ 64	53,733	$\geq 1,048$	76
≥ 128	43,634	$\geq 2,097$	36
≥ 256	34,753	$\geq 4,194$	18
≥ 512	26,808	$\geq 8,388$	8
$\geq 1,024$	19,701	$\geq 16,777$	5
$\geq 2,048$	13,927	$\geq 33,554$	1
$\geq 4,096$	9,386	$\geq 67,108$	0
$\geq 8,192$	6,118		

Figure 8.1: English Gigaword: Number of Words Given Minimum Counts

Qualitatively, the corpus contains a large number of misspellings. For example, the white grape `gewürztraminer` has two acceptable variants, `gewurztraminer` and `gewuerztraminer` (there are no umlauts in the gigaword corpora), as well as misspellings `gewrztraminer`, `gewuetztraminer`, `gewurzstraminer`, `gewurrtztraminer`, and `gewurtztraminer`.

There are a large number of “run on” words, that is, two words that were not separated in the corpus, such as `hikearizona`, `himto` and `ambassadorto`. The count of these is so high that we suspect conversion errors on the part of the LDC. The corpus also contains a large number of Spanish words, such as `muerto`, `librar`, and `nombrada`. Some words appear to be acronyms, such as `mscd`. Other words seem to be names of web sites, etc., such as `mrshowbiz` with 24 instances.

The `nyt-en` corpus contains 336,277 unique words according to the definition in the previous section. A table listing the number of words with counts above a given threshold is provided in Figure 8.1.

8.4 Counts and Independence

In the code distribution, we’ve provided various routines for providing count-based corpus analyses. These are quite straightforward plots of counts. They

include a Zipf plot of the raw words, stems (prefixes), suffixes and arbitrary substrings. They include counts weighted by word frequency and where each word is taken to have count 1.

8.5 Binomial Hypothesis Testing for Paradigms

Given the above counts, it is quite straightforward to generate groupings of suffixes based on their co-occurrence with stems. Given two suffixes τ_1 and τ_2 , we consider the counts of stems σ such that $\sigma\tau_1$ is a word and $\sigma\tau_2$ is a word. We count the number of strings σ that (a) precede both τ_1 and τ_2 in the corpus, (b) precede just τ_1 , (c) precede just τ_2 , and (d) precede neither.

	$+\tau_1$	$-\tau_1$
$+\tau_2$	p_{++}	p_{-+}
$-\tau_2$	p_{+-}	p_{--}

The hypothesis we test is that τ_1 and τ_2 are independent. This is just the standard binomial hypothesis test from the four-way table given above. The idea is that we would expect the probability of a stem σ attaching to both τ_1 and τ_2 to be the product of the probabilities of attaching to τ_1 and the probability of attaching to τ_2 .

We let $p_{+?} = p_{++} + p_{+-}$ be the probability of a stem attaching to τ_1 ; $p_{?+}$ is defined similarly:

$$t = \frac{N \cdot (p_{++} - p_I)}{\sqrt{N \cdot p_I \cdot (1 - p_I)}} = \sqrt{N} \frac{p_{++} - p_I}{\sqrt{p_I \cdot (1 - p_I)}}$$

where N is the corpus size and p_I is the probability of a stem applying to both τ_1 and τ_2 under the assumption that τ_1 and τ_2 are independent. The probability p_I is defined by:

$$p_I = (p_{++} + p_{+-}) \cdot (p_{++} + p_{-+})$$

where $p_{++} + p_{+-}$ is the probability of a stem combining with τ_1 and $p_{++} + p_{-+}$ is the probability of a stem combining with τ_2 . Note that the denominator, $N \cdot p_I \cdot (1 - p_I)$ is just the variance of N samples drawn from a binomial distribution with probability p_I ; taking the square root scales it to a deviation. Note that confidence in non-independence grows with the square root of the number of samples.

We found the results were better when words were not weighted by their frequency counts. The top 10 pairs for Portuguese are listed in figure 8.2. As noted by Hammerström (2006) and others, these methods are not sensitive

τ_1	τ_2	t
ou	ando	314
ou	aram	310
ou	ava	295
ndo	ram	287
ar	ou	284
rem	sse	280
ou	arem	276
ado	ou	276
ou	avam	274
ndo	rem	273

Figure 8.2: Portuguese: Least independent pairs of suffixes

τ_1	τ_2	τ_3	t
ssem	rão	riam	2516
ssem	riam	rmos	2408
rá	ssem	riam	2312
sse	ssem	riam	2308
ssem	rão	rmos	2303
sse	ssem	rão	2280
rá	ssem	rão	2268
sse	ssem	rmos	2266
sse	rão	riam	2252
sse	rá	ssem	2218

Figure 8.3: Portuguese: Least independent triples of suffixes

to shared material, thus it’s likely to see **aram** and **ram** as suffixes in lists like these.

We extend these tests to triples, quadruples, etc. of stems by considering the least independent pairwise evaluation of subsets.

Finally, we are able to evaluate single suffixes similarly by considering whether their appearance in suffix position is much more likely than we would expect given the overall distribution of the sequence. For a given suffix τ , we consider $p_{\text{suff}}(\tau)$, the probability of a word having τ as a suffix versus the expected likelihood if τ showed up by chance, which we represent

as $p_{\text{substr}}(\tau)$, the substring probability of τ . The hypothesis test is then:

$$t = \frac{N \cdot (p_{\text{suff}}(\tau) - p_{\text{substr}}(\tau))}{\sqrt{N \cdot p_{\text{substr}}(\tau) \cdot (1 - p_{\text{substr}}(\tau))}}$$

The top suffixes by this ranking in Portuguese are:

s, o, a, os, e, as, m, es, do, am, r, te, ão, em, ia, mos, da, ar, nte

Chapter 9

Concatenative Stem/Suffix Model

In this chapter, we apply the expectation/maximization (EM) algorithm to the problem of estimating a statistical stemmer given only a corpus of unlabeled text data in a target language. With EM, the segmentation is treated as missing data, and inferred along with the maximum likelihood model parameters.

9.1 Model Overview

9.1.1 Stemming Model

In section 9.2, we introduce a naive generative model of words, which we take to be composed of independently generated stems and suffixes. Given such a model, it is straightforward to estimate the joint probability of a word being segmented into a specified stem and suffix. These joint probability estimates may be conditionalized to estimate the probability that a word has a given suffix. These conditional probability estimates are used in the expectation step in the EM algorithm.

In section 9.3, we elaborate the structure of the stem and suffix models. Stems and suffixes are generated with independent models of roughly the same structure. First, the length of the stem or suffix is generated from their length distributions, both of which are Poisson distributions with adjustments for zero length cases. Then the specific characters are using independent character language models for stems and suffixes.

9.1.2 Maximum Likelihood Model Estimation

In section 9.5, we show how to extract the maximum a posterior (MAP) likelihood model given complete training data consisting of words segmented into stem and suffix. The various priors effectively smoothe the training data, allowing the resulting models to apply to unseen data.

9.1.3 Expectation-Maximization (EM) Algorithm

In this paper, we assume that we are given training data consisting of a set of words drawn from a language. The analysis of a word into a stem and a suffix can be treated as missing training data. We use the expectation/maximization (EM) algorithm to infer a posterior mode (locally maximum likelihood model) for the combination of model parameters and missing data. Starting with a random model, EM alternately (a) estimates expected values for the missing stem/suffix data and then (b) uses the estimates to train a maximum likelihood model, terminating when the estimates no longer change significantly.

9.2 Naive Suffixation Model

The naive model of stemming generates a word $\omega = \sigma\tau$ by generating its stem σ according to distribution p_r and then independently generating its suffix τ according to distribution p_s ; mnemonically, r is for root and s for suffix. The joint probability of generating a word $\omega = \sigma\tau$ with stem σ and suffix τ defined to be:

$$p(\sigma, \tau) = p_r(\sigma) \cdot p_s(\tau)$$

For instance, $p(\text{eat}, \text{ing})$ is the probability of the word **eating** generated with stem **eat** and suffix **ing**.

We calculate the probability $p(\omega)$ of the word ω by marginalization, which sums over each possible analysis of the word into stem and suffix.

$$p(\omega) = \sum_{\{ \sigma, \tau \mid \omega = \sigma\tau \}} p(\sigma, \tau)$$

For example, the marginal probability of the word **runs** is:

$$p(\text{runs}) = p(\text{runs}, \epsilon) + p(\text{run}, \text{s}) + p(\text{ru}, \text{ns}) + p(\text{r}, \text{uns})$$

The conditional probability of a stem σ given a word $\omega = \sigma\tau$ is derived by dividing the probability of the stemming by the marginal probability of

the word:

$$p(\sigma, \tau | \omega) = \frac{p(\sigma, \tau)}{p(\omega)}$$

For example, the probability that the word **runs** has stem **run** and suffix **s** is:

$$p(\mathbf{run}, \mathbf{s} | \mathbf{runs}) = \frac{p(\mathbf{run}, \mathbf{s})}{p(\mathbf{runs}, \epsilon) + p(\mathbf{run}, \mathbf{s}) + p(\mathbf{ru}, \mathbf{ns}) + p(\mathbf{r}, \mathbf{uns})}$$

Because of this explicit renormalization during the computation of conditional estimates, we may use unnormalized probabilities $p(\sigma, \tau)$ in the computation — any multiplicative constants simply cancel out.

The conditional probability estimates induce a first-best stemmer that returns the most likely stem:

$$\begin{aligned} \text{stem}(\omega) &= \operatorname{argmax}_{\sigma : \sigma\tau = \omega} p(\sigma, \tau | \omega) \\ &= \operatorname{argmax}_{\sigma : \sigma\tau = \omega} \frac{p(\sigma, \tau)}{p(\omega)} \\ &= \operatorname{argmax}_{\sigma : \sigma\tau = \omega} p(\sigma, \tau) \end{aligned}$$

The notation $\operatorname{argmax}_{\sigma : \sigma\tau = \omega}$ indicates the value of σ which together with the τ for which $\sigma\tau = \omega$ produces the largest value of the following expression. As in the computation of the conditional probability, we may use unnormalized estimates for $p(\sigma, \tau)$ in the first-best maximization.

For instance, we have $\text{stem}(\mathbf{runs}) = \mathbf{run}$ if $p(\mathbf{run}, \mathbf{s})$ is larger than $p(\mathbf{runs}, \epsilon)$, $p(\mathbf{ru}, \mathbf{ns})$, and $p(\mathbf{r}, \mathbf{uns})$.

9.3 Stem and Suffix Models

We model the stem and suffix probabilities generatively. First, we generate a length for the stem or suffix, then we generate the specified number of characters. Letting $\omega = c_0 \cdots c_{n-1}$ be an n character sequence, we define the root probability distribution p_r in terms of a length distribution $p_{\text{rlen}}(n)$ and a character distribution p_{rchar} that depends on n .

$$p_r(c_0 \cdots c_{n-1}) = p_{\text{rlen}}(n) \cdot p_{\text{rchar}}(c_0 \cdots c_{n-1} | n)$$

Given n , the conditional probabilities should sum to one:

$$\sum_{c_0, \dots, c_{n-1}} p_{\text{rchar}}(c_0 \cdots c_{n-1} | n) = 1$$

The suffix distribution is defined with the same general model:

$$p_s(c_0 \cdots c_{n-1}) = p_{\text{slen}}(n) \cdot p_{\text{schar}}(c_0 \cdots c_{n-1} | n)$$

Continuing our running example:

$$\begin{aligned} p(\mathbf{run}, \mathbf{s}) &= p_r(\mathbf{run}) \cdot p_s(\mathbf{s}) \\ &= p_{\text{rlen}}(3) \cdot p_{\text{rchar}}(\mathbf{run} | 3) \cdot p_{\text{slen}}(1) \cdot p_{\text{schar}}(\mathbf{s} | 1) \end{aligned}$$

9.3.1 Stem and Suffix Length Models

The stem and suffix length models are slightly different, both modifying the behavior of Poisson distributions for length zero outcomes.

We begin with stems. We rule out zero-length stems by offsetting a Poisson distribution¹ by one:

$$p_{\text{rlen}}(n) = \text{Poisson}(n - 1 | \lambda_r - 1)$$

The model parameter λ_r is the average length of a root. Zero-length roots are not allowed, so $p_{\text{rlen}}(0) = 0$, and the resulting distribution is proper.

We allow zero-length suffixes, but roughly follow Jansche (2003) in modeling them separately with probability ζ :

$$p_{\text{slen}}(n) = \begin{cases} \zeta & \text{if } n = 0 \\ (1 - \zeta) \cdot \text{Poisson}(n - 1 | \lambda_s - 1) & \text{otherwise} \end{cases}$$

The parameter λ_s is the average length of non-zero-length suffixes. Including zero-length suffixes, average suffix length is $(1 - \zeta) \cdot \lambda_s$.

9.3.2 Character String Models

We model the character probabilities for roots, $p_{\text{rlen}}(c_0 \cdots c_{n-1} | n)$, and for suffixes, p_{slen} , using the same class of character language models.

We will be employing bounded character language models which take into account the known length constraint n by generating a final end-of-string character we will write as EOS. Further, we will assume that the first

¹Recall that for non-negative integers $n \geq 0$, the Poisson distribution is defined by a single parameter λ :

$$\text{Poisson}(k | \lambda) = \frac{\lambda^k \cdot e^{-\lambda}}{k!}$$

The mean and variance of $\text{Poisson}(k | \lambda)$ are both λ .

character is generated from the context consisting of the fixed begin-of-word character, which we write as **BOS**. Thus we take:

$$\begin{aligned} p(c_0 \cdots c_{n-1} | n) &= \frac{p(c_0 \cdots c_{n-1} \cdot \text{EOS} | \text{BOS})}{\sum_{c'_0, \dots, c'_{n-1}} p(c'_0 \cdots c'_{n-1} \cdot \text{EOS} | \text{BOS})} \\ &\propto p(c_0 \cdots c_{n-1} \cdot \text{EOS} | \text{BOS}) \end{aligned}$$

The unnormalized estimates are too low because of the non-zero probability assigned to **EOS** in non-final position as well as the probability assigned to characters other than **EOS** in final position. Computationally, the proportional value suffices for computing the conditional probability estimates for the EM algorithm.

We apply the chain rule to calculate the joint probability in terms of the conditionals:

$$\begin{aligned} &p(c_0 \cdots c_{n-1} \cdot \text{EOS} | \text{BOS}) \\ &= p(\text{EOS} | \text{BOS} \cdot c_0 \cdots c_{n-1}) \cdot \prod_{i < n} p(c_i | \text{BOS} \cdot c_0 \cdots c_{i-1}) \end{aligned}$$

We further make the n th-order Markovian (i.e. n -gram) assumption:

$$p(c_i | c_0 \cdots c_{i-1}) = p(c_i | c_{i-n+1} \cdots c_{i-1})$$

which restricts the conditioning contexts to the previous $n - 1$ characters.

9.4 Training Data

The known training data consists of two aligned sequences: a sequence of words $word_i$ and a sequence of non-negative finite weights $weight_i$. The missing training data consists of a prefix $stem_i$ of $word_i$ indicating the stem of the word. Note that the word and the stem together uniquely define the suffix $suff_i$. Together, the known and the missing data make up the complete data.

An example of a three-word training set is shown in figure 9.1. In general, the same word may show up with non-zero weight with two different stems. Such training sets arise either with ambiguity in training data or uncertainty arising from using expected values of the missing data in the

9.5 Maximum a Posteriori Model Selection

Our models involve parameters which we collectively denote ϕ and a complete set of data we denote \mathbf{z} . The likelihood function $p(\mathbf{z} | \phi)$, which indicates

	<i>Known Data</i>		<i>Missing Data</i>
i	$word_i$	$weight_i$	$stem_i$
0	run	1.00	run
1	talking	2.10	talk
2	sleeper	0.72	sleep

Figure 9.1: Known and Missing Components of the Complete Training Data

how likely the data \mathbf{z} is under the model ϕ , was specified in the previous section. Before estimation, we establish a prior distribution $p(\phi)$ over the parameters; we discuss our actual priors below.

Our goal is to estimate given our training data \mathbf{z} , the probability $p(stem, suffix|\mathbf{z})$ of a given word being composed of a specified stem and suffix. The standard Bayesian inference involves integrating out our posterior uncertainty in the parameters ϕ :

$$p(stem, suffix|\mathbf{z}) = \int p(stem, suffix|\phi) \cdot p(\phi|\mathbf{z}) d\phi$$

In words, the integral averages over parameter settings ϕ , weighting them by their posterior likelihood $p(\phi|\mathbf{z})$. Unfortunately, this approach is far too computationally intensive given the number of parameters in our models and their structure.

Instead of integrating, we replace the probability weighted average over parameters expressed by the integral with a point estimate derived from the most likely ϕ given the priors and training data; this value is called the maximum a posterior (MAP) estimate:

$$\begin{aligned} \operatorname{argmax}_{\phi} p(\phi|\mathbf{z}) &= \operatorname{argmax}_{\phi} \frac{p(\phi) \cdot p(\mathbf{z}|\phi)}{p(\mathbf{z})} \\ &= \operatorname{argmax}_{\phi} p(\phi) \cdot p(\mathbf{z}|\phi) \end{aligned}$$

Because this leaves us with a single maximum a posteriori model set of model parameters ϕ , we simply use these parameters for (approximate) inference.

$$p(stem, suffix|\mathbf{z}) \approx p(stem, suffix|\operatorname{argmax}_{\phi} p(\phi|\mathbf{z}))$$

A family of distributions forms a conjugate prior for a likelihood model if the posterior distribution $p(\phi|\mathbf{z}) \propto p(\phi) \cdot p(\mathbf{z}|\phi)$ over parameters for any data \mathbf{z} falls into the same model family as the prior distribution $p(\phi)$. In all the

cases we consider here (binomial, multinomial and Poisson), the conjugate prior amounts to starting with a set of counts for outcomes. The values indicate the starting point and the count indicates how strong the prior is. Proper priors, that is ones with finite integrals over the space of parameters, provide non-zero counts for all outcomes.

9.5.1 Estimating Length Parameters

Three length parameters are estimated as part of the model. $\hat{\lambda}_r$ estimates the average length of roots, $\hat{\lambda}_s$ the average length of non-zero length stems, and $\hat{\zeta}$ the probability of a stem being length zero.

Estimating Poisson Length Parameters

The Poisson length components are estimated the same way for stems and suffixes, using the class of gamma distributions as conjugate priors. A gamma prior involves two parameters which can be thought of as specifying a prior average sample length and a prior number of samples.

The Poisson length estimates only depend on the length of segments being trained and their weights. For stems, the lengths are given by $|stem_i|$ and their weights $weight_i$; for suffixes, lengths are given by $|suff_i|$. For prior count $priorCount$ and prior root length $priorLen_r$ and prior suffix length $priorLen_s$, the maximum likelihood posterior point estimates are:

$$\hat{\lambda}_r = \frac{priorCount \cdot priorLen_r + \sum_{i < k} weight_i \cdot |stem_i|}{priorCount + \sum_{i < k} weight_i}$$

$$\hat{\lambda}_s = \frac{priorCount \cdot priorLen_s + \sum_{i < k} weight_i \cdot |suff_i|}{priorCount + \sum_{i < k} weight_i}$$

In words, we add the weighted total of data word lengths to the weighted prior average length, dividing by the prior count plus the data count. The larger the parameter $priorCount$, the the more training samples will be required to move the posterior estimate of the mean away from the prior estimate $priorLen_s$.

For the case listed above for estimating stem average length $\hat{\lambda}_s$:

$$\sum_{i < 3} weight_i \cdot |stem_i| = 1.0 \cdot 3 + 2.1 \cdot 4 + 0.72 \cdot 5 = 15$$

and

$$\sum_{i < 3} weight_i = 1.0 + 2.1 + 0.72 = 3.82$$

If the prior count is 1 and the prior average length is 2, then

$$\hat{\lambda}_s = \frac{1 \cdot 2 + 15}{1 + 3.82} \approx 3.53$$

If the prior count was 100 instead of 1, the posterior point estimate would be $(100 \cdot 2 + 15)/(100 + 3.82) = 2.07$, which is much closer to the prior length of 2.

Estimating Zero-Length Stem Probabilities

The probability ζ of a zero-length stem may be estimated as a binomial from a set of training data. The conjugate prior distributions for binomials are beta distributions. These essentially take the form of a prior zero probability estimate *priorZeroProb* and a number of samples *priorCount*. The prior maximum likelihood estimate is thus *priorZeroProb*.

Given our training data, the posterior maximum likelihood estimate of the probability of a stem being length zero is:

$$\hat{\zeta} = \frac{\text{priorCount} \cdot \text{priorZeroProb} + \sum_{i < n} \text{weight}_i \cdot \text{I}(\text{word}_i = \text{stem}_i)}{\text{priorCount} + \sum_{i < n} \text{weight}_i}$$

where I is an indicator function taking the value 1 if the boolean argument is true and 0 otherwise. In this case, it is true if the suffix is of length zero, or as written, if the stem is the same as the word.

As before, the strength of the prior is determined by the prior count. For the example data we have:

$$\begin{aligned} & \sum_{i < 3} \text{weight}_i \cdot \text{I}(\text{word}_i = \text{stem}_i) \\ &= 1.0 \cdot \text{I}(3 = 3) + 2.1 \cdot \text{I}(7 = 4) + 0.72 \cdot \text{I}(7 = 5) \\ &= 1.0 \cdot 1 + 2.1 \cdot 0 + 0.72 \cdot 0 \\ &= 1 \end{aligned}$$

and as before:

$$\sum_{i < 3} w_i = 3.82$$

If the prior zero count is 2 and the prior zero probability is 0.5, then the maximum likelihood posterior estimate is:

$$\hat{\zeta} = \frac{2 \cdot 0.5 + 1}{2 + 3.82} = 0.344$$

Had the prior count been larger, the posterior maximum likelihood estimate would be closer to the prior estimate. For instance, if the prior count was 200, then the posterior estimate would be $(200 \cdot 0.5 + 1) / (200 + 3.82) = 0.495$.

9.5.2 Estimating Character Language Models

We estimate character language models by mixtures of maximum likelihood models of different orders.

The maximum likelihood conditional character estimate \hat{p}^{ml} of order i over a given set of training data is defined by:

$$\hat{p}^{\text{ml}}(c_i | c_1 \cdots c_{i-1}) = \frac{\text{freq}(c_1 \cdots c_{i-1})}{\text{freqExt}(c_0 \cdots c_{i-1} \cdot c_i)}$$

with $\text{freq}(c_i \cdots c_j)$ indicating the frequency of the character sequence as a subsequence of the training data, and

$$\text{freqExt}(c_0 \cdots c_{i-1}) = \sum_{c'} \text{freq}(c_0 \cdots c_{i-1} \cdot c')$$

providing the count of the number of times the character sequence was extended by one character in the training data.

The character language models are structured as parametric interpolations over maximum likelihood estimators for each context size.

$$\begin{aligned} \hat{p}(c_i | c_{i-n+1} \cdots c_{i-1}) &= \lambda(c_{i-n+1} \cdots c_{i-1}) \cdot \hat{p}^{\text{ml}}(c_i | c_{i-n+1} \cdots c_{i-1}) \\ &\quad + (1 - \lambda(c_{i-n+1} \cdots c_{i-1})) \cdot \hat{p}(c_i | c_{i-n+2} \cdots c_{i-1}) \end{aligned}$$

with the interpolation ratio $\lambda(c_{i-n+1} \cdots c_i)$ depending on the estimation context of $n - 1$ previous characters.

We ground the recursion by interpolating with a uniform estimate:²

$$\hat{p}(c_i) = \lambda() \cdot \hat{p}^{\text{ml}}(c_i) + (1 - \lambda()) \cdot \text{Uniform}(c_i | \text{numChars})$$

where numChars is the number of unique characters in our training and test sets.

² Recall that the discrete uniform distribution over M elements is defined by:

$$\text{Uniform}(x | M) = 1/M$$

The interpolation parameters are defined as a parametric version of Witten-Bell interpolation (see, e.g., Carpenter 2004):

$$\lambda(c_0 \cdots c_{n-1}) = \frac{\text{freqExt}(c_0 \cdots c_{n-1})}{\text{freqExt}(c_0 \cdots c_{n-1}) + \kappa \cdot \text{numExt}(c_0 \cdots c_{n-1})}$$

where $\kappa \geq 0$ is a free hyperparameter determining the degree of smoothing, and where the number of extensions of a character sequence is defined by:

$$\text{numExt}(c_0 \cdots c_{n-1}) = |\{c \mid \text{freq}(c_0 \cdots c_{n-1} \cdot c) > 0\}|$$

Note that for any $\kappa \geq 0$, the resulting mixture estimates are maximum likelihood estimates for the model class because each component is a maximum likelihood estimator. **Is this right?** The value of κ determines the relative weighting between individual maximum likelihood estimates and lower order estimates. In the limit, the result is either the maximum context maximum likelihood estimate or the uniform estimate:

$$\lim_{\kappa \rightarrow 0} \hat{p} = \hat{p}^{\text{ml}}$$

$$\lim_{\kappa \rightarrow \infty} \hat{p} = \text{Uniform}(c_i | \text{numChars})$$

Quantization

The character language model package in LingPipe only accepts integer-valued training data. In order to work with floating point weights, we multiply all weights by a large multiplier and round off. By multiplying the interpolation ratio by the same factor, estimates work out to be the same as if floating point values of up to the multiplier level of accuracy were allowed. Note that the size of the multiplier will define the arithmetic accuracy of language model estimation.

Mixture Interpolation vs. Hierarchical Empirical Bayes

Rather than interpolating, we could have taken a more standard hierarchical Bayesian approach in which each character estimator is a multinomial with a Dirichlet prior, which like our other priors, specifies a prior number of counts per outcome. We could then either take the order $n - 1$ estimator as a prior for the order n estimator, or we could take an “empirical Bayes” approach following (MacKay and Peto 1994) in which an optimal prior is approximated. Like the approach in this paper, MacKay and Peto take point estimates for the priors rather than averaging over their uncertainty

in proper Bayesian fashion; as usual, this is because effective sampling is too expensive computationally. In previous experiments, the differences between the approaches has been minimal and the Witten-Bell interpolation model is by far the easiest to implement.

Informative Priors

Even in the Witten-Bell setting, we could begin with a collection of prior counts for any of the estimators. These would then combine with the data counts to produce a posterior estimate. If we provided every unigram outcome a nonzero count (a proper Dirichlet prior), we could eliminate the interpolation with the uniform estimate. In practice, the uniform estimate plays almost no role with reasonable settings of the hyperparameter κ . For instance, with 500,000 outcomes and 26 characters, and $\kappa = 8$ (a reasonable setting), the ratio of weight applied to the uniform estimates are only $1 - 500000 / (500000 + 8 \cdot 26) = 0.00041$. A more interesting approach would be to set prior counts for low-order n -grams based on a large general language corpus.

9.6 Expectation Maximization (EM)

In this section, we apply the expectation maximization (EM) algorithm to estimating the parameters of our model from data in which the segmentations into stem and suffix are unknown. Recall that our complete data consists of three sequences, *word*, *stem* and *weight*, where a triple $word_i$, $stem_i$ and $weight_i$ is taken as a training instance. Although we only know the *word* and *weight* data ahead of time, we may use EM to infer the most likely values for *stem*.

We have shown how to find the parameters ϕ which maximize $p(\phi | word, stem, word)$. We have also shown how to find the expected stem likelihoods given a model and a word $p(stem | \phi, word)$. This is all we need for EM.

The EM algorithm for our case is sketched in figure 9.2. To make this sketch more specific, we initialize the parameters ϕ to the prior maximum likelihood, which amounts to a model that has not seen any training data. Because our priors are proper, the result will be well defined. Often EM is run with multiple initial parameterizations to explore the possibility of multiple modes in the data, which amount to locally maximal likelihood estimates of ϕ that are not global maximums.

In the expectation step, we're computing our best guess for stems and their weights given our current parameter estimates ϕ . In the maximization

Initialization: Guess initial parameter estimates ϕ .

Repeat

- (a) *Expectation:* Compute the expected stem weights from the words, word weights and current parameter estimates:

$$E(\text{word}_i, \text{stem}|\phi) = \text{weight}_i \cdot p(\text{stem}|\text{word}_i, \phi)$$

- (b) *Maximization:* Reset ϕ to the maximum likelihood estimate using the expected stem counts as training weights:

$$\phi := \operatorname{argmax}_{\phi'} p(\phi'|\{\text{word}, \text{stem}, E(\text{word}, \text{stem}|\phi)\})$$

until the estimate ϕ converges.

Figure 9.2: EM Algorithm for Naive Stem Model

step, we use these expected stem weights to train a new model.

We continually alternate expectation and maximization steps until the model parameters ϕ converge. Because there are so many parameters, we instead measure the complete log data probability estimate:

$$\log_2 p(\mathbf{z}|\phi) = \sum_{i, \text{stem}} \text{weight}_i \cdot \log_2 p(\text{word}_i, \text{stem}|\phi)$$

As usual in numerical methods, we measure relative convergence. That is, given a threshold δ , we iterate until the difference between the value under the previous parameters, ϕ^{old} converge to those in the reestimated parameters, ϕ :

$$\frac{|\log_2 p(\mathbf{z}|\phi) - \log_2 p(\mathbf{z}|\phi^{\text{old}})|}{|\log_2 p(\mathbf{z}|\phi)| + |\log_2 p(\mathbf{z}|\phi^{\text{old}})|} < \delta$$

For instance, if $\delta = 0.001$, we iterate until the log probability changes by less than 1/10th of 1 percent in an iteration. We have to be careful here not to set the threshold below our arithmetic accuracy or we may never converge. Our arithmetic accuracy is in part determined by the quantization approximation for language model estimation.

In figure 9.3, we show the data log likelihood as a function of iteration, as well as the average word and character likelihoods.

In figure 9.4, we show the estimates for a few words at several different iterations. For this example, we used a word count threshold of 32, produc-

	<i>Iteration</i>				
	1	2	4	8	16
<i>Total</i>	-3,461,500	-2,041,800	-1,899,500	-1,835,300	-1,822,500
<i>Word</i>	-52.392	-30.905	-28.750	-27.778	-27.585
<i>Char</i>	-6.4081	-3.7800	-3.5165	-3.3976	-3.3740

Figure 9.3: Cross-Entropies by Iteration: Corpus, Word Rate and Character Rate

<i>Word</i>	<i>Stem</i>	<i>Suffix</i>	<i>Expected Count in Iteration</i>				
			1	2	4	8	16
eats	e	ats	.074	.032	.003	.000	.000
	ea	ts	.297	.486	.130	.000	.000
	eat	s	.297	.411	.852	.995	.996
	eats	€	.332	.071	.015	.005	.004
runners	r	unners	.004	.000	.000	.000	.000
	ru	nners	.038	.004	.000	.000	.000
	run	ners	.153	.077	.002	.000	.000
	runn	ers	.307	.632	.959	.940	.904
	runne	rs	.307	.220	.006	.000	.000
	runner	s	.123	.056	.032	.059	.095
	runners	€	.069	.009	.000	.000	.000

Figure 9.4: EM Estimation Step Examples

ing 66,068 total words. We set a prior stem average length to 4 and the prior suffix average non-zero length to 2.5, the prior count to 500,000 (effectively turning off training for lengths) and the zero-length suffix probability to 0.20. We assumed 8-grams with an interpolation ratio of $\kappa = 8$, and a character set of size 26.

Continuing the example, note that even in the first iteration, the estimated counts are not uniform. This is because of the length priors, which favor putting extra characters in the stem.

For any given stem or suffix, we may compute its expected count in the corpus by simply summing the expected counts contributed by each word.

$$E(\text{count}(\text{stem}) | \text{weight}, \text{word}, \phi) = \sum_i \text{weight}_i \cdot p(\text{stem}(\text{word}) = \text{stem})$$

Consider again the expectations shown in figure 9.4. For example, in the second iteration, the word **runners** contributes 0.056 count for **s** and a

<i>Iteration</i>									
1		2		4		8		16	
<i>Stem</i>	<i>Count</i>	<i>Stem</i>	<i>Count</i>	<i>Stem</i>	<i>Count</i>	<i>Stem</i>	<i>Count</i>	<i>Stem</i>	<i>Count</i>
€	5394	€	6852	€	10410	€	12621	€	13684
s	1741	ing	3995	ed	5972	ed	6288	s	6600
ing	1498	ed	3805	ing	5435	s	6175	ed	6293
ed	1495	s	2430	s	4524	ing	5508	ing	5508
ng	1149	es	2408	es	3874	e	4436	e	5139
es	1112	er	1658	er	2421	es	4259	es	4321
e	784	ers	1347	e	2342	er	2424	er	2403
d	747	e	1001	ers	1843	ers	1784	ers	1759
er	711	y	729	y	1166	y	1373	y	1436
rs	549	ly	668	ly	838	ly	943	ly	1045
ers	540	ts	664	al	662	ion	813	ion	953
y	513	ted	579	ness	652	o	702	a	728
r	450	ness	487	able	570	ness	653	o	712
g	439	tion	478	ation	559	a	647	ness	672
ts	435	ng	475	o	488	al	617	able	585
t	422	able	424	ts	477	able	586	ic	506
ly	414	al	422	ion	461	ation	572	al	497
ted	375	ion	417	ic	452	ic	550	ions	497
n	368	rs	412	a	389	ions	456	ation	467
ion	356	d	384	os	388	ity	404	ity	429

Figure 9.5: Top Suffixes Ordered by Expected Counts

0.632 count for **ers**; the word **eats** in the second iteration contributes a 0.411 count for **s** and no count for **ers**. In figure 9.5, we show the top 5 stems and top 25 suffixes by count for the specified iterations. Note that as the probabilities become more sharply attenuated, the counts for dominant suffixes become greater. If we had set our termination condition to require greater relative accuracy, the trend would continue.

Some examples of stretches of wordsw with their analysis are shown in figure 9.6.

	amphora.		genuflect.	question.
zig.s	amphora.e		genuflect.ing	question.able
zigzag.	amphoteric.in		genuflect.ion	question.ably
zigzagg.ed	ampl.e	box.	genuin.e	questionair.e
zigzagg.ing	ampl.ia	boxcar.	genuin.ely	question.ed
zigzag.s	ampli.ar	boxcar.s	genuin.eness	question.er
zilch.	amplificat.ion	boxcutt.ers	genu.s	question.ers
zill.ion	amplifi.ed	box.ed	geo.	question.ing
zillionair.e	amplifi.er	box.er	geobiolog.ist	questionnair.e
zillionair.es	amplifi.ers	box.ers	geocach.ing	questionnair.es
zill.ions	amplifi.es	box.es	geochem.ist	question.s
zin.	amplify.	box.es	geochemistry.	quest.s
zinc.	amplify.ing	boxful.	geocit.ies	queu.e
zin.e	ampli.o	box.ing	geodes.ic	queu.ed
zin.es	amplitud.e	box.like	geograph.er	queue.ing
zinfandel.	amp.ly	boxscor.e	geograph.ers	queu.es
zinfandel.s	amp.s	boxwood.	geograph.ic	queu.ing
z.ing	amputat.e	boxwood.s	geograph.ical	qui.
zing.ed	amputat.ed	box.y	geograph.ically	quibbl.e
zing.er	amputat.ing	boy.	geograph.ies	quibbl.ed
zing.ers	amputat.ion	boycott.	geograph.y	quibbl.es
zing.ing	amputat.ions	boycott.ed	geolog.ic	quibbl.ing
zing.s	ampute.e	boycott.ers	geolog.ical	quich.e
zing.y	ampute.es	boycott.ing	geolog.ically	quich.es
zinn.ia	am.s	boycott.s	geolog.ist	quick.
zinn.ias	amstephens.on	boyfriend.	geolog.ists	quicken.
zip.	amtrak.	boyfriend.s	geolog.y	quicken.ed
zipp.ed	amuck.	boyhood.	geomagnet.ic	quicken.ing
zipp.er	amulet.	boyish.	geometr.ic	quicken.s
zipper.ed	amulet.s	boyish.ly	geometr.ical	quick.er
zipp.ers	amus.e	boyish.ness	geometr.ically	quick.est
zipp.ier	amus.ed	boy.s	geometr.ics	quick.ie
zipp.ing	amus.ement	bozo.	geometr.ies	quick.ies
zipp.o	amus.ements	bozo.s	geometr.y	quick.ly
zipp.y	amus.es		geophysic.al	quick.ness
zip.s	amus.ing		geophysic.ist	quicksand.
	amus.ingly		geophysic.ists	quicksilv.er
			geophysic.s	

Figure 9.6: Examples of first-best analyses after convergence

Chapter 10

Affixation Boundary Effects

10.1 Problem for Pure Suffix Model

The pure suffixation model fails to account for the morpho-phonemic boundary effects found in affixation. For instance, the present participle verb suffix **-ing** behaves differently depending on the form of the stem to which it attaches. For instance, consider the verbs **walk**, **run** and **pace**.

```
walk + -ing = walking
run + -ing = running
pace + -ing = pacing
```

With **walk**, the suffixation is just a pure suffixation. With **run**, the result **running** has an additional **n** inserted (an orthographic form of gemination). With **pace**, the result is **pacing**, in which the final **e** in the stem is dropped when combined with the suffix.

The pure suffixation approach has to decide where to draw the boundary. No problem for **walk.ing**, but what is the right analysis for **running**, **run.ning** or **runn.ing**? With the former, there is a spurious suffix, **ning**, whereas with the latter, there is a spurious stem, **runn**. In the EM approach derived in the last chapter, some words went one way, and some another, depending on the statistics. With **pacing**, the only sensible analysis is **pac.ing**, but the result is the spurious stem **pac**. In practice, the EM implementation tends to treat English verbs ending in **e** as if the **e** in the base form is itself a suffix, thus **pace** is analyzed as **pac.e**, thus sharing a stem with **pac.ing**. The problem with this is that it induces false positives; for instance, **caring** (the verb) and **cars** (the plural noun) wind up erroneously sharing the stem **car**.

10.2 Modeling Boundary Effects

Wicentowski’s (2004) word frame model captures the intuition that **walking**, **pacing**, and **running** share the suffix **-ing**, while deriving the stems **walk**, **pace** and **run**. Recall that the model allows a “boundary effect”. The boundary effect for **running** is a doubling of **n**, whereas the boundary effect for **pacing** is the dropping of an **e**.

Our model follows Wicentowski’s in taking the context for estimating boundary effects to be the full suffix and the last letter of the stem. For instance, $p(\mathbf{nn}|\mathbf{n}, -\mathbf{ing})$ is the probability of a final **n** being realized as **nn** when followed by the suffix **-ing**, $p(\mathbf{k}|\mathbf{k}, -\mathbf{ing})$ is the probability of **k** being realized as itself, and $p(\epsilon|\mathbf{e}, -\mathbf{ing})$ is the probability of dropping a final **e** (we use ϵ to represent the empty string).

We limit realizations of the final stem character to zero, one or two surface characters. The first component of the model is a global estimate of this three-valued discrete distribution; this is set as a parameter rather than estimated online. Like stem/suffix length, simple maximum likelihood with no prior tends to fall into a degenerate solution where all final characters are dropped, because shorter stems are cheaper to encode all else being equal.

The actual realization probabilities are then treated as conditional on the length:

$$\hat{p}(\sigma|\rho, \tau) = \hat{p}_{len}(\text{len}(\sigma)) \cdot \hat{p}_n(\sigma|\rho, \tau)$$

where the final term is really three distributions, \hat{p}_0 , \hat{p}_1 and \hat{p}_2 for the three different possible outcome lengths.

We also place a prior on the \hat{p}_i distributions which favors exact matching in \hat{p}_1 and favors duplication in \hat{p}_2 . These priors make sense as we know ahead of time that these are more likely outcomes. The priors are fairly weak, and wind up dominated by the data, while still keeping the solutions out of degenerate local maxima. This allows us to overcome some of the sparse data problem associated with inducing these mappings from scratch. We further restrict the length two outputs to have the same initial character as the length one inputs. Thus **m** could map to **mm**, or it might map to **me**, but it could not map to **nn** or even **nm**. This restriction may need to be relaxed for languages with more complex affixation boundary effects.

The underlying EM procedure does not change, but the space of possible analyses grows enormously. Rather than simply considering the n break points for a word of length n , the algorithm must consider all possible characters as possible substitutions or insertions.

10.3 Results

Overall, results are disappointing in that the boundary effects learned through EM seem to be rather random. While **zipped** is properly analyzed as **zip+ed**, **zoology** is analyzed **zoologi+y**, presumably because the prefix **zoologi** is so common in words like **zoological**.

Insertions are very difficult to generate. Making them too likely in the distribution of outcomes causes degenerate behavior, whereas leaving them at realistic prior settings finds them undergenerating. For instance, **spaced** is analyzed as **spac+ed**, and **spacing** is analyzed as **spac+ing**; the three words with root **space** (**space**, **spacemen**, **spaces**, are just not enough to tip the balance in favor of the insertion.

10.4 Conclusions

Although Wicentowski (2004) clearly demonstrates these models are learnable from supervised training data, we have not been able to find settings that generate linguistically “natural” results, nor even settings that generate morphologically meaningful clusters.

Part IV

**Future Directions and
Conclusions**

Chapter 11

Generalized Sequence Model

The general solution to the morphology problem must go beyond a simple stem, suffix and boundary effect model. The simplest generalization is to simply allow stemming to iterate. Thus rather than analyzing a word as a stem plus a suffix, a word may consist of a root and any number of prefixes and suffixes. In order to handle compounding, as found in many Germanic languages and other agglutinative languages such as Turkish, it is also necessary to allow multiple stems. This basic view of morphology has a long tradition in linguistics, going under the name “item and arrangement”.

The simplest way to generalize what we have done now is to simply allow a stem and any number of additional stems, suffixes and prefixes. For instance, the word **unfortunately** might be analyzed as a prefix **un-**, a stem **fortune**, a suffix **-ate**, and a suffix **-ly**. As soon as we allow boundary effects, we begin to blur the traditional boundary between “item and arrangement” (pure concatenation) and the completely general “item and process” approach under which affixes have completely general effects on the stems to which they apply.

Computationally, things become more difficult. Instead of being able to simply enumerate the break points and possible boundaries, every subset of break points needs to be considered. Because this leads to an exponential growth in number of full-string analyses, some sort of dynamic programming is required. The traditional solution for a purely concatenative solution would be to use the forward-backward algorithm. If we do not discriminate stems from affixes, a simple one-state hidden Markov model would suffice. We could even use the states of the HMM to naturally partition the results into states for prefixes, suffixes, and stems (and possibly modifying stems if the model will produce a head stem and modifiers for compounding).

To account for boundary effects, a customized HMM-like forward-backward algorithm would be needed. This could be naturally implemented in a generalized weighted finite-state transducer toolkit, or implemented directly without too much trouble.

The only problem is that it is even less likely that a reasonable solution will be found by purely unsupervised methods. Luckily, it is possible to use semi-supervised methods with HMMs or other generative models (e.g. see Nigam et al.'s 1998 survey of Naive Bayes and EM for semi-supervised classification).

Chapter 12

Recursive Categorical Model

The final model we discuss is the most general, and the closest to that used by semi-theoretical linguists. In a recursive model, each stem will have a part-of-speech or other category tag, and each suffix or affix will have a modifying effect. For instance, **dine** may be tagged as a verb and **-er** as a suffix that maps verbs to nouns (e.g. **diner**).

The recursive model assumes that each word may be analyzed hierarchically as a sequence of compounding, prefixation or suffixation operations. Each word is either an atomic stem, or is the result of compounding two words, attaching a prefix to a word, or attaching a suffix to a word. For instance, **unhappily** may be analyzed as the prefix **un-** plus the word **happily**, which in turn may be analyzed as the word **happy** plus the suffix **-ly**. With a hierarchical model, we are able to distinguish the two analyses of **unfortunately**, corresponding to $((\text{un}, (\text{fortune}, \text{ate})), \text{ly})$ (in an unfortunate manner), and to $(\text{un}, ((\text{fortune}, \text{ate}), \text{ly}))$ (not in a fortunate manner).

This basic view of morphology was first put forward in this form by Hoeksema (1985). Such a generatively powerful model is frowned upon these days in the theoretical literature, but still widely used among computational linguists; if the phrase structure has a particular form, the lexical component of Koskenniemi's two-level morphology can be viewed in this way.

By adding boundary effects to the model, the result is a kind of phrase-structure analysis of a word with corresponding morphological effects (see, e.g., Mastroianni and Carpenter 1994), which bears a strong similarity to the kind of surface realization models found in head-driven phrase structure grammar (HPSG) (Reape 1989).

Morphology has a syntactic effect on the part-of-speech tag of the result.

Ideally, this should be part of the model, as the part-of-speech tag provides a great deal of information about possible and likely patterns of word formation. Although possible to use local co-occurrence information to infer syntactic relatedness, it is nearly impossible to extract useful part-of-speech information in an unsupervised manner. The closest anyone has come is to cluster words that appear in similar contexts. For instance, modeling the distribution of words before a word and the distribution of words after a word and using a distance such as symmetricized Kullback-Liebler distance or Shannon-Jensen divergence.

Theoretically, the sequence models underlying HMMs may be generalized to the context-free models underlying context-free grammars. A full EM approach is still possible using the inside-outside algorithm (see, e.g., Lari and Young 1990). In practice, the computations become much more expensive, but perhaps still feasible with careful coding. The real problem is convergence to an interesting model. Results in pure inducing CFGs have not been promising, but results have been much better in the case of semi-supervised input (see Pereira and Schabes 1992).

As noted by Paice (1994), it is difficult to draw the line in practice between historical morphology and active morphology in a language. For instance, the word **happy** was introduced in the early 14th century as a derived form of the noun **hap** with the adjectival suffix **-y**. What's confusing is that **-y**-suffixation is still productive in modern English, with the same meaning as it had in the 1300s; for instance, we might describe a computer as **Windowsy** or **Macintoshy**. As another example, the word **ugly** is historically derived from the stem **ugg** and adverbial suffix **-ly**.

A further cause of confusion is morphological ambiguity. The word **unhappily** actually has two valid derivations. In one, it is **unhappy** plus suffix **-ly**, whereas in the other, it is prefix **un** plus **happily**. The former is a manner adverb meaning roughly "done in an unhappy manner", whereas the latter derivation means roughly "not done in a happy manner". This is a subtle, yet nevertheless real distinction, just as in syntactic attachment. Even more difficult cases involve expressions such as **Montague grammarian**, which seems to involve the root **Montague grammar** and the suffix **-ian**.

Chapter 13

Annotation Tool

Although a fully automated approach to morphology induction is tantalizingly promising, we feel the current state of the art demands a semi-supervised approach.

For the Indo-European languages, the distribution of tokens follows a Zipf distribution (inverse power law). In figure 13.1, we show the distributions for several languages with Western scripts, as reported in the Leipzig corpus distribution (Quasthoff et al. 2006). This strongly skewed distribution shows that annotating the top 1000 entries, which takes almost no time at all and handles most of the highly irregular cases, will be clearly worthwhile. We believe that annotating anywhere from 10,000 to 100,000 examples by hand and bootstrapping the rest of the analysis in a semi-supervised fashion is the most promising direction for future work.

These numbers are highly correlated with average word form length, as shown in figure 13.2, also derived from the Leipzig corpora distribution (Quasthoff et al. 2006).

To support the annotation task, we developed a simple annotation graphical user interface with which it is quite easy to tag over 500 words/hour, or 4000 words/day, or 20,000 words/week. We believe one week’s manual tagging effort will increase accuracy from the 70% range into the mid-to-high 90% range, depending on language, as demonstrated for the most sophisticated model in the literature, Wicentowski’s word frame model (Wicentowski 2004).

A screen shot of the tool is provided in figure 13.3. It shows the word to annotate at the top. Below the word as title is a list of radio buttons used to select the type of the analysis. For instance, in the screen shot, the analysis “suffix” is selected, indicating the word is composed of a stem

<i>Language</i>	<i>100</i>	<i>1K</i>	<i>10K</i>	<i>100K</i>
Catalan	46.2%	64.9%	86.4%	97.7%
English	42.8%	67.0%	91.3%	99.3%
Finnish	20.3%	36.8%	61.5%	85.4%
French	47.7%	68.0%	89.6%	n/a
German	38.0%	58.2%	78.5%	92.8%
Italian	41.9%	63.4%	86.5%	98.6%
Swedish	39.2%	59.4%	79.9%	94.7%
Turkish	19.2%	37.7%	65.7%	91.0%

Figure 13.1: Out-of-Vocabulary Rates for top n Words in Various Languages

<i>Language</i>	<i>Length</i>
Catalan	8.46
English	8.49
Finnish	11.74
French	8.24
German	12.35
Italian	8.89
Swedish	10.33
Turkish	9.16

Figure 13.2: Average Word Form Length in Various Languages

plus a suffix. Based on the word type, fields to the right are available for entering the components. For the word **years**, that’s a stem **year** and suffix **s**. Note that the stem and suffix do not need to form the word by pure concatenation; for radical cases (e.g. **are** stemming to **be** or **shake** stemming to **shook**, there is an “irregular” analysis type). Annotators may also mark a word as a non-word to correct errors in corpus extraction.

To help with the annotation, the application downloads the dictionary definition from **dictionary.com** over the web and displays it in HTML. This is often useful for exploring the etymology of words whose morphological structure is unclear.

The analyzer is designed to perform a single analysis of a word at a time. Thus a word like **ridiculously** would be analyzed as stem **ridiculous** and suffix **-ly**, rather than a triple of **ridicule** + **-ous** + **ly**. To make sure full analyses are carried out for every word, all stems generated in the tool must

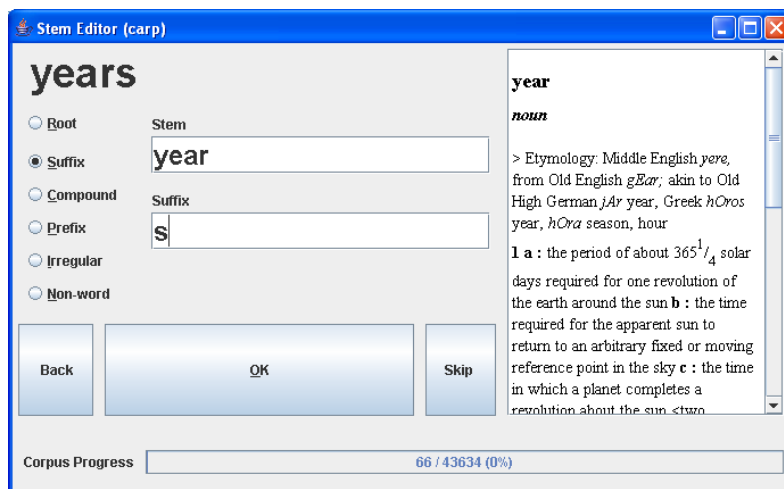


Figure 13.3: Screen shot of editor tool

be added to the set of words to analyze.

Once the analysis of a word is complete, the annotator clicks the “OK” button, which is nice and fat in compliance with Fitts’ law (time to move the mouse to an area is inversely proportional to target size). Annotators may skip entries to return to later and there’s a “Back” button in order to navigate the previous analysis stack; this is useful because click brainos are not uncommon when annotating at speed. At the very bottom is a list of progress on a corpus.

The annotator is instrumented to populate analyses with the automatic stem+suffix analyses produced by our unsupervised stemmers. Elements to annotate are chosen in order of frequency. With active learning, these would be chosen by a risk estimate, which we would define by:

$$\text{risk}(\sigma) = \text{count}(\sigma) \cdot (1 - p(\text{firstBest}(\sigma)|\sigma))$$

This is just the expected number of errors arising from analyzing the word σ . Here $\text{firstBest}(\sigma)$ is meant to indicate the first-best analysis of σ , and the probability term is the model’s probability estimate for the first best analysis

given the input. Thus one minus this number is the model's estimate of the chance of error in assigning the first-best analysis to σ . This is all then scaled by the count of σ , yielding the expected number of errors arising from the word. Note that this count may be scaled to a probability to make the risk figures comparable across corpora of different counts.

Every time a word is analyzed as anything other than a root or non-word, its count is added to the stem count. Note that compounds include two stems and thus **basketball** would add its count to **basket** and **ball**, as getting the full analysis right depends on getting both subanalyses correct.

Training and decoding the general model is straightforward given training data (see, e.g., Wicentwoski 2004).

Chapter 14

Conclusions

For information retrieval applications, simple affix-stripping and compounding algorithms work as well, on average, as more sophisticated morphologically sound approaches. Without further detailed evaluations in the context of particular languages and particular applications, it's unclear whether the more sophisticated stem-and-suffix sensitive EM methods developed here will actually improve performance.

Our overall recommendation is to use multiple stems, if possible, along with the original word in a TF/IDF setting. This kind of indexing is supported by Apache Lucene's `TokenStream` interface to indicate that full forms and their stem occupy the same token position. In this way, the high recall of a stemming approach is preserved, while also preserving most of the precision benefits of full-form indexing. These benefits hold not only for single-word queries, but also for phrasal queries. The only drawback is index size and query retrieval speed.

For developing high quality morphological analyses of languages, including highly sensitive stemmers, we believe labeled training data is necessary. Given the low cost of annotating 20,000 examples (one person week), it seems an obvious win for accuracy under any approach. From Wicentowski's (2004) word frame model, it appears this may result in a very high accuracy stemmer (in the mid 90% accuracy range across a variety of languages).

In terms of supervision, the "unsupervised" methods work without labeled training data, but require a more subtle form of input: tuning by a statistical computational linguist sensitive to the facts of the language at hand. Because we didn't know any of the target languages other than English, the actual performance figures for our work in non-English languages is subject to finding evaluation data and tuning the various EM parameters.

References

- Adamson, G. and J. Boreham. 1974. The use of an association measure based on character structure to identify semantically related pairs of words and document titles. *Information Storage and Retrieval* **10**.
- Ando, R. K. and Lillian Lee. 2000. Mostly unsupervised statistical segmentation of Japanese: application to Kanji. *NAACL 2000*.
- Antworth, Evan L. 1990. PC-KIMMO: a two-level processor for morphological analysis. *Occasional Publications in Academic Computing* **16**. Summer Institute of Linguistics, Dallas.
- Argamon, Shlomo, Navot Akiva, Amihoud Amir and Oren Kapah. 2004. Efficient unsupervised recursive word segmentation using minimum description length. *Proceedings of COLING*.
- Baroni, Marco, Johannes Matiassek, and Harald Trost. 2002. Unsupervised discovery of morphologically related words based on orthographic and semantic similarity. *ACL SIGPHON Workshop*.
- Bauer, L. 1983. *English Word Formation*. Cambridge University Press.
- Baayen, R. R. Piepenbrock, and L. Gulikers. 1995. The CELEX lexical database. Release 2. CD-ROM.
- Bernhard, Delphine. 2005. Unsupervised morphological segmentation based on segment predictability and word segments alignment. In *Proceedings of MorphoChallenge 2005*.
- Bordag, Stefan. 2005a. Unsupervised knowledge-free morpheme boundary detection. *Proceedings of RANLP 05*.
- Bordag, Stefan. 2005b. Two-step approach to unsupervised morpheme segmentation. *Proceedings of MorphoChallenge 2005*.
- Brent, Michael. 1999. An efficient, probabilistically sound algorithm for

- segmentation and word discovery. *Machine Learning* **34**:71–106.
- Buckwalter, Tim. 2004. Issues in Arabic orthography and morphology analysis. *Proceedings of the Workshop on Computational Approaches to Arabic Script-based Languages, COLING 2004*.
- Carpenter, Bob. 2004. Scaling language models to gigabytes. In *Proceedings of the ACL Software Workshop*.
- Chan, Erwin. 2006. Learning probabilistic paradigms for morphology in a latent class model. *ACL SIGPHON Workshop*. 69–78. Brooklyn.
- Chomsky, Noam and Morris Halle. 1968. *The Sound Patterns of English*. Harper and Row.
- Church, Kenneth and Patrick Hanks. 1990. Word association norms, mutual information, and lexicography. *Proceedings of ACL 27*, 76–83.
- Church, Kenneth Ward. 1995. One term or two? *ACM SIGIR 18*:310–318.
- Cretuz, Mathias and Krista Lagus. 2002. Unsupervised discovery of morphemes. *ACL SIGPHON Workshop*, 21–30.
- Cretuz, Mathias and Krista Lagus. 2005. Inducing the morphological lexicon of a natural language from unannotated text. In *Proceedings of the International and Interdisciplinary Conference on Adaptive Knowledge Representation and Reasoning*, 106–113.
- Dalrymple, M., R. Kaplan, L. Karttunen, K. Koskenniemi, S. Shaio, and M. Westcoat. 1987. Tools for morphological analysis. *CSLI Technical Report CL-87-108*. Stanford University.
- Dawson, J. 1974. Suffix removal and word conflation. *ALLC Bulletin*.
- Déjean, Hervé. 1998. Morphemes as necessary concepts for structure discovery from untagged corpora. In *NeMLaP3/CoNLL98 Workshop on Paradigms and Grounding in Natural Language Learning*, 295–299.
- Deligne, Sabine and Frédéric Bimbot. 1997. Inference of variable-length linguistic and acoustic units by multigrams. *Speech Communication* **23**:223–241.
- Dempster, A. P., N. M. Laird, D. B. Rubin. 1977. Maximum likelihood from incomplete data via the EM algorithm. *Journal of the Royal Statistical Society, Series B* **39**:1–38.
- Dhillon, I. S., S. Mallela, and D. S. Modha. 2003. Information-theoretic

co-clustering. Technical Report TR-03-12, Dept. of Computer Science, U. Texas at Austin.

Dzeroski, S. and T. Erjavec. 1997. Induction of Slovene nominal paradigms. In *Proceedings of the 7th International Workshop on Inductive Logic Programming*.

Namer, Fiammetta and Pierre Zweigenbaum. 2004. Acquiring meaning for French medical terminology: contribution of morphosemantics. In *Proceedings of Medinfo. 2004* **11**:535–539.

Freitag, Dayne. 2005. Morphology induction from term clusters. *Proceedings of CoNLL*.

Gaussier, Éric. 1999. Unsupervised learning of derivational morphology from inflectional lexicons. *ACL 1999 Workshop: Unsupervised Learning in NLP*.

Gelman, Andrew, John B. Carlin, Hal S. Stern and Donald B. Rubin. 2003. *Bayesian Data Analysis (2nd Edition)*. Chapman and Hall.

Goldsmith, John. 2001. Unsupervised learning of the morphology of a natural language. *Computational Linguistics* **27**(2):153–198.

Goldwater, Sharon and Mark Johnson. 2005. Representational bias in unsupervised learning of syllable structure. *9th Conference on Natural Language Learning*:112–119. Ann Arbor.

Graff, David. 2003. *English Gigaword*. Catalog No. LDC2003T05. Linguistic Data Consortium, Philadelphia.

Hammarström, Harald. 2006. A naive theory of affixation and an algorithm for extraction. *SIGPHON Workshop at NAACL 2006*.

Hafer, Margaret A. and Stephen F. Weiss. 1974. Word segmentation by letter successor varieties. *Information Storage and Retrieval* **10**:371–385.

Hakkani-Tür, D. Z., G. Oflazer and G. Tür. 2000. Statistical morphological disambiguation for agglutinative languages. *Proceedings of COLING*.

Harman, Donna. 1991. How effective is suffixing? *Journal of the American Society for Information Science* **42**(1):7–15.

Harris, Zellig. 1955. From phoneme to morpheme. *Language* **31**:190–222.

Harris, Zellig. 1967. Morpheme boundaries within words: report on a computer test. *Transformation and Discourse analysis Papers 73*, University of

Pennsylvania.

Hoeksema, Jacob. 1985. *Categorial Morphology*. New York, Garland.

Hu, Yu, Irina Matveeva, John Goldsmith and Colin Sprague. 2005a. Refining the SED heuristic for morpheme discovery: a look at Swahili. *ACL Psycho- and Computational-Linguistics Workshop*.

Hu, Yu, Irina Matveeva, John Goldsmith and Colin Sprague. 2005b. Using morphology and syntax together in unsupervised learning. *ACL Workshop on Psycho-Computational Linguistics*.

Hull, D. A. 1996. Stemming algorithms: a case study for detailed evaluation. *Journal of the American Society for Information Science* **47**(1): 70–84.

Karagol-Ayan, Burcu, David Doermann, and Amy Weinberg. 2006. Morphology induction from limited noisy data using approximate string matching. *ACL SIGPHON Workshop*. 60–68. Brooklyn.

Jacquemin, Christian. 1997. Guessing morphology from terms and corpora. In *Proceedings of the ACM SIGIR Conference*. 156–165.

Jain, Anil K. and Richard Dubes. 1988. *Algorithms for Clustering Data*. Prentice-Hall.

Jain, Anil K., M. Narasimha Murty and P. J. Flynn. 1999. Data clustering: a review. *ACM Computing Surveys* **31**(3):264–323.

Jansche, Martin. 2003. Parametric models of linguistic count data. *41st Meeting of the Association for Computational Linguistics*.

Johnson, C. Douglas. 1972. *Formal Aspects of Phonological Description*. Mouton.

Johnson, Howard and Joel Martin. 2003. Unsupervised learning of morphology for English and Inuktitut. *ACL-HLT*.

Jordan, Chris, John Healy and Vlado Keselj. 2005. Swordfish: using ngrams in an unsupervised approach to morphological analysis. In Kurimo, Mikko, ed., *MorphoChallenge 2005*.

Karttunen, Lauri. 1983. KIMMO: A general morphological processor. *Texas Linguistic Forum* **22**: 165–186.

Karttunen, L., K. Gaál, and A. Kempe. 1997. *Xerox Finite-State Tool*. Xerox Research Center, Grenoble.

Karttunen, Lauri and Kenneth R. Beesley. 2001. History of two-level mor-

phology. *ESSLLI 2001*.

Kazakov, D. 1997. Unsupervised learning of naive morphology with genetic algorithms. *ECML/Mlnet Workshop on Empirical Learning of NLP Tasks*.

Kazakov, Dimitar and Suresh Manadhar. 2001. Unsupervised learning of word segmentation rules with genetic algorithms and inductive logic programming. *Machine Learning* **43**:121–162.

Kit, Chunyu and Yorkick Wilks. 1999. Unsupervised learning of word boundary with description length gain. *Conference on Natural Language Learning*.

Kleinberg, J. M. 1998. Authoritative sources in a hyperlinked environment. In *Proceedings of the ACM-SIAM Symposium on Discrete Algorithms*.

Koskenniemi, Kimmo. 1983. *Two-Level Morphology: A General Computational Model for Word-Form Recognition and Production*. Ph.D. thesis, University of Helsinki.

Krovetz, Robert. 1993. Viewing morphology as an inference process. *SIGIR* **16**:191–202.

Kurimo, Mikko, Mathias Creutz, Matti Varjokallio, Ebru Arisoy and Murat Saraclar. . 2005. Unsupervised segmentation of words into morphemes – Challenge 2005:an introduction and evaluation report. In *EU Pascal Challenge: Unsupervised segmentation of words into morphemes*.

Lari, K. and S. Young. 1990. The estimation of stochastic context-free grammars using the Inside-Outside algorithm. *Computer Speech and Language*, **4**:35-56.

Lennon, M., D. S. Pierce, B. D. Tarry, P. Willet. 1981. An evaluation of some conflation algorithms for information retrieval. *Journal of Information Science* **3**:177–183.

Ling, C. X. 1994. Learning the past tense of English verbs: the symbolic pattern associator vs. connectionist models. *Journal of Artificial Intelligence Research* **1**:209–229.

Lovins, J. B. 1968. Development of a stemming algorithm. *Mechanical Translation and Computational Linguistics* **11**.

Lovins, J. B. 1971. Error evaluation for stemming algorithms as clustering algorithms. *Journal of the American Society for Information Sciences* **22**:28-40.

- Luk, Robert W. P., K. F. Wong, K. L. Kwok. 2002. Hybrid term indexing: an evaluation. *Proceedings of the 3rd NTCIR Workshop*
- MacKay, David J. C., Linda C. Bauman Peto. 1994. A hierarchical Dirichlet language model. *Natural Language Engineering* **1**(3):1–19.
- MacWhinney, B. 1978. *The Acquisition of Morphophonology*. Monographs of the Society for Research in Child Development **43**.
- Manning, Christopher D. 1998. The segmentation problem in morphology learning. In *NeMLaP3/CoNLL98 Workshop on Paradigms and Grounding in Language Learning*, 299–305.
- Marchand, H. 1969. *The Categories and Types of Present-Day English Word Formation*. C. H. Beck.
- Martin, W., B. Al, and P. Sterkenburg. 1983. On the processing of a text corpus: from textual data to lexicographical information. In R. Hartman (ed.), *Lexicography, Principles and Practice*, 77–87. Academic Press.
- Mastroianni, Michael and Bob Carpenter. 1994. Constraint-based Morphophonology. In *Proceedings of the First ACL SIGPhon Workshop*. Las Cruces, New Mexico.
- Maxwell, Mike. 2002. Resources for morphology learning and evaluation. *LREC*.
- McNamee, Paul. 2003. Knowledge-light Asian language text retrieval at the NTCIR-3 workshop. *Proceedings of the 3rd NTCIR Workshop*.
- McCray, A., A. Browne, and D. Moore. 1988. The semantic structure of neo-classical compounds. In *SCAMC'88 – Proceedings of the 12th Annual Symposium on Computer Applications in Medical Care*, 165–168.
- Mikheev, A.. 1997. Automatic rule induction for unknown-word guessing. *Computational Linguistics* **23**(3):405–423.
- Monson, Christian. 2004. A framework for unsupervised natural language morphology induction. *2004 ACL Student Workshop*.
- Monson, Christian, Alon Lavie, Jaime Carbonell and Lori Levin. 2004. Unsupervised induction of natural language morphology inflection classes. *ACL SIGPHON*.
- Neuvel, Sylvain and Sean A. Fulop. 2002. Unsupervised learning of morphology without morphemes. In *ACL SIGPHON Workshop*, 31–40.

- Nevill-Manning, Craig G. and Ian H. Witten. 1997. Identifying hierarchical structure in sequences: a linear time algorithm. *Journal of Artificial Intelligence Research* **7**:67–82.
- Nie, Jian-Yun, Jiangfeng Gao, Jian Zhang and Ming Zhou. 2000. On the use of words and n-grams for Chinese information retrieval. *Proceedings for the 5th International Workshop on Information Retrieval with Asian Languages*.
- Oflazer, K., S. Nirenberg and M. McShane. 2001. Bootstrapping morphological analyzers by combining human elicitation and machine learning. *Computational Linguistics* **27**(1):59–84.
- Paice, C. D. 1996. Method for evaluation of stemming algorithms based on error counting. *Journal of the American Society for Information Science* **47**(8):632–649.
- Peng, Fuchun and Dale Schuurmans. 2001. A hierarchical EM approach to word segmentation. *6th NLP Pacific Rim Symposium*.
- Pereira, Fernando and Yves Schabes. 1992. Inside-outside reestimation from partially bracketed corpora. In *ACL*.
- Porter, Martin. 1980. An algorithm for suffix stripping. *Program* **14**(3):130–137.
- Quasthoff, U.; M. Richter; C. Biemann. 2006. Corpus Portal for Search in Monolingual Corpora, In *Proceedings of the Fifth International Conference on Language Resources and Evaluation*, LREC 2006, Genoa.
- Quasthoff, Uwe and Christian Wolff. 2002. The Poisson collocation measure and its applications. In *Proceedings of the Second International Workshop on Computational Approaches to Collocations*.
- Rayson, Paul and Roger Garside. 2000. Comparing corpora using frequency profiling. In *Proceedings of the ACL Workshop on Comparing Corpus*, 1–6.
- Reape, Michael. 1989. A Logical Theory of Semi-Free Word Order and Bounded Discontinuous Constituency. In *Proceedings of the EACL*. Manchester, England.
- Redlich, A. Norman. 1993. Redundancy reduction as a strategy for unsupervised learning. *Neural Computation* **5**:289–304.
- Rissanen, Jorma. 1983. A universal prior for integers and estimation by minimum description length. *Annals of Statistics* **11**(2):416–431.

- Rissanen, Jorma. 1989. *Stochastic Complexity in Statistical Inquiry*. World Scientific Publishing.
- Rumelhart, D. E. and J. L. McClelland. 1986. On learning the past tenses of English verbs. In J. L. McClelland and D. E. Rumelhart (eds.), *Parallel Distributed Processing: Explorations in the Microstructure of Cognition*, Volume 2, 216–271. MIT Press.
- Saffran, Jenny R., Elissa L. Newport, and Richard N. Aslin. 1996. Word segmentation: the role of distributional cues. *Journal of Memory and Language* **35**(4):606–621.
- Salton, Gerald and M. Lesk. 1968. Computer evaluation of indexing and text processing. *Journal of the ACM* **15**(1):8–36.
- Schone, Patrick and Daniel Jurafsky. 2001. Knowledge-free induction of inflectional morphologies. *NAACL 2001*.
- Schulz, Stefan, Martin Honeck and Udo Hahn. 2002. Biomedical text retrieval in languages with a complex morphology. In *Proceedings of the ACL Workshop on NLP in the Biomedical Domain*, 61–68.
- Schützenberger, M. P. 1961. A remark on finite transducers. *Information and Control* **4**:185–196.
- Snover, Matthew G. and Michael R. Brent. 2002. A probabilistic model for learning concatenative morphology. *NIPS 2002*.
- Sproat, Richard. 1992. *Morphology and Computation*. MIT Press.
- Su, K. M. Wul and J. Chang. 1994. A corpus-based approach to automatic compound extraction. *32nd ACL*. Las Cruces, NM.
- Swanson, Don. 19??.
- Theron, P. and I. Cloete. 1997. Automatic acquisition of two-level morphological rules. *Proceedings of the Fifth ANLP*, 103–110.
- Wicentowski, Richard. 2004. Multilingual noise-robust supervised morphological analysis using the WordFrame model. *ACL SIGPHON 2004*. Barcelona.
- Xanthos, Aris, Yu Hu, and John Goldsmith. 2006. Exploring variant definitions of pointer length in MDL. *ACL SIGPHON Workshop*. 32–40. Brooklyn.
- Xu, Jinxi and W. Bruce Croft. 1998. Corpus-based stemming using cooccur-

- rence of word variants. *ACM Transactions on Information Systems* **16**(1):
- Yarowsky, David and Richard Wicentowski. 2000. Minimally supervised morphological analysis by multimodal alignment. *Proceedings of ACL*, 207–216.
- Yarowsky, David, Grade Ngai, and Richard Wicentowski. 2001. Inducing multilingual text analysis tools via robust projection across aligned corpora. *Proceedings of HLT*, 161–168.
- Yu, Hua. 2000. Unsupervised word induction using MDL criterion. *ICSLP*. Beijing.
- Zipf, George Kingsley. 1968. *The Psychobiology of Language. An Introduction to Dynamic Philology*. Second Edition. (First edition, 1935) M. I. T. Press/Cambridge.
- Zipf, George Kingsley. 1949. *Human Behavior and the Principle of Least Effort*. Addison-Wesley.
- Zweigenbaum, Pierre and Natalia Grabar. 1999. Contribution of medical terminology to medical language processing: experiments in morphological knowledge acquisition from thesauri. In *Proceedings of IMIA Workshop on Natural Language Processing and Medical Concept Representation*.